

Unauthorized Image Generation, Paid Usage Drain, Context Failures, Duplicate Closures, Nine Days of Silence, and What “Faster and Unlimited” Really Means

From June 21 through July 20, 2026, I repeatedly filed formal reports with OpenAI concerning product defects in ChatGPT and the way its support organization handled those reports.

These were not subjective complaints about an image having the wrong mood or a sentence failing to match my preferred style.

ChatGPT launched the image-generation tool during conversations in which I had not requested an image. It did so again after I explicitly prohibited image generation. It prioritized stale context over current instructions, importing locations, objects, and entire work products from previous tasks. It generated the wrong aspect ratio despite exact numerical specifications. It stopped showing image-generation progress and sometimes failed to return completed results. When I refreshed the page to check whether the image had finished, ChatGPT displayed a “Too many requests” error and interfered with unrelated work.

When I reported these problems to OpenAI, cases were closed as “duplicates” without identifying where they had supposedly been merged. Support requested information that had already been submitted. It announced that specialists would respond “within a few days,” then provided no substantive response for nine days, eleven days, or twenty days. OpenAI could remain silent for weeks, while the customer was warned that the case would be closed unless they replied within one day.

This article documents that record using the emails I retained, the case numbers OpenAI assigned, the responses issued under the OpenAI Support name, and the language displayed on OpenAI’s pricing page.

I am not claiming to know the undisclosed technical cause of these failures, nor am I making a final legal determination. I am separating documented events, OpenAI's own explanations, and the questions raised by those records.

Once the record is assembled, however, the pattern cannot reasonably be reduced to one accidental bad answer.

The product failed to retain the user's current intent.

Support failed to retain the substance of the user's current complaint.

The product returned to obsolete context.

Support returned complaints to unrelated prior cases.

The product ignored explicit instructions and performed a different action.

Support ignored the complaint and issued an unrelated response.

The system intended to investigate the product appeared to reproduce the same failure modes reported in the product.

Case 1: I Did Not Request Image Generation. The Paid Feature Ran Anyway, and OpenAI Went Twenty Days Without a Substantive Response

On June 21, 2026, I formally reported that ChatGPT had repeatedly launched image generation during ordinary conversation, commentary, analysis, and requests to write text prompts for use in an image generator.

The most serious element was not the first erroneous invocation.

After the first incident, I explicitly prohibited image generation. ChatGPT itself acknowledged that it should not generate images during casual conversation, analysis, or prompt drafting, and that the tool should be used only when I clearly requested an actual image.

The same failure then occurred again.

This was not merely a matter of an unwanted picture appearing on the screen.

Image generation is subject to usage limits, rate limits, time constraints, and other restrictions. When the system generates an image the user did not request, the user may lose capacity and opportunities that should have remained available for intended work.

I requested an investigation into the tool-call history, usage consumption, cause of the erroneous invocation, possible restoration or compensation, permanent corrective measures, and preservation of the relevant records.

OpenAI's AI-assisted support system created Case Number 10337718. It characterized the report as a possible defect in which image generation had been triggered unintentionally and may have affected the user's available image-generation usage.

OpenAI then stated that the matter had been escalated to a support specialist and that a response was expected "within the next few days."

No such response arrived within the next few days.

On July 1, I requested the investigation status, the responsible department, the status of any usage review, whether compensation was being considered, and whether the issue had been shared with an engineering team.

By July 11, twenty days had passed without a substantive answer concerning the logs, usage consumption, technical cause, compensation, or prevention of recurrence.

On July 11, a response signed by Warner of OpenAI Support finally acknowledged that the issue was not about image quality. OpenAI stated that it had recorded the matter as an unintended invocation of the image-generation tool, including instances in which the user had explicitly instructed ChatGPT not to generate images.

Even then, OpenAI did not provide a concrete investigative finding.

It did not say whether the relevant logs had been reviewed.

It did not determine whether usage had been consumed.

It did not address compensation.

It did not identify the responsible team.

It did not provide a completion date.

Then the case moved backward.

On July 15, I replied that nothing had been resolved and explicitly instructed OpenAI not to close the case until meaningful corrective action had been taken.

The following day, OpenAI again asked me to provide:

A description of the problem.

Steps to reproduce it.

Screenshots.

A screen recording.

OpenAI also warned that if I did not respond within one day, the ticket would be closed.

Those matters had been explained repeatedly and in detail since June 21. Warner's own prior response had demonstrated that OpenAI understood the issue.

Nevertheless, the case had effectively been reset to a state preceding the initial intake.

On July 16, I formally objected to the failure to review the case history, the repeated demand for previously submitted information, and the imposition of a twenty-four-hour deadline on the customer.

On July 18, OpenAI sent essentially the same message again. It requested a description, reproduction steps, screenshots, and a recording, and again warned that the case would be closed if I failed to reply within one day.

OpenAI had taken twenty days to issue even a partial acknowledgment.

The customer was given one day to prevent closure.

OpenAI had already stated that the case remained under review, yet the burden of keeping that investigation alive was shifted to the person who had reported the problem.

This was more than a delayed reply. It raised questions about continuity of complaint handling, preservation of case history, internal handoffs, and whether any identifiable person or department had actually assumed responsibility for the investigation.

Case 2: I Reported a Context-Management Failure. OpenAI Support Then Lost the Context and Could Not Identify the Case It Had Merged

On July 12, 2026, I submitted a separate report concerning context management.

I expressly distinguished this report from the unauthorized image-generation case.

Within a long working thread, ChatGPT had lost track of the current work product and returned to an old “flower-arrangement video prompt” task, producing an entire video prompt that was no longer the subject of the conversation.

The thread also had an established working instruction: when I said “仕事しろ,” meaning “do the work,” ChatGPT was supposed to stop explaining, apologizing, or analyzing and immediately produce the corrected deliverable.

Despite repeated corrections, ChatGPT sometimes responded with explanations and meta-commentary rather than completing the active task.

The report stated that the issue concerned:

Management of conversational context.

Identification of the current editing target.

Improper return to obsolete context.

Failure to retain user-established working rules.

It explicitly stated that unauthorized image generation was being handled in a different case.

Approximately one minute after I sent the report, Case Number 11319313 was closed as a duplicate of an existing case.

I objected that the two cases concerned different defects.

OpenAI then created Case Number 11319385 and stated that it could not determine, from the email before it, which existing case had supposedly caused Case 11319313 to be classified as a duplicate.

OpenAI asked me to provide the number of the case into which it had allegedly merged my report.

OpenAI had made the duplicate determination.

OpenAI had closed the case.

OpenAI nevertheless claimed it could not identify the supposed duplicate.

I replied that OpenAI had made the decision and should organize its own cases.

Case 11319385 was then also closed as a duplicate.

A new Case Number, 11319521, appeared.

I asked OpenAI to identify the destination case.

OpenAI replied that the current thread could not access the destination case number. It instructed me to search my own inbox for prior OpenAI emails and provide the case number myself.

In other words, the system that performed the merge could not identify where the case had gone and attempted to outsource that search to the customer.

Only after I resubmitted a large volume of prior email history did OpenAI finally state that the destination was Case Number 10837305.

That did not resolve the problem.

A later “support summary” omitted the independent context-management complaint. It reduced the matter back to the separate allegation that ChatGPT had triggered image generation after being told not to.

I had repeatedly requested that the context-management failure be treated as a distinct issue.

By the end of the process, that distinction had disappeared from OpenAI’s own summary.

I reported that ChatGPT could not remember what it was currently editing.

OpenAI Support could not remember what problem it was currently handling.

I reported that ChatGPT had returned to the wrong prior context.

OpenAI Support returned the report to the wrong prior case.

The support process reproduced the structure of the product defect.

Case 3: I Requested an Investigation Into Tool Errors and Context Contamination. OpenAI Sent Instructions for Deleting Conversations and the Entire Account

On July 2, 2026, I reported two problems that occurred while using ChatGPT’s “High” reasoning setting.

First, ChatGPT launched image generation even though my request explicitly stated:

“This is not an image-generation request.”

“Output only the Japanese prompt text to be given to an image generator.”

Second, I instructed ChatGPT to work with a scene set in a geothermal-pipeline maintenance building in Reykjavík. Instead, it imported material from the previous task, which had been set at a lakeside glamping site near Mount Fuji.

The current location, environment, props, actions, and line of sight had been specified. The output nevertheless prioritized stale conversational material over the current input.

OpenAI’s first response explained that “High” referred to GPT-5.5 Thinking with extended reasoning.

That did not answer the report.

I had not asked what the label meant. I had asked why that setting had been associated with an explicit instruction violation, an unintended tool call, and contamination from obsolete context.

OpenAI requested the conversation URL and the time of the incident.

I provided the shared conversation link.

OpenAI then requested the approximate time.

I identified the affected period as approximately June 20 through the present.

OpenAI responded that the period was too broad and asked me to isolate a single execution time.

It also stated that if image generation had occurred, the resulting image should appear in the Images area, and that if the image was no longer wanted, deleting the conversation would also remove it from Images.

To be precise, OpenAI did not explicitly order me to destroy evidence.

The problem is contextual.

I was requesting an investigation into whether the tool had been invoked, whether paid usage had been consumed, and whether the related records would be preserved. In that setting, instructions explaining how to remove the relevant conversation and generated asset pointed in the opposite direction from evidence preservation.

I argued that OpenAI should be able to search its own records by account, conversation, date range, and tool type.

Shortly afterward, Case Number 10837305 was closed as a duplicate.

On July 12, I submitted another detailed message in the same broader thread. This time, I explicitly excluded image-tool invocation from the request and focused solely on context management, current-task identification, improper reversion to past context, and retention of working instructions.

OpenAI created Case Number 11319858.

The first response provided links explaining how to:

Export ChatGPT data.

Delete uploaded files and conversations.

Delete the OpenAI account.

The request had been for an investigation into context management.

The answer concerned data export, conversation deletion, and account deletion.

Approximately one minute later, the case was closed as a duplicate.

I had reported a product defect in which ChatGPT prioritized irrelevant context over the user's current request.

OpenAI Support then prioritized irrelevant deletion material over the complaint it had just received.

Again, the support system reproduced the reported defect.

Case 4: I Reported That ChatGPT Would Not Honor a 4:5 Aspect Ratio. Support Misidentified the Reference Image, Escalated the Case, and Then Went Silent for Nine Days

On July 11, 2026, I reported that ChatGPT's image generator was failing to honor specified aspect ratios.

My prompts do not merely contain the words "vertical" or "4:5."

I specify:

A width-to-height ratio of exactly 4:5.

Width divided by height equals 0.8.

Height divided by width equals 1.25.

Permitted output dimensions of 1024 by 1280, 1536 by 1920, 2048 by 2560, or another integer resolution with the identical ratio.

I prohibit 1:1, 3:4, 2:3, 9:16, and 16:9.

I prohibit generating at a different ratio and cropping afterward.

I prohibit automatic expansion, outpainting, and ratio changes made to accommodate the subject or background.

The generator nevertheless returned images at other aspect ratios.

For my use, aspect ratio is not an aesthetic preference. These images are produced for 4:5 promotional portraits, advertisements, and social-media assets. An image at the wrong ratio cannot be used, even if its content is otherwise good.

Each generation took approximately two minutes.

Every ratio failure therefore required another generation, another wait, another inspection, and another opportunity for the same failure to occur.

OpenAI's AI-assisted support requested the environment, model, actual prompt, generated image dimensions, reference-image dimensions, and approximate start date.

I provided:

Chrome.

GPT-5.6 Instant.

The full production prompt.

The period in which the problem had increased.

A reference image.

My reply clearly stated, "The reference image is attached."

OpenAI then treated the attached reference image as though it were the generated result that had violated the aspect ratio. It instructed me to inspect that image's pixel dimensions.

I corrected the record.

The attached file was the input reference image, not the failed output.

The failed images had been generated within OpenAI's own service. Their generation time, asset record, dimensions, model, canvas configuration, and tool history should have existed in OpenAI's systems.

I asked OpenAI to stop making an AI system that had misunderstood the attachment repeat the same questions. I requested escalation to a human support representative and a technical team capable of reviewing the relevant assets and logs.

At 3:45 a.m. on July 11, OpenAI stated that Case Number 11271585 had been escalated to a support specialist and that a response should arrive within the next few days.

As of July 20, nine days had passed.

No investigative result had arrived.

No status update had arrived.

OpenAI had not corrected its misidentification of the attachment.

It had not stated whether engineering had received the report.

It had not addressed compensation.

By that point, I had supplied the environment, the model, the full prompt, the reference image, and the timeframe. OpenAI had stated that a human specialist would review the matter.

The result was nine days of silence.

Case 5: Image-Generation Progress and Results Disappeared. Refreshing Triggered Restrictions. OpenAI Closed the Report as a Duplicate Within One Minute

At 2:51 a.m. on July 11, I reported a new failure that had appeared on or after July 10.

After image generation began, the browser no longer displayed progress.

There was no reliable way to determine whether the job had been accepted, was still processing, had stopped, had failed, or had completed.

Sometimes the finished image did not appear at all.

To determine whether anything had happened, I refreshed the page.

After repeated refreshes, ChatGPT displayed “Too many requests.”

The restriction then interfered not only with image generation but also with ordinary ChatGPT work.

My report set out the causal sequence:

1. Image generation is initiated.
2. No progress state is displayed.
3. No result appears after waiting.
4. The user refreshes the page to determine the status.
5. Repeated refreshes trigger a “Too many requests” restriction.
6. Unrelated work is then blocked.

I requested an investigation into any July 10 UI or deployment changes, frontend state synchronization, result delivery, generation and error logs, the relationship between refreshes and rate limiting, an official recovery

procedure, compensation, and escalation to the responsible technical team.

At 2:52 a.m., OpenAI created Case Number 11269846.

One minute after the original report, OpenAI closed the case as a duplicate.

It did not identify the destination case.

It did not state whether the missing progress indicator, failed result delivery, or rate-limit interaction had been added to any active investigation.

Because the report resembled an earlier case submitted approximately forty minutes before, it may have been merged into Case Number 11268075.

That is an inference, not a confirmed fact.

OpenAI never identified the destination.

The only explicit action was closure.

Nine days later, there had been no substantive response and no confirmation that the technical issues described in the report remained under investigation.

Case 6: OpenAI Treated the Missing-Image and Rate-Limit Problem as a Local Environment Issue, Escalated It, and Then Went Silent for Nine Days

At 2:10 a.m. on July 11, before Case 5 was filed, I had reported the underlying problem as Case Number 11268075.

I explained that I was using ChatGPT Pro in Chrome and operating two image-generation pages in parallel.

The image jobs appeared to complete, but the generated images did not load into the browser.

I refreshed the page every few minutes to see whether the results had arrived.

After repeated refreshes, ChatGPT displayed “Too many requests” and interfered with ordinary work.

I asked OpenAI to investigate:

Why completed images were not displayed.

Whether refreshes were counted as excessive requests.

The formal conditions under which Pro users receive the restriction.

The official policy on using two image-generation pages in parallel.

A browser-side correction that would eliminate the need to refresh.

OpenAI’s initial AI-assisted response recommended disabling extensions, VPNs, proxies, and secure DNS. It suggested using an incognito window or another browser.

It also advised waiting thirty to sixty seconds, selecting “Stop,” choosing “Regenerate,” and using a hard refresh if necessary.

That response contained a basic contradiction.

The complaint was that refreshing the page to recover missing results was contributing to a usage restriction.

One of the proposed remedies was another form of refresh.

The recommendation to stop and regenerate also risked duplicating a job that might already have completed on the server but had failed to appear in the interface.

I replied that the problem had existed since the prior month and that focusing on user-side actions reversed cause and effect.

The display failure forced the refreshes. The refreshes then produced the restriction.

I requested immediate escalation to a human representative and the relevant technical team. I asked for review of the account logs, generation-processing logs, frontend delivery, rate-limit conditions, multi-tab policy, correction timeline, workaround, and possible compensation.

At 2:19 a.m., OpenAI stated that the case had been escalated to a support specialist and that a response should arrive within the next few days.

Nine days passed.

No investigative result arrived.

No progress update arrived.

OpenAI did not state whether Case 5 had been merged into this case.

The phrase “within the next few days” remained in the record. The technical explanation did not.

Case 7: I Reported a Broad GPT-5.5 Thinking Performance Regression in English. OpenAI Closed the Case as a Duplicate Two Minutes Later

At 3:04 p.m. on July 5, 2026, I sent OpenAI an English-language request for a serious review of GPT-5.5 Thinking.

The report identified five recurring failures:

ChatGPT invoked image generation despite explicit instructions not to generate an image.

It imported obsolete scene material, such as carrying a San Francisco hotel-backdoor setting into a new prompt set at the Port of Antwerp in Belgium.

It repeatedly lost the established rule that “promotional portrait” meant a 4:5 vertical, upper-body composition.

It misunderstood a straightforward physical-layout correction. When a diner-booth table appeared to protrude into the aisle, ChatGPT absurdly reduced the real-world table itself instead of correcting the camera framing, crop, and perspective.

It inserted explanations, apologies, and meta-commentary even when instructed to output only the final production text.

I did not present these as isolated creative disagreements.

I asked OpenAI to investigate a possible regression, routing issue, degraded configuration, memory or context-handling failure, tool-use bug, or account-level behavioral problem.

I explicitly stated that I was not requesting a generic explanation that language models sometimes make mistakes.

At 3:06 p.m., OpenAI created Case Number 10976796 and immediately closed it as a duplicate.

Two minutes had elapsed.

OpenAI did not identify the supposed duplicate.

It did not confirm that all five symptoms had been transferred into another case.

It did not say whether the reported GPT-5.5 Thinking regression had been shared with engineering.

When I submitted individual incidents in Japanese, I received peripheral questions, boilerplate responses, or irrelevant instructions.

When I consolidated the incidents into a structured English-language regression report, the case was closed as a duplicate within two minutes.

Whether the problems were divided into separate reports or consolidated into one comprehensive report, they did not appear to reach a substantive technical investigation.

Case 8: OpenAI's Pro Pricing Page Says "Fast and Unlimited Image Creation." OpenAI Support Says Pro Is Not Always Faster Than Plus and Remains Subject to Restrictions

On July 2, I contacted OpenAI to clarify the image-generation benefits of ChatGPT Pro.

After upgrading from Plus to Pro, I had observed little or no noticeable difference in the time required to complete an individual image.

I was not initially reporting an account-specific outage.

I was asking what OpenAI's Pro-plan claims of "faster" and "unlimited" image generation actually meant, what Pro contractually provided beyond Plus, and what OpenAI did or did not guarantee.

OpenAI's initial AI-assisted response stated that there was no public fixed specification guaranteeing that Pro image generation would always be visibly faster than Plus, nor any published breakdown of the alleged speed difference.

It stated that under the same model and similar system load, Pro and Plus could appear to have essentially the same generation speed.

A later response signed by Razel of OpenAI Support provided a more detailed explanation.

According to that response:

"Faster" did not guarantee that every individual image would complete noticeably faster than it would on ChatGPT Plus.

The primary benefits were higher usage limits, priority access during heavy demand, reduced likelihood of congestion-related delays, and improved service availability.

During normal operating conditions, image-generation performance could appear very similar between Plus and Pro.

“Unlimited,” or “virtually unlimited,” did not mean the absence of all restrictions.

It referred to the absence of a fixed daily or monthly image quota for normal use, while remaining subject to fair-use protections, temporary rate limiting, throttling, anti-abuse systems, safety systems, maintenance, and infrastructure-related controls.

OpenAI later reaffirmed that explanation.

As of July 20, 2026, however, OpenAI’s Japanese pricing page displayed the Pro benefit as:

“高速かつ無制限の画像作成”

A natural English translation is:

“Fast and unlimited image creation.”

The Pro column, priced from ¥16,800 per month, included a note stating that unlimited use remained subject to safeguards against abuse.

The same comparison panel did not state that “fast” primarily meant priority during periods of high demand.

It did not state that ordinary image-rendering times could be essentially the same as Plus.

It did not state that faster completion of individual images was not guaranteed.

The language presented to the prospective customer was short and strong:

“Fast and unlimited image creation.”

The limitations presented after a detailed support inquiry were longer and materially narrower:

Not every image will be faster than Plus.

Normal-load performance may appear similar.

The principal advantage is priority during congestion.

Unlimited use remains subject to multiple operational restrictions.

That discrepancy raises a legitimate consumer-protection question.

The issue is not simply that one image felt slow.

The issue is whether the overall presentation of the higher-priced plan could lead an ordinary consumer to believe that individual image generation would be faster than Plus and that image creation would be effectively unrestricted, when OpenAI’s own support explanation defines those claims more narrowly.

In Japan, the relevant questions may include whether the representation could constitute a misleading claim concerning quality or performance, or a misleading claim concerning transaction terms under the Act Against Unjustifiable Premiums and Misleading Representations.

A final legal determination would require examination of the entire pricing page, the placement and clarity of limitations, the basis for the “fast” and “unlimited” claims, the actual operational differences between Plus and Pro, and the overall impression received by an ordinary consumer.

I am not declaring that OpenAI has violated Japanese law.

I notified OpenAI that the discrepancy warranted submission to Japan’s Consumer Affairs Agency as a suspected misleading representation requiring regulatory review.

OpenAI did not withdraw its explanation.

It stated that my concerns had been documented and shared internally with the appropriate teams.

It did not identify any planned correction to the pricing language.

It did not disclose whether legal or advertising review had been initiated.

It did not provide the substantiation behind the “fast” claim.

It did not provide a timetable for any clarification.

ChatGPT

概要機能学ぶCodexビジネス料金ダウンロード

ログイン無料でサインアップ

無料版

日常のタスクに役立つインテリジェンス

¥0 / 月

無料版をはじめよう

✓ GPT-5.5 Instant の利用制限あり

✓ メッセージ数とアップロード数に上限あり

✓ 画像生成の回数に上限あり、速度も制限あり

✓ Deep Research の利用に上限あり

✓ メモリとコンテキストに上限あり

✓ Codex へのアクセスに上限あり

✓ デスクトップでの ChatGPT ワークへの制限付きアクセス

すでにプランをご利用の場合は、請求に関するヘルプをご覧ください。

Go

拡張アクセスで会話を継続

¥1,400 / 月

Goをはじめよう

◆ 無料版のすべてに加え、次の機能を利用できます

✓ GPT-5.5 Instant の上限拡大

✓ より多くのメッセージ

✓ より多くのアップロード

✓ より多くの画像生成

✓ より長いメモリ

このプランには広告が表示される場合があります。詳細はこちら

Plus

高度なインテリジェンスで、さらに多くのことを実現

¥3,000 / 月

Plusをはじめよう

◆ Go のすべてに加え、次の機能を利用できます

✓ GPT-5.6 による高度なリズニングモデル

✓ メッセージ送信とアップロードの上限拡大

✓ より高度で正確な画像生成

✓ Deep Research の拡張

✓ メモリとコンテキストの拡張

✓ プロジェクト、スケジュール読み取り、カスタム GPT

✓ Codex の利用上限が拡大

✓ デスクトップ、ウェブ、モバイルで ChatGPT ワークへのアクセスが拡大

✓ 新機能への先行アクセス

制限が適用されます

Pro

生産性を最大限に高める

月額 ¥16,800 から

Proをはじめよう

◆ Plus のすべてに加え、次の機能を利用できます

✓ 利用量が6倍または20倍に拡大

✓ GPT-5.6 Sol Pro による Pro 推論

✓ Codex の最大タスク数

✓ 高速かつ無制限の画像作成

✓ Deep Research の最大活用

✓ メモリとコンテキストの最大利用

✓ プロジェクト、タスク、カスタム GPT の拡張

✓ 新機能の研究プレビュー

無制限で利用できますが、不正利用防止のための安全対策が適用されます。詳細はこちら

The Same Processing Pattern Runs Through All Eight Cases

Viewed separately, these incidents may appear to concern different parts of the product:

Unauthorized image generation.

Context contamination.

Aspect-ratio violations.

Missing progress indicators.

Failed result delivery.

Rate limiting.

Plan advertising.

Viewed together with OpenAI's support responses, however, a consistent processing pattern emerges.

1. The Substance of the Complaint Is Reclassified Into a More Convenient Category

A tool-control problem is treated as an image-quality issue.

A context-management problem is merged into an image-generation case.

A complaint about current-task identification receives links for exporting data and deleting an account.

A causal problem involving missing results and rate limits is reframed as an extension, VPN, or browser problem.

The user's actual complaint appears to be replaced by whatever preexisting template the support system is prepared to process.

2. OpenAI Requests That the Customer Reconstruct Information OpenAI Should Already Possess

OpenAI repeatedly requested:

Conversation URLs.

Incident times.

Screenshots.

Screen recordings.

Generated-image dimensions.

Prior case numbers.

Destination case numbers after a merge.

Some supplemental information is reasonably necessary in technical support.

The problem arises when the customer is expected to rediscover information that should exist in OpenAI's own conversation history, tool-call records, generated assets, support records, and attachment metadata.

The clearest example is the duplicate-case sequence in which OpenAI performed a merge, could not identify the destination, and instructed the customer to search their own inbox for it.

3. “Duplicate” Functions as an Answer-Termination Mechanism Rather Than a Case-Management Tool

Consolidating duplicate reports can be reasonable.

A legitimate consolidation should identify, at minimum:

The destination case number.

The issues transferred into that case.

The scope of the ongoing investigation.

The responsible team.

The next expected response.

In several of these cases, OpenAI simply stated that the case was a duplicate and closed it.

No destination was identified.

No transferred issue list was provided.

No confirmation was given that the new technical allegation had been preserved.

The duplicate label appeared to terminate the conversation rather than organize an investigation.

4. Escalation Is Announced, but the Responsible Process Disappears

OpenAI repeatedly sent the phrase:

“Your case has been escalated to a support specialist.”

“A response is expected within the next few days.”

After that announcement, nine days, eleven days, or twenty days could pass without a substantive response.

The responsible department remained unknown.

The investigation stage remained unknown.

Whether logs had been reviewed remained unknown.

The next response date remained unknown.

The word “escalation” appeared in the email, but no accountable escalation process became visible.

5. OpenAI’s Deadlines and the Customer’s Deadlines Are Radically Asymmetrical

OpenAI announced a response “within a few days,” then remained silent for twenty days.

The customer was warned that the case would be closed unless they replied within one day.

That warning was issued even after the customer had replied the previous day and explicitly requested that the unresolved case remain open.

The provider's silence was treated as open-ended.

The customer's silence was given a twenty-four-hour expiration period.

The structure shifted the burden of maintaining an unresolved investigation away from the provider and onto the person who reported the defect.

6. AI Support Reproduced the Failures Reported in ChatGPT

I reported that ChatGPT had lost the current context.

Support lost the current complaint.

I explained that an attachment was a reference image.

Support treated it as the failed output.

I submitted a new issue.

Support returned to an older, unrelated case.

I provided the same information repeatedly.

Support repeatedly asked for it again.

OpenAI placed an AI system at the entrance to the process for investigating AI failures. That AI system could misunderstand the complaint, lose the case context, issue irrelevant instructions, classify reports as duplicates, and close them.

Under those conditions, product defects and support defects do not merely coexist. They amplify one another.

“AI Systems Can Make Mistakes” Does Not Resolve These Cases

I do not expect a generative model to be infallible.

I am not calling every creative disagreement, stylistic difference, or imperfect visual interpretation a product defect.

The issues documented here are more basic and more objectively testable.

The user wrote, “This is not an image-generation request,” and the image tool ran.

The user specified an exact 4:5 ratio, and another ratio was produced.

The user specified the current location, and a previous location appeared.

The user repeated the same straightforward question twice, and the system remained trapped in the wrong task.

The progress indicator and completed result disappeared.

Refreshing to determine the job status triggered a usage restriction.

OpenAI said it would respond within a few days, then remained silent for nine or twenty days.

These are not disagreements over creativity.

They concern whether a tool was invoked.

The numerical dimensions of an output.

The text of the current instruction.

The visible state of a processing job.

The date of a promised response.

Each can be examined against records.

Questions OpenAI Still Needs to Answer Publicly

Following publication of this report, I continue to request answers to the following questions.

In Case Number 10337718, did OpenAI confirm from its logs that the image-generation tool ran in response to messages that did not request image generation?

Did OpenAI confirm that the behavior recurred after an explicit prohibition?

Did those invocations consume image-generation capacity, time windows, rate-limit allowances, or other paid usage opportunities?

Did OpenAI consider restoring usage, issuing credits, extending the subscription, or adjusting the fee?

In Case Number 11271585, did OpenAI correct the record to reflect that the attached image was a reference image rather than the failed output?

Did OpenAI compare the requested 4:5 ratio with the actual dimensions stored in its generated-asset history?

In Case Number 11268075, did OpenAI investigate why completed image results failed to appear in the browser?

Did page refreshes contribute to the rate-limit restriction?

What is OpenAI's formal policy on operating two image-generation pages simultaneously?

Where were duplicate Cases 11269846, 10976796, 11319313, 11319385, and 11319858 merged?

Were the distinct issues in those cases preserved in the destination cases?

When OpenAI states that a matter has been escalated to a support specialist, what department receives it, what information is transferred, and what event constitutes actual acceptance of responsibility?

Why did OpenAI repeatedly state that responses were expected within a few days and then provide no substantive response for nine, eleven, or twenty days?

If the Japanese Pro pricing page says “Fast and unlimited image creation,” but “fast” primarily means congestion priority rather than consistently shorter individual rendering times, why is that limitation not stated directly in the pricing comparison?

What rate limits, fair-use standards, throttling rules, and other conditions apply to “unlimited” image generation, and how clearly are they disclosed before a customer selects the higher-priced plan?

The Record Remains After the Explanation Disappears

OpenAI can update an interface.

It can change the language on a pricing page.

It can close a support case.

It can summarize a complaint in a way that removes a disputed issue.

What it cannot do is retroactively make the underlying record disappear.

The emails remain.

The case numbers remain.

The dates remain.

The screenshots remain.

The names attached to support responses remain.

The answers OpenAI issued remain.

So do the answers it did not issue.

This report does not disclose OpenAI's confidential internal reasoning or proprietary system information.

It documents the service a paying customer actually received.

It documents the complaints that customer actually submitted.

It documents the replies OpenAI actually sent.

It documents the replies that never came.

If OpenAI sells paid plans under the language of safety, trust, advanced intelligence, and maximum productivity, the reliability of the service cannot be measured solely by product demonstrations.

It must also be measured by what happens when the product fails.

Can OpenAI establish what occurred?

Can it retain the substance of the customer's complaint?

Can it maintain continuity across a case?

Can it restore paid usage consumed by an erroneous tool invocation?

Can it explain material plan limitations before purchase rather than after escalation?

Can an accountable human representative provide a specific response within a defined period?

Those functions are not separate from the product.

They are part of the paid AI service itself.