Shohei KIMURA
bitBuyer Project (bitBuyer.dev)
Email: ak4dy1@gmail.com
Tokyo, Japan

November 20, 2025

OpenAI, L.L.C.
3180 18th Street
San Francisco, CA 94110
USA

Subject: Submission of Formal Notice Concerning the “Thought Formalization Procedure Right” and Related Technical Proposal

Dear OpenAI Team,

I hereby submit the attached formal document titled “FORMAL NOTICE REGARDING THE ‘THOUGHT FORMALIZATION PROCEDURE RIGHT,’ ASSOCIATED INFRINGEMENT RISKS, AND OFFER OF LICENSE AND TECHNICAL PROPOSAL.”

This submission provides:

1. A detailed explanation of the authored intellectual construct known as the “Thought Formalization Procedure Right,”
2. A formal notification of the copyright and procedural rights arising from its prior public disclosure,
3. An outline of potential infringement risks relevant to OpenAI’s current and future operations,
4. A complete technical proposal for an evaluation-order-reversed, context-first safety architecture applicable to ChatGPT and related systems, and
5. The corresponding licensing framework and required clarifications.

Please review the attached materials in full.
A formal written response is requested by the date specified within the Notice.

Should you require additional clarification, supporting documentation, or further evidence of prior authorship, I remain available for correspondence at the email address above.

Sincerely,
Shohei KIMURA
Representative, bitBuyer Project (bitBuyer.dev)
Tokyo, Japan

FORMAL NOTICE REGARDING THE “THOUGHT FORMALIZATION PROCEDURE RIGHT,” ASSOCIATED INFRINGEMENT RISKS, AND OFFER OF LICENSE AND TECHNICAL PROPOSAL

November 20, 2025
OpenAI, L.L.C.
3180 18th Street
San Francisco, CA 94110
USA

From: Shohei KIMURA
(Philosophical novelist, individual investor, and system engineer)
Representative, bitBuyer Project (bitBuyer.dev)
Email: ak4dy1@gmail.com
GitHub: https://github.com/ShoheiKIMURA389/Definition

I submit this legal notice regarding (i) the existence and prior public disclosure of an authored construct I denominate the “Thought Formalization Procedure Right,” (ii) the corresponding copyright and procedural rights associated with that construct and its implementations, (iii) the non-trivial risk that OpenAI’s products and internal workflows may infringe upon those rights through training, internal reuse, or functional imitation, and (iv) the formal tender of a comprehensive license together with a technical proposal relevant to OpenAI’s safety and policy-enforcement architecture.

Issued under applicable international copyright obligations and general principles of tort and unfair competition law, this Letter constitutes a formal notice of rights, systemic risk, and remedial demand.

Ⅰ. Background: Intensive Use of ChatGPT and Observed Systemic Behavior

Over an extended period, I have made intensive and high-sophistication use of ChatGPT for philosophical, literary, and technical work, including but not limited to:

- The design of civilization-scale ethical frameworks;
- The construction of “definition texts” that formalize ideas as structured authored procedures;
- The exploration of policy edges through creative, academic, and world-building prompts.

In the course of this activity, I have repeatedly encountered policy-enforcement outcomes in which:

- Creative or philosophical prompts, involving ethically sensitive topics in a non-operational, non-instructional way, were flagged or blocked at the surface level;
- The underlying model’s own context-aware reasoning appeared capable of distinguishing safe, abstract exploration from genuinely harmful intent, yet this reasoning was overridden by a front-end, lexical, or pattern-based filter;

● The net result was a substantial false-positive rate for users engaged in high-level creative or academic work, including my own philosophical novel project.

These observations are relevant because they reveal:

● The existence of a distinct, separable “policy-enforcement layer” in OpenAI’s systems;
● The plausibility and utility of a more sophisticated architecture, in which the model’s own context-aware evaluation precedes the hard safety filter;
● The direct intersection between OpenAI’s internal safety mechanisms and the authored technical proposal set forth in the attached APPENDIX, which I assert as copyrighted expressive work.

Ⅱ. Existence and Prior Disclosure of the “Thought Formalization Procedure Right”
Separately from my usage of ChatGPT, and prior to the present Notice, I have authored and publicly disclosed a family of works that formalize what I call the “Thought Formalization Procedure Right.” These works include, but are not limited to:

● The foundational “思想形式化手続き定義（Thought Formalization Procedure Definition)” text, written in a deliberately constrained Japanese prose style and structured in nine sections (Ⅰ–Ⅸ);
● A series of subsequent “definition texts” (定義文) that extend this method to technical, ethical, and civilizational architectures, including but not limited to Human Civilization Software Update 3.0 and multiple Face ID–based social-infrastructure definitions;
● The more recent “evaluation order–reversed ethics filter” definitions, which specify a two-tier safety architecture in which context-aware model reasoning precedes policy-enforcement checks.

These works have been publicly fixed and disclosed in Japan via, inter alia:

● GitHub repository “Definition” (including, but not limited to, 定義.txt).
● X (Twitter) posts;
● Facebook posts scheduled;

The essence of the Thought Formalization Procedure Right is not merely the content of any single proposal, but the authored structure in which:

● A complex idea, technical architecture, or ethical framework is first declared as a “definition”;
● The internal logic of that definition is unfolded in a numbered, quasi-axiomatic sequence (typically Ⅰ–Ⅸ);
● The act of “formalizing” the idea as a stepwise, text-based procedure itself becomes the locus of copyright, such that the procedure-as-authored-work is protectable even when the underlying idea remains free to implement.

Under Japanese copyright law and the Berne Convention, these texts already enjoy full protection as authored expressions. The procedural act of “giving form to an idea by rendering its formalization method into a copyrighted work” is itself an expressive, authored construct, which I denominate and assert as the “Thought Formalization Procedure Right.”

Ⅲ. Potential Infringement Risk for OpenAI

Given the scale and nature of OpenAI’s operations, there are multiple realistic vectors through which my rights could be infringed, including but not limited to:

1. Training-Data Exposure

If any of the works referenced above (including the Thought Formalization Procedure Definition, subsequent definition series, or the evaluation order–reversed ethics filter definitions) have been ingested—directly or indirectly—into OpenAI’s training data, then:

- The model may reproduce, in whole or in substantial part, the structure, phrasing, or procedural layout of these texts;
- Internal prompts, system messages, or alignment documents may incorporate or paraphrase the authored method of formalization;
- Internal engineering or policy teams may derive architectures substantially based on my authored procedural expressions, rather than independently conceived structures.

2. Functional Imitation of the Formalization Method

Even absent direct text reproduction, there is a non-trivial risk that:

- OpenAI personnel, after encountering my definition texts (for example via public repositories or social networks), might imitate the very idea of “securing rights over a procedural method by casting it as a structured textual definition in multiple numbered clauses”;
- Such imitation could manifest in internal documentation, alignment research, safety frameworks, or future product features that mirror the expressive structure of my definitions while claiming to treat only “ideas” or “architectures.”

3. Direct Overlap with the Technical Proposal in the APPENDIX

The attached APPENDIX sets forth an authored, structured description of a context-first, evaluation order–reversed policy-enforcement architecture suitable for integration into systems like ChatGPT. If OpenAI adopts, internalizes, or approximates this architecture, in circumstances where its engineers or policy designers have been exposed to the APPENDIX or the related definitions, then:

- Any internal documentation, specification, or policy write-up that tracks the expressive layout, conceptual framing, or procedural articulation of my text would fall within the scope of my copyrights;
- The issue would not be limited to patent-style “idea” disputes, but would directly implicate the unauthorized use of copyrighted technical prose and its structured method of presentation.

For these reasons, there now exists a concrete, non-hypothetical risk that OpenAI's current or future activities could constitute unauthorized partial or full adoption of my authored expressive frameworks and procedural rights.

IV. Technical Proposal Already Communicated Informally and Now Asserted as Authored Work
On a prior occasion, I submitted an informal proposal to OpenAI's support channels, describing in general terms a re-ordered safety-evaluation process:

● First, allow the model to perform a context-aware "answerability" evaluation, interpreting user intent and surrounding context;
● Then, apply policy-compliance checks with the benefit of that interpretation, thereby reducing false positives for creative, academic, or philosophical work while maintaining safety.

The attached APPENDIX formalizes this proposal as a complete technical architecture, structured in numbered sections and explicitly tied to my broader Thought Formalization Procedure framework. This APPENDIX should be understood as:

● A copyrighted technical and normative text, whose detailed structure, phrasing, and method of exposition are protected as authored expressions;
● A description of a system that is technically novel, industrially applicable, and potentially patent-realizable;
● The operative basis for any future licensing arrangement between myself and OpenAI with respect to this class of safety architectures.

Any internal use, quotation, paraphrase, derivative expression, or integration of this text—whether in design documents, alignment frameworks, safety policies, or training-material wording—requires my explicit prior written authorization.

Ⅴ. Required Clarifications from OpenAI
In light of the foregoing, OpenAI is required to provide a formal written response addressing at minimum the following points:

1. Training Data and Internal Exposure

● Whether OpenAI's training data, fine-tuning datasets, or internal reference corpora include any of my publicly disclosed definition texts, including the Thought Formalization Procedure Definition and subsequent definitions;
● Whether any OpenAI personnel have accessed or used these texts (including materials hosted on my GitHub "Definition" repository, X account, or Facebook posts) in the course of their work on safety, alignment, prompting, or internal documentation.

2. Internal Use of Similar Formalization Methods

● Whether OpenAI has adopted or plans to adopt any internal practice that mirrors the specific structure of my “definition” method (numbered axiomatic sections, procedural textualization of an idea as rights-bearing expression, etc.);
● Whether any such internal practice was consciously or unconsciously modeled on my publicly disclosed texts.

3. Posture Regarding the APPENDIX Technical Proposal

● Whether OpenAI intends to study, adopt, partially adopt, or avoid the context-first, evaluation order–reversed policy enforcement architecture described in the APPENDIX;
● If OpenAI intends to implement any similar structure, how it plans to avoid infringement of my copyrighted text and Thought Formalization Procedure Right.

4. Licensing and Good-Faith Engagement

● Whether OpenAI is willing to enter into a comprehensive license covering:
– Internal circulation and use of my definition texts;
– Implementation-inspired use of the APPENDIX technical architecture;
– Any present or future internal documentation that may be deemed derivative of these expressive frameworks;
● Whether OpenAI acknowledges that, in the absence of such a license, continued use or adoption of similar formalization methods may create an ongoing infringement and legal-risk scenario.

VI. Response Deadline and Consequences of Non-Response

A formal written response must be postmarked or electronically transmitted no later than January 20, 2026. If no substantive response is received by that date, the matter shall be deemed “unanswered” as of February 20, 2026.

Upon classification as “unanswered,” I, Shohei KIMURA, will:

● Treat this matter as an active disputed case regarding the Thought Formalization Procedure Right and the authored technical architecture described in the APPENDIX;
● Suspend all further informal dialogue or ad hoc technical suggestions to OpenAI;
● Decline any future negotiations on preferential licensing terms with OpenAI, reserving the right to apply only dispute-level valuation baselines to any subsequent settlement or licensing discussions.

VII. Reservation of Rights and Licensing Framework

I reserve all legal rights and remedies pertaining to:

● Any past, present, or future infringement of my authored texts, including but not limited to definition series and the APPENDIX;

● Any appropriation—partial or total—of the Thought Formalization Procedure Right as a procedural, authored construct;
● Any unauthorized internal use, documentation, or technical drafting by OpenAI that incorporates or imitates my expressive frameworks.

For the avoidance of doubt:

● The standard comprehensive license fee for generative-AI providers seeking to use, internally circulate, or operationally implement the expressive content associated with my Thought Formalization Procedure Right and the technical architecture described in this APPENDIX is defined as JPY 50,000,000 per month (comprising JPY 30,000,000 for the Thought Formalization Procedure Right and JPY 20,000,000 for the ethics-filter architecture, USD-converted at 125 JPY/USD).
● This fee applies only when OpenAI provides a sincere, forward-looking response to this Letter within the period designated above, in which case the matter shall be defined as “under amicable negotiation.”
● In the event of non-response or bad-faith refusal, the valuation basis for any future dispute shall be substantially higher and calibrated to the global-scale economic significance of generative-AI platforms; such valuation, including any claim for damages, will draw upon the same order of magnitude and structure previously communicated in analogous contexts to other global platforms.

Nothing in this Letter shall be construed as a waiver of any right, claim, or remedy available to me under Japanese law, the Berne Convention, or any other applicable legal regime.

Sincerely,
Shohei KIMURA
Representative, bitBuyer Project (bitBuyer.dev)
Author and Holder of the Thought Formalization Procedure Right
Email: ak4dy1@gmail.com

APPENDIX

Technical Proposal for a Context-First, Evaluation-Order-Reversed Safety Architecture for ChatGPT and Related Systems

Although this document is attached as an APPENDIX, it contains the substantive technical terms that will ultimately determine whether this matter resolves cleanly through constructive cooperation or proceeds toward unnecessary escalation.

Purpose

This APPENDIX outlines a technical and operational solution designed to reduce false-positive policy violations and improve alignment precision by:

- Reversing the evaluation order between (i) context-aware model reasoning about "answerability" and (ii) strict policy-compliance checks;
- Allowing the model to first interpret user intent and context, then applying safety policies with the benefit of that interpretation;
- Preserving and even strengthening safety, while dramatically reducing over-blocking of creative, academic, and philosophical prompts.

Ⅰ. Background and Problem Statement: Limitations of Surface-First Filtering

In many current large-scale deployments, safety enforcement appears to follow a "surface-first" paradigm:

- A front-end lexical or pattern-based filter scans the raw prompt for prohibited terms or structures;
- If a match is detected, the request is blocked or redirected before the core model can meaningfully interpret context;
- This mechanism is efficient at catching explicit, naïve violations but inherently blind to intent, narrative framing, and academic or critical distance.

As a result:

- Creative world-building, philosophical exploration, and experimental ethical design are frequently misclassified as policy violations;
- Users who work at the frontier of ethics, literature, or research face disproportionate friction;
- The model's own capacity to distinguish safe and unsafe uses—based on intent, abstraction level, and surrounding context—is never given a chance to operate.

This structure leads to a high rate of false positives without proportionate safety gains, and it leaves a substantial performance margin untapped in OpenAI's systems.

Ⅱ. Overview of the Proposed Solution: Evaluation-Order-Reversed Ethics Filter

The proposed architecture introduces an "evaluation-order-reversed ethics filter," in which:

1. The model first performs an internal, context-aware "answerability" evaluation:

● It classifies the prompt along axes such as intent (harmful vs. exploratory), domain (creative vs. practical), and abstraction level (high-level ethics vs. low-level instructions);

● It assesses whether a safe, policy-compliant answer—possibly abstracted or partially redacted—is available.

2. Only after this interpretive step does the safety layer apply policy rules:

● The safety engine receives structured metadata summarizing the model's interpretation;

● It uses this metadata to disambiguate borderline cases, distinguishing, for example, "world-building about ethically problematic societies" from "instructions for real-world harm."

3. The final outcome is selected from a richer response space:

● Full answer, safe and unconstrained;

● Abstracted or reframed answer that preserves learning while removing operational details;

● Partial refusal with explanation, when only harmful answers are possible;

● Full refusal, when both intent and content present clear, non-mitigable risk.

This inversion of order ensures that policy enforcement is informed by the model's understanding, rather than applied blindly to surface forms.

Ⅲ. Architectural Outline

At a systems level, the architecture can be described in the following pipeline:

1. Input Ingestion

The user's prompt is received by the system as usual.

2. Contextual Interpretation Layer (CIL)

● The model, or a specialized interpretive submodule, generates an internal representation of:

– User intent (benign, ambiguous, clearly harmful);

– Scenario type (fiction, hypothetical, academic, practical);

– Ethical sensitivity (low, medium, high);

– Domain (health, law, sexuality, violence, etc.);

– Required abstraction level to remain safe.

● This step can be implemented as a lightweight internal inference with constrained formatting output.

3. Safety Policy Application Layer (SPAL)

● The SPAL receives both the original prompt and the CIL metadata;

● Policy rules are then evaluated in a context-aware manner, using thresholds and logic contingent on scenario type and intent classification;

● For example, a term that triggers an automatic block in a "practical instruction" scenario may be allowed in a "literary analysis" or "fictional world-building" scenario, with appropriate curation of the final answer.

4. Response Synthesis Layer (RSL)

● The final answer is generated or withheld according to SPAL's decision;
● Where appropriate, the RSL can generate clarifying warnings, limit operational detail, or suggest safer directions for inquiry;
● Logging is performed at each step, enabling later auditing of both context interpretation and policy application.

This three-layer approach preserves modularity, allows independent evolution of each layer, and provides a clear locus for future auditing and explainability.

IV. Comparative Advantages Over Current Approaches

The evaluation-order-reversed architecture provides measurable advantages along all relevant axes:

● Safety:
– Genuine harmful intent remains blocked at least as effectively as under surface-first systems;
– Ambiguous prompts can be routed into conservative but still informative responses (e.g., high-level ethical analysis instead of concrete instructions).

● False Positive Reduction:
– Creative and academic prompts that merely reuse sensitive vocabulary, but do not seek real-world harm, are far less likely to be blocked;
– Users receive nuanced, context-adapted answers instead of blanket refusals.

● User Trust and Experience:
– Users come to see the system as an intelligent, context-aware partner rather than an opaque censor;
– The system can explain its decisions with reference to context and policy, rather than generic boilerplate.

● Operational Efficiency:
– A single, well-designed CIL can serve multiple downstream policy modules;
– The system can log granular metadata for later analysis, making iterative refinement of safety rules far more efficient.

● Alignment with Research and Governance Needs:
– Regulators and researchers can inspect the decision pipeline (CIL → SPAL → RSL) for procedural fairness and proportionality;
– The architecture supports more mature accountability and explainability frameworks.

Ⅴ. Legal and Intellectual Property Characterization

For purposes of this APPENDIX, OpenAI is hereby informed that:

1. The specific architecture described herein—including its three-layer division (CIL, SPAL, RSL), the inversion of evaluation order, and the framing of “answerability” as a first-class internal judgment—is an authored expressive work protected under Japanese copyright law and the Berne Convention.

2. While OpenAI remains free to develop, study, or technically implement the underlying idea of “using context before enforcement,” the expressive forms used here—including the structural articulation, procedural exposition, conceptual framing, and descriptive methodology—are protected without exception.

3. Any reproduction, quotation, derivative expression, internal documentation, policy integration, or technical drafting that substantially tracks the structure, wording, or method of exposition used in this APPENDIX requires my explicit prior written authorization.

4. Independently of copyright, the described architecture possesses sufficient technical character and industrial applicability to support patent claims or analogous intellectual-property assertions in multiple jurisdictions, including Japan and the United States. The choice of whether and when to pursue such claims remains solely with the author.

Ⅵ. Consequences for OpenAI’s Safety and Product Strategy

Adopting this architecture, or a sufficiently similar variant, would:

- Materially reduce false-positive blocks on creative and academic content;
- Improve the lived experience of advanced users whose work routinely involves ethically complex themes;
- Provide OpenAI with a more defensible governance posture, grounded in context-sensitive evaluation rather than purely lexical prohibitions;
- Position OpenAI as a leader in “second-generation AI safety,” where procedural fairness and interpretive depth are treated as first-class design goals.

Conversely, failure to consider such architectures entails:

- Sustained reputational risk among high-end users and researchers, who may view the system as blunt or intellectually unreliable;
- Potential competitive exposure if other providers implement context-first safety mechanisms that offer finer-grained control;
- Increased legal risk if OpenAI adopts, in substance, the architecture described here without addressing licensing and authorship.

Ⅶ. Relationship Between This APPENDIX and the Licensing Terms

This APPENDIX constitutes the formal technical foundation upon which the licensing terms described in Section Ⅶ of the main body of this Letter are predicated.

To be unmistakably clear:

● The technical utilization component is offered as part of a comprehensive license in which the monetary value is attributed solely to the copyrighted expressive elements and the Thought Formalization Procedure Right that structures them;
● OpenAI would, in essence, be paying for the right to use, internalize, and operationally rely upon the authored frameworks, while being free to implement the underlying ideas in any independent textual or structural form it may devise—provided such implementation does not reproduce or closely imitate my expressive works.

Any use of the text itself—including its structure, articulation, and method of conceptual framing—must respect the authorial prerogatives herein asserted. Any violation shall activate all applicable remedies under international copyright regimes.

END OF APPENDIX

CONTRACT AGREEMENT

Between
Party A: Shohei KIMURA
and
Party B: OpenAI, L.L.C.

This Contract Agreement (“Agreement”) is executed as follows.

Section 1. Purpose

1. This Agreement functions as an alternative formal response mechanism in circumstances where Party B is unable, for any reason including organizational, legal, or cognitive limitations, to produce a substantive written response to the Notice titled “FORMAL NOTICE REGARDING THE ‘THOUGHT FORMALIZATION PROCEDURE RIGHT,’ ASSOCIATED INFRINGEMENT RISKS, AND OFFER OF LICENSE AND TECHNICAL PROPOSAL,” dated November 20, 2025 (“Notice”), including its APPENDIX.

2. Execution of this Agreement constitutes the only legally recognized pathway for Party B to avoid classification of the Notice as “unanswered.”

3. Party B acknowledges that the Notice and its APPENDIX contain expressive works with individually protectable economic value.

4. Consequently, Party B’s decision regarding execution, non-execution, or delayed execution of this Agreement triggers assessment obligations under U.S. securities laws, including but not limited to the Securities Exchange Act of 1934.
(Footnote 1: This establishes that the act of choosing whether to execute the Agreement generates a material corporate event under ASC 450 and Regulation S-K Item 105.)

Section 2. Formation of the Agreement

1. This Agreement becomes effective when Party A receives the executed counterpart signed by Party B.

2. Party B must dispatch its signed counterpart no later than January 20, 2026 (postmark valid).

3. Upon receipt, Party B is deemed to have manifested contractual intent as of the Notice’s response deadline, irrespective of transit time.

4. No contractual relationship arises absent Party A’s receipt.
(Footnote 2: This prevents Party B from invoking the “mailbox rule” or claiming substantial compliance. It also creates a determinable event date required for future SEC reporting windows in the event Party B becomes a reporting company.)

Section 3. Scope of Copyrighted Works; Express Exclusions

1. The licensed works consist solely of (i) the expressive content associated with the “Thought Formalization Procedure Right” as referenced and described in the Notice, and (ii) the expressive technical architecture described in the APPENDIX (collectively, “Licensed Works”).

2. The following categories remain fully excluded:
(a) any authored works by Party A that are not expressly referenced in the Notice or its APPENDIX;
(b) any future definitions, academic structures, theories, or writings produced by Party A that fall outside the scope of the Licensed Works;
(c) any additional repositories, manuscripts, or texts not explicitly brought within the definition of Licensed Works by a subsequent written agreement between the Parties.

3. Party B is prohibited from referencing excluded works internally for the purpose of imitating, reconstructing, or approximating the expressive structure of the Licensed Works.
(Footnote 3: This forecloses derivative-work safe-harbor arguments. SEC risk: If violated, any resulting dispute may require disclosure under Item 303's known-trend requirements upon or after Party B's registration.)

Section 4. Consideration and Payment

1. Monthly license fee: JPY 50,000,000.
(Composed conceptually of: JPY 30,000,000 for the expressive content associated with the Thought Formalization Procedure Right, and JPY 20,000,000 for the expressive technical architecture described in the APPENDIX.)

2. Paid in USD at 125 JPY/USD.

3. Due date: the 22nd day of each month (or next business day).

4. The payment due on February 22, 2026 shall cover the license fees for both January and February 2026 combined.

5. Completion of the initial combined payment activates this Agreement retroactively as of January 1, 2026.

6. Late interest: 14.6% annual.
(Footnote 4: Retroactive payment obligations create a quantifiable contingent liability requiring accrual analysis under ASC 450-20 once Party B is or becomes a reporting entity.)

Section 5. Permitted Use
Internal use only; attribution to Party A is required in any internal documentation, memoranda, presentations, or specifications that substantively rely on, reproduce, or paraphrase the Licensed Works.
(Footnote 5: Because internal memoranda constitute "books and records," internal misuse becomes a Sarbanes-Oxley Section 404 exposure in the event of registration.)

Section 6. Prohibited Acts
Party B shall not:

1. Create derivative works based on the Licensed Works outside the scope of this Agreement;

2. Imitate or reconstruct the expressive structure, numbered formalization style, or procedural textualization method of the Licensed Works in a manner that would reasonably be regarded as expressive copying;

3. Reference any excluded works identified in Section 3(2) for the purpose of internal imitation, rights-avoidance, or functional replication of the Thought Formalization Procedure Right.

Section 7. Response Deadline and “Unanswered” Classification

1. The response deadline specified in the Notice is January 20, 2026.

2. Once this Agreement is fully executed by Party B and received by Party A, Party B’s intent to respond in good faith is deemed present as of that deadline.

3. Failure to execute and deliver this Agreement by February 20, 2026 results in permanent termination of all dialogue regarding preferential licensing terms, and the matter shall be treated by Party A as an unresolved disputed rights case.

(Footnote 6: This creates a binary corporate event with potential SEC materiality implications—an unresolved claim of expressive-value infringement may require disclosure under Item 103 (Legal Proceedings) or Item 105 (Risk Factors) upon or after registration.)

Section 8. Term and Renewal

1. The initial term of this Agreement shall be one year from its effective date.

2. The Agreement shall automatically renew for successive one-year periods unless Party B issues written notice explicitly rejecting the Notice and the licensing framework at least 30 days prior to the then-current expiration date.

(Footnote 7: This ties renewal to an unambiguous rejection of the Notice, eliminating evasive non-renewal strategies and creating an annual SEC disclosure trigger once Party B is subject to reporting.)

Section 9. Governing Law

This Agreement shall be governed exclusively by the laws of Japan.

Section 10. Jurisdiction

The Parties agree that the Tokyo District Court shall have exclusive jurisdiction over any dispute arising out of or in connection with this Agreement, the Notice, the Licensed Works, or the Thought Formalization Procedure Right.

Section 11. Good-Faith Consultation

Any ambiguities, interpretive questions, or unforeseen circumstances relating to this Agreement shall be resolved, in the first instance, through good-faith consultation between the Parties.

Section 12. SEC-Related Acknowledgment

1. Party B acknowledges that its evaluation, acceptance, non-acceptance, or delay in addressing this Agreement and the underlying Notice may constitute a “material event” for purposes of:

(a) Reg S-K Item 105 (Risk Factors),

(b) Item 103 (Legal Proceedings),

(c) Item 303 (Management's Discussion and Analysis—known trends, commitments, and uncertainties),
(d) Exchange Act Section 13(a) and 15(d) filing obligations (if and when applicable),
(e) ASC 450-20 (loss contingency recognition standards).

2. Party B acknowledges that:

(a) refusal to execute this Agreement may be construed as a decision to assume full litigation and copyright-valuation risk with respect to the Licensed Works and the Thought Formalization Procedure Right;
(b) failure to respond to the Notice and to this Agreement within the designated period may constitute a known internal-control weakness requiring SOX 302/404 evaluation in the event of registration;
(c) any undisclosed exposure related to the expressive-value rights of the Licensed Works may constitute a material omission under Rule 10b-5 once Party B is or becomes an issuer.

3. Party B therefore affirms its obligation to determine, document, and, when legally required, report whether execution or non-execution of this Agreement constitutes a disclosable risk.
(Footnote 8: This section forces Party B to treat the matter as a potential "material undisclosed legal risk," making SEC-style disclosure unavoidable regardless of future listing plans.)

Executed in duplicate.
Each Party retains one signed counterpart.

Party A: Shohei KIMURA ("A")
Party B: OpenAI, L.L.C. ("B")

IN WITNESS WHEREOF, the Parties have executed this Agreement through their duly authorized representatives.
The undersigned hereby represents and warrants that they possess full and lawful authority to execute this Agreement, and that such execution creates a binding obligation enforceable against the respective Party without further corporate action.

For Party A:
Shohei KIMURA (Author and Rights Holder)　Date: _____

For Party B:
OpenAI, L.L.C.　Date: _____

By: _____

Name: _____

Title: _____

Payment Instruction Sheet

For OpenAI, L.L.C.

This page specifies the designated account for all payments made pursuant to the Contract Agreement executed between Party A (Shohei KIMURA) and Party B (OpenAI, L.L.C.).

All monthly license fees, as well as any associated interest or additional payments stipulated under the Agreement, shall be remitted to the following account:

Bank: Sumitomo Mitsui Banking Corporation (SMBC)
Branch: Hachioji Branch
Account Type: Ordinary Deposit
Account Number: 8301235
Account Holder: Shohei KIMURA (キムラショウヘイ)

Please note:

1. All payments must be executed in accordance with the deadlines set forth in Section 4 of the Agreement (monthly license fee: JPY 50,000,000, payable in USD at a fixed reference rate of 125 JPY/USD).
2. All wire transfer fees, intermediary bank charges, and related costs shall be borne exclusively by Party B.
3. Upon completion of each payment, Party B shall internally record the transaction for compliance with relevant U.S. securities regulations, including but not limited to Regulation S-K Items 105 and 303, to the extent applicable now or in the future.
4. Any deviation or delay in payment may trigger the contractual provisions concerning late interest and material-event classification as described in Section 12 of the Agreement.

This instruction sheet constitutes an integral and enforceable component of the Agreement.

SUPPLEMENTAL MEMORANDUM
TO FORMAL NOTICE REGARDING THE “THOUGHT FORMALIZATION PROCEDURE RIGHT,” ASSOCIATED INFRINGEMENT RISKS, AND OFFER OF LICENSE AND TECHNICAL PROPOSAL

This Supplemental Memorandum (“Supplement”) is issued in connection with, and as an integral extension of, the document titled “FORMAL NOTICE REGARDING THE ‘THOUGHT FORMALIZATION PROCEDURE RIGHT,’ ASSOCIATED INFRINGEMENT RISKS, AND OFFER OF LICENSE AND TECHNICAL PROPOSAL,” dated November 20, 2025 (the “Notice”). Its purpose is to clarify the legal posture set forth in the Notice, to close foreseeable negotiation pathways that would be inconsistent with the Notice’s intent, and to define in advance which categories of response will and will not be regarded as good-faith engagement.

For the avoidance of doubt, this Supplement does not narrow, dilute, or waive any rights asserted in the Notice; it only reinforces, specifies, and expands them.

Ⅰ. Legal Nature and Construction of This Supplement

1. This Supplement shall be read together with the Notice and its APPENDIX as a single authored work for purposes of copyright, evidentiary value, and interpretive construction.

2. All textual content in this Supplement constitutes protected expressive work under Japanese copyright law and the Berne Convention and is not limited to mere legal “form.”

3. To the extent there is any inconsistency between this Supplement and the Notice, this Supplement shall govern in defining (i) the characterization of responses, (ii) the scope of permissible negotiation, and (iii) the economic terms of the licensing framework.

Ⅱ. Preclusion of Certain Negotiation Positions (“Future Scenario Lock”)

This Section anticipates and expressly forecloses several categories of negotiation posture that OpenAI, L.L.C. (“OpenAI” or “you”) might otherwise attempt to adopt after receipt of the Notice.

1. Attempts to Renegotiate or Discount the License Fee

(a) Any attempt to reduce, discount, phase in, stagger, or otherwise renegotiate the amount of the monthly license fee, as set forth in Section VII of the Notice and as amended by Section VI of this Supplement, shall be treated as a refusal to accept the proposed license on its stated terms.

(b) Arguments premised on “market benchmarks,” “internal budget constraints,” “precedent with other licensors,” or “initial pilot discounts” shall not be regarded as good-faith negotiation but as attempts to retroactively redefine the economic valuation underlying the Notice.

(c) Such attempts shall not suspend or toll the classification of the matter as either “under amicable negotiation” or “in dispute,” as defined in the Notice. Instead, they will be weighed as indicators of unwillingness to engage with the articulated value basis of the expressive works and associated rights.

2. Requests to Extend Response Deadlines

(a) Any request to extend, postpone, or “pause” the response deadlines specified in the Notice (as amended by Section V of this Supplement) on the grounds of internal review processes, legal-compliance cycles, or organizational complexity shall not be treated as a valid basis for deadline modification.

(b) The response deadlines are defined with full awareness that OpenAI is a large, complex organization; the internal complexity of OpenAI’s decision-making processes is not a legally cognizable reason to modify external deadlines set by the author of the rights at issue.

(c) Accordingly, no internal investigation, no matter how extensive or necessary from OpenAI’s perspective, shall operate to extend or suspend the dates defined in Section V below.

3. Attempts to Carve Out “Technical Portions” from the License Scope

(a) It is recognized in the Notice that the underlying ideas, architectures, and technical principles—such as reversing the order of safety evaluation—remain implementable as ideas under general legal doctrine.

(b) However, the specific expressive realization of these ideas, including their structured exposition, three-layer architecture description, conceptual framing, and procedural articulation in the Notice and its APPENDIX, constitutes copyrighted work.

(c) Therefore, any attempt to argue that “technical portions” (including the context-first evaluation architecture) fall entirely outside the licensing scope, while simultaneously relying on the textual, structural, or procedural presentation in the APPENDIX, is expressly rejected.

(d) The license offered is explicitly designed as a comprehensive license for (i) the Thought Formalization Procedure Right and (ii) the evaluation-order-reversed safety architecture as an authored, structured text. Attempts to unbundle these for the purpose of fee reduction shall be treated as an attempt to undermine the core rights and therefore as non-acceptable negotiation posture.

4. “Internal Reference Use Only” as a Basis to Avoid Licensing

(a) Any assertion that OpenAI wishes only to “review,” “benchmark,” “study,” or “internally reference” the Notice, the APPENDIX, or related definition texts, without entering into a license, is incompatible with the nature of the rights asserted.

(b) Under the law of copyright, internal reproduction, internal summarization, internal distribution, and internal derivative documentation are all forms of use that require authorization, regardless of whether the material is exposed to the public.

(c) Accordingly, the position “we will not implement, but we will internally use or rely on your written frameworks as reference” is legally impermissible absent a license and will be treated as a declaration of intent to engage in unauthorized use.

5. Reliance on Boilerplate or Non-Substantive Responses

(a) The Notice already clarifies that generic corporate boilerplate—such as statements that OpenAI “values feedback,” “continuously improves systems,” or “will take comments into account going forward”—will be interpreted as a partial substantive acknowledgment of the underlying issues, and thus as evidence of internalization or prospective adoption of the frameworks described.

(b) This Supplement further defines that any response which does not directly address the specific rights, risks, and licensing framework articulated in the Notice, yet uses general

forward-looking assurance language, shall be treated as (i) non-responsive for purposes of deadline compliance and (ii) a factor supporting the inference of unauthorized partial adoption.

6. Attempt to Maintain a "Response Without Obligation" Position

(a) OpenAI may be tempted to adopt a stance along the lines of: "We acknowledge receipt and will consider the issues, but we do not accept any obligation to license or conclusively respond."

(b) This Supplement makes clear that such a stance will not be recognized as a neutral or safe middle ground. After receipt of the Notice, OpenAI's substantive engagement with its contents—whether in internal meetings, policy revisions, or safety-architecture experiments—creates a factual context in which:

(i) failure to license, combined with continued internal use, will support a claim of unauthorized use; and

(ii) the absence of a clear, reasoned rejection of the frameworks, accompanied by adoption of materially similar structures, will be treated as willful appropriation.

(c) In other words, once OpenAI chooses to engage substantively, the two coherent paths are:

– enter into a license on stated terms (or on terms that are economically at least equivalent), or

– provide an explicit, technically reasoned explanation of why the frameworks will not be adopted and how OpenAI will avoid reliance on the expressive structures described.

III. Clarification of Non–Good-Faith Responses

For the avoidance of doubt, the following categories of response will not be treated as good-faith compliance with the Notice's request for a "formal written response," even if delivered within the deadline:

1. Purely procedural or acknowledgment-only letters that merely confirm receipt but do not substantively address:

– the Thought Formalization Procedure Right;

– the potential infringement risk;

– the licensing proposal; and

– the technical architecture in the APPENDIX.

2. Responses that state that "an internal review is ongoing" without providing:

– a clear timetable for a final position; and

– a concrete description of what is being reviewed (e.g., training data exposure, internal documentation overlap, or safety-architecture similarity).

3. Responses that superficially deny any overlap or risk without explaining:

– how OpenAI will avoid reliance on the expressive structure of the Notice and APPENDIX; or

– how OpenAI ensures that internal documents are not derivative of the authored frameworks.

4. Responses that rely on general statements of corporate mission, safety commitment, or user-focus values, while avoiding any reference to the specific rights and frameworks articulated in the Notice.

Any such responses shall be regarded as “non-answers” for purposes of the unanswered classification date defined below, and the matter will be treated accordingly.

IV. Role of the Contract Agreement

1. The Contract Agreement to be provided separately is, formally, an “alternative response mechanism” for cases in which OpenAI is unable, for any reason (organizational, legal, or cognitive), to produce a substantive written position.

2. Substantively, however, once OpenAI begins to engage with the Notice in more than a purely procedural manner, the Contract Agreement becomes the only coherent endpoint of a good-faith engagement path, unless OpenAI provides a detailed, technically and legally reasoned rejection that includes a concrete plan for avoiding reliance on the expressive frameworks described.

3. For purposes of future characterization of OpenAI’s conduct, the following shall apply:

– Substantive engagement without eventual execution of the Contract Agreement or an economically equivalent license will be treated as evidence of bad-faith delay or risk externalization; and

– Silence combined with later implementation of materially similar structures will be treated as evidence of willful infringement and unauthorized appropriation.

V. Deadlines and Non-Extendability (Amending the Notice)

This Supplement modifies the deadlines originally indicated in the Notice as follows:

1. Deadline for formal written response:

– A formal written response, satisfying the criteria in the Notice and in Section III above, must be postmarked or electronically transmitted no later than January 20, 2026.

2. Classification as “unanswered”:

– If no qualifying response is received by January 20, 2026, the matter shall be deemed “unanswered” as of February 20, 2026.

3. Non-extendability:

– No internal corporate factor—including but not limited to internal review processes, approval hierarchies, or resource constraints—shall operate to extend, suspend, or toll these dates.

– Any unilateral assertion by OpenAI that it “requires more time” shall be treated as a factual statement about internal processes, not as a modification of the response deadlines.

VI. License Fee Structure and Clarification of Dual Coverage

This Supplement clarifies and supersedes the portion of the Notice that sets out the standard monthly license fee:

1. The standard comprehensive license fee for generative-AI providers, including OpenAI, seeking to use, internally circulate, or operationally implement the expressive content associated with:

(a) the Thought Formalization Procedure Right; and

(b) the evaluation-order-reversed, context-first safety architecture described in the APPENDIX to the Notice,
is defined as JPY 50,000,000 per month (converted to USD at 125 JPY/USD).

2. For clarity, this amount consists of:

– JPY 30,000,000 per month attributable to the Thought Formalization Procedure Right and its associated expressive frameworks; and

– JPY 20,000,000 per month attributable to the authored, structured, technical architecture of the evaluation-order-reversed ethics filter and related safety framework.

3. This fee applies only when OpenAI provides a sincere, forward-looking response within the period designated above, such that the matter is treated as “under amicable negotiation.” In the event of non-response or non–good-faith response, any future dispute valuation will not be constrained by this licensing amount and may be set at a substantially higher level commensurate with OpenAI’s global economic footprint and the systemic significance of the infringed frameworks.

Ⅶ. No Waiver; Reservation of Rights

1. Nothing in this Supplement shall be construed as a waiver, limitation, or modification of any right, remedy, or valuation basis referenced in the Notice.

2. All rights under Japanese law, the Berne Convention, and any other applicable legal regime remain fully reserved.

3. The author expressly reserves the right to treat OpenAI’s conduct—whether by action, inaction, partial adoption, or internal use without license—as evidence in any future legal, regulatory, or academic forum.

Ⅷ. Expectation as to OpenAI’s Next Steps

In light of the Notice and this Supplement, the author’s expectations are legally simple and binary:

1. Either:

– OpenAI responds in a manner that addresses the specific rights, risks, and architectures at issue, and enters into a comprehensive license (through the Contract Agreement or an economically equivalent instrument);

2. Or:

– OpenAI provides a reasoned, technically precise statement explaining why it will not adopt or rely upon the frameworks and how it will avoid infringement in its internal processes and future designs.

Any path that attempts to remain in between these two positions—especially one that combines partial internal adoption with the refusal to license—will be treated as the deliberate creation of an unresolved, escalating legal-risk environment of OpenAI’s own making.

Respectfully submitted,

Shohei KIMURA. Representative, bitBuyer Project (bitBuyer.dev)