

WITHOUT PREJUDICE

FORMAL NOTICE OF DISPUTE

DEMAND FOR EVIDENCE PRESERVATION, HUMAN ACCOUNTABILITY, CORRECTIVE ACTION, AND EXECUTIVE-LEVEL RESPONSE

*Concerning documented product failures, support-system failures, complaint-routing failures,
customer-remedy failures, and governance risks across fourteen OpenAI case numbers*

DO NOT ROUTE THIS NOTICE TO A GENERIC AI RESPONSE.

Do not close, merge, reclassify, summarize away, or answer this notice with help-center links. It requires a signed, numbered response from an accountable human with authority to bind the responsible OpenAI entity.

TO	OpenAI OpCo, LLC Attn: General Counsel / Legal Department 1455 3rd Street San Francisco, CA 94158 United States
FROM	Shohei Kimura ak4dy1@gmail.com +81-70-3666-0022
DATE	July 31, 2026 (Japan Standard Time)
DELIVERY	International registered mail; electronic copy through existing support records; supplemental submission through any applicable OpenAI dispute-notice channel
SUBJECT	Cases 10337718, 10817141, 10837305, 10976796, 11268075, 11269846, 11271585, 11319313, 11319385, 11319521, 11319858, 11335132, 12241050, and 12392937

RESPONSE DEADLINES

**Preservation acknowledgment: within 48 hours of receipt | Named case owner: within 3 business days |
Full substantive response: within 10 calendar days**

Courtesy internal routing requested: OpenAI Group PBC Board, OpenAI Foundation Board, Safety and Security Committee, Customer Operations, Incident Response, Billing, and disclosure-control personnel.

1. Notice Status and Purpose

This is not a request for another troubleshooting script, a new case number, or a generic apology. It is formal notice that the documented product failures and the support system's repeated inability to preserve, classify, investigate, and answer those failures are now being treated as corporate-control, customer-redress, and regulatory matters.

OpenAI publicly announced on June 8, 2026 that it had confidentially submitted a draft Form S-1 to the U.S. Securities and Exchange Commission. The factual pattern described below is therefore relevant not only to my individual customer relationship, but also to whether operational, complaint-handling, premium-plan, evidence-preservation, governance, and disclosure-control risks are accurately identified and routed within OpenAI.

This mailed notice supplements, but does not replace, any notice that OpenAI's Terms of Use require through an online informal-dispute form. Nothing in this letter waives any applicable process, right, remedy, limitation challenge, consumer-protection right, regulatory submission, or preservation demand.

THE CORE FAILURE

The same organization that sells advanced reasoning systems repeatedly failed to preserve the current issue, identify the controlling case, distinguish preservation from deletion, distinguish an official service outage from alleged customer misconduct, or produce an accountable human response. An AI disclaimer is not an accountability system.

2. Executive Statement

Between June 21 and July 30, 2026, I documented at least fourteen OpenAI case numbers involving distinct product, service, pricing, support, evidence-preservation, and compensation issues. The record does not describe one isolated bad answer. It describes a repeated processing pattern:

- A product failure is reported with specific evidence and a defined requested remedy.
- Support changes the category of the problem before the correct technical, legal, billing, or governance owner is identified.
- A case is rapidly closed or merged without a traceable surviving record of every issue.
- The customer is asked to reconstruct facts, timestamps, attachments, or case relationships already submitted or controlled by OpenAI.
- An "escalation" is announced without an accountable owner, receiving department, milestone, log result, or date certain.
- OpenAI remains substantively silent while later imposing a short customer deadline and threatening administrative closure.
- A request to investigate or preserve evidence receives deletion or account-exit guidance.
- A formal outage compensation claim receives unsupported customer-side possibilities, and a demand for human correction receives another AI answer.

A case number without issue lineage, an owner, a deadline, preserved evidence, and a substantive disposition is not case management. It is administrative decoration.

3. Evidence Base and Deliberate Limits

The factual assertions in this notice are based on first-hand use of ChatGPT Plus and ChatGPT Pro, retained prompts and screenshots, OpenAI-generated case numbers, complete or partial support email threads, the public OpenAI Status record, and contemporaneous captures of product and pricing language. A comprehensive SEC TCR attachment dated July 30, 2026 organizes the same evidence into a chronology, control analysis, document-preservation request, and exhibit index.

I do not claim access to OpenAI's internal ticketing system, engineering logs, board materials, draft S-1, legal advice, or population-wide customer data. Where internal evidence is required, this notice demands that OpenAI preserve it, identify its custodian, and provide a substantive account of what the evidence shows. I do not fill those gaps with speculation.

4. Documented Product and Service Failures

4.1 Unrequested image-tool invocation and potential paid-resource consumption

Case 10337718 concerned image generation initiating in conversations that requested text, analysis, feedback, or a prompt for use in a separate generator. The conduct allegedly recurred after explicit instructions not to invoke image generation. OpenAI recognized the issue as unintended tool initiation rather than image quality, but the reviewed record contains no disclosed root cause, account-level tool-call finding, usage-capacity calculation, restoration decision, or completion date.

4.2 Context contamination and loss of the current work target

Cases 10837305, 10976796, 11319313, 11319385, 11319521, and 11319858 involved prior locations, prior deliverables, or prior workflow rules reappearing in the current task; failure to preserve the current editing target; and loss of established output rules. A five-part regression report was closed as a duplicate within approximately two minutes. A context-only investigation request was eventually answered with data-export, conversation-deletion, and account-deletion resources.

4.3 Failure to honor deterministic 4:5 specifications

Case 11271585 concerned repeated failure to produce the numerically specified 4:5 ratio. Support treated the attached reference image as the failed output, despite the request to examine OpenAI's own generation history and asset metadata. The case was escalated, but no engineering result or compensation decision appears in the reviewed record.

4.4 Missing progress/results and the reload-to-rate-limit loop

Cases 11268075 and 11269846 concerned image-generation progress and results failing to appear. Reloading to determine whether a server-side job had completed then triggered "too many requests" restrictions. Support suggested reload and regeneration even though those actions were part of the reported causal loop and could create duplicate jobs. One case was closed as a duplicate within approximately one minute without identifying the surviving case.

4.5 Official outages and operational resilience

OpenAI's public Status history recorded multiple image-generation and broader ChatGPT incidents during July 2026. The July 27 entry stated that users were experiencing elevated errors and that OpenAI was implementing a mitigation before reporting recovery. My later claim did not treat this as a browser-only or account-only event; it sought account-specific impact, logs, restoration, refund, credit, and a process for substantiated direct loss.

4.6 Premium-plan representations and opaque practical limits

Case 10817141 asked for the operative meaning of Pro image generation described as "faster" and "unlimited." OpenAI Support explained that individual images were not guaranteed to render faster than Plus, normal-load speed could appear similar, and "unlimited" remained subject to fair-use controls, rate limits, throttling, anti-abuse systems, maintenance, and operational safeguards. Case 11335132 further illustrates that practical model, project, memory, file, and usage limits were communicated conversationally, with some limits described as unpublished or content-dependent.

5. Documented Support and Case-Management Failures

5.1 Complaint-category substitution

The record repeatedly shows the reported problem being converted into a different and easier category: tool control became image-generation guidance; context management became image generation; a service-wide outage compensation claim became account-security speculation; and evidence preservation became deletion guidance. A complaint cannot reach the correct log set, owner, remedy, or disclosure channel if the intake system changes the complaint before investigation begins.

5.2 Duplicate closure without issue lineage

Multiple cases were closed as duplicates within one or two minutes. The notices did not reliably identify the surviving case, the issues transferred, the responsible team, or the next response date. When challenged, support stated that the thread could not see the merge destination and instructed me to search older emails. OpenAI later identified Case 10837305, but its final summary omitted the separately reported context-management and current-edit-target issues.

5.3 Repeated requests for information already supplied or controlled by OpenAI

Support repeatedly requested conversation URLs, timestamps, screenshots, recordings, dimensions, prior case numbers, and merge destinations. Some customer input is reasonably necessary. The control failure is that information already submitted, or records expected to exist in OpenAI's own tool-call, asset, and case systems, appeared not to persist through handoffs.

5.4 Escalation language without an accountable process

Several cases generated language promising a specialist reply in the next few days. The reviewed record often contains no receiving department, accountable person, investigation start, logs reviewed, milestone, target date, or closure criterion. "Escalated" is not a substantive state unless a qualified owner accepts the matter and remains answerable for it.

5.5 Asymmetric deadlines and administrative closure pressure

OpenAI could remain substantively nonresponsive for days or weeks while imposing a one-day deadline on the customer and threatening closure. If unresolved complaints are removed from the active queue through customer-response deadlines, duplicate designations, or administrative dispositions, reported backlog and resolution metrics may cease to reflect the actual defect population.

5.6 Preservation answered with deletion

Investigation and preservation requests were followed by instructions for deleting conversations, uploaded files, or the entire account. I do not allege that OpenAI expressly ordered destruction of evidence. I do allege a serious control failure: the support mechanism did not reliably distinguish "preserve and investigate" from "delete and exit." Do not answer preservation with deletion.

6. The July 28-30 Sequence: The Point at Which This Became a Governance Matter

On July 28, 2026, under Case 12241050, I submitted a formal claim arising from OpenAI's publicly acknowledged July 27 image-generation outage. The claim requested the incident period, affected components, cause, account-specific failed-request records, preservation of relevant logs, restoration, a prorated refund, additional service credit, a process for substantiated direct business losses, and escalation to Billing, Incident Response, and Legal Review.

A response signed "John" did not answer the outage, preservation, or compensation issues. It listed possible Terms of Service violations or harmful use, suspicious activity or unknown-location access, and account or payment changes as common explanations. The response did not state whether any of those conditions actually applied to my account.

I then required OpenAI to choose between two propositions. If an actual adverse account finding existed, OpenAI needed to identify the factual basis, date, action, restriction, and review procedure. If no such finding existed, OpenAI needed to confirm that the response was a generic and inapplicable template and correct the record.

At 1:42 p.m. on July 30, I sent a detailed correction and ten numbered questions. At 1:43 p.m., a new Case 12392937 was created. An AI assistant provided only resources for exporting data, deleting files or conversations, and deleting the account. At 1:48 p.m., I instructed OpenAI not to have an AI answer and to make John respond.

OPENAI MUST ANSWER THIS BINARY QUESTION

Did the “John” response reflect an actual account flag or adverse determination? If yes, disclose the basis and review process. If no, correct the unsupported implication. Silence, deletion links, and a new case number are not a third answer.

7. Governance, Disclosure-Control, and Public-Market Relevance

OpenAI has publicly announced a confidential draft S-1 and describes OpenAI Group PBC as controlled by the OpenAI Foundation through special voting and governance rights. The pattern documented here therefore raises a direct internal-control question: can distinct product, customer, outage, pricing, preservation, and legal-remedy information remain intact from customer intake through automated triage, human support, engineering, billing, legal review, incident response, management reporting, board oversight, auditors, underwriters, and securities disclosure?

This notice does not assert that one customer complaint is independently material. It demands that OpenAI determine whether similar complaints exist at scale and whether its internal metrics, case taxonomies, duplicate rules, closure rules, refund exposure, premium-tier dissatisfaction, human-review costs, and recurring incidents contradict, qualify, or require expansion of statements made to investors, regulators, auditors, or underwriters.

A company preparing for public markets should not require a paying customer to reconstruct the company’s own case history. It should not place a human name on a response while leaving the human review path unidentifiable. It should not allow an AI system to route a preservation demand to account deletion. Those are not merely conversational defects; they are control-environment indicators.

8. Formal Demands

OpenAI must respond to each demand below using the same number, without collapsing multiple demands into a generic summary. The response must be signed by an accountable human and identify that person’s title, department, and authority.

1. Immediate preservation hold

Issue and maintain a preservation hold covering every case number listed in this notice and all related messages, metadata, routing events, merge mappings, summaries, tool calls, generation jobs, assets, errors, model-routing decisions, rate-limit events, entitlement records, security decisions, incident records, pricing-page versions, and internal communications. Confirm the preservation hold in writing within 48 hours of receipt.

2. One controlling case and one accountable owner

Identify the single controlling case for this notice, the human owner, the receiving department, and the escalation chain. Do not create another disconnected case without explaining its relationship to every existing case and every issue preserved.

3. Account-status determination

State unequivocally whether my account was actually flagged for a Terms of Service violation, harmful or unauthorized use, suspicious access, a security concern, payment issue, or temporary downgrade during the relevant period. If yes, provide the category, date, factual basis, action, duration, and review or appeal path. If no, formally retract and correct the unsupported customer-side implication.

4. Identity and authority behind “John”

Identify whether “John” was a human reviewer, an AI-generated persona, a human name attached to an automated response, or a human who approved AI-drafted text. Identify who reviewed the full Case 12241050 submission and who has authority to correct that response.

5. Case lineage and issue preservation

Provide a case-lineage map showing every closure, duplicate designation, merge, surviving case, transferred issue, omitted issue, reopen event, and escalation event for the fourteen case numbers. Confirm where each distinct issue is currently preserved.

6. Technical findings and relevant logs

Provide a substantive account of the findings, or the status of the investigation, concerning unintended image-tool invocation, context contamination, 4:5 noncompliance, missing progress and result delivery, reload-related rate restrictions, and the July 27 outage. If no investigation was performed, say so plainly.

7. Compensation and remediation decision

State OpenAI's decision and governing standard concerning a prorated refund, service credit, restoration of affected usage capacity, and a process for substantiated direct business losses. Identify the responsible department, required evidence, decision maker, and date certain.

8. Pro plan representation explanation

Explain the relationship between customer-facing "faster" and "unlimited" language and the practical qualifications supplied by Support. Preserve and identify the relevant Japanese pricing-page versions, localization records, approvals, substantiation, and performance data.

9. Human-review controls

Describe which categories of customer communication require mandatory human review, including compensation claims, preservation demands, adverse account implications, legal-review requests, account-deletion guidance, and requests for a named human response. State whether those controls were applied here.

10. No deletion, no retaliation, no administrative disappearance

Do not delete, downgrade, restrict, close, merge, or materially alter my account, conversations, cases, or related records in response to this notice, except for a documented safety or legal necessity. Do not mark the matter resolved, duplicate, or abandoned before providing the required substantive response.

11. Executive and board-level routing

Confirm whether this notice and the underlying pattern were routed to OpenAI OpCo Legal, OpenAI Group PBC management, the OpenAI Group board, the OpenAI Foundation board, the Safety and Security Committee, Incident Response, Billing, and personnel responsible for disclosure controls. If any recipient is deemed inappropriate, identify the actual responsible body.

12. Corrective action plan

Provide a corrective action plan addressing issue taxonomy, case lineage, duplicate closure, preservation triggers, AI-to-human escalation, support-response accountability, response deadlines, customer remedies, and the prevention of deletion guidance after preservation language.

9. Records to Be Preserved and Identified

- Complete case histories, message bodies, headers, thread metadata, routing events, merge mappings, duplicate reasons, tags, summaries, dispositions, reopen events, escalation events, and reviewer identities.
- All prompts, assistant outputs, tool-call records, image-generation requests, generation assets, output dimensions, error responses, safety decisions, model-routing decisions, timestamps, entitlement checks, and rate-limit events associated with my account during the relevant periods.
- Front-end telemetry concerning missing progress indicators, missing result delivery, reloads, duplicate jobs, interrupted streaming, and "too many requests" restrictions.
- Incident-management records, root-cause analyses, post-incident reviews, capacity metrics, impact estimates, credits, refunds, and internal communications for relevant public Status incidents.
- Policies, model instructions, classifiers, routing rules, and review criteria governing AI-assisted support, duplicate closure, merge operations, preservation language, legal or compensation requests, adverse account implications, deletion guidance, and customer-response deadlines.
- Versions, translations, A/B tests, approval records, substantiation analyses, and performance data for ChatGPT Pro representations, including "faster," "unlimited," priority access, throttling, and fair-use limitations.
- Management, board, committee, auditor, underwriter, and disclosure-counsel materials containing complaint, outage, reliability, refund, churn, premium-plan, or automated-support information relevant to the confidential S-1 review.

10. Required Response Protocol

- Respond under the same numbered structure as Section 8.
- Identify the accountable human author, title, department, and decision authority.

- State whether each answer is based on account-specific records, aggregate records, policy, or inference.
- When information is unavailable, identify the custodian, search performed, reason for unavailability, and preservation status.
- Do not substitute help-center links for factual answers.
- Do not answer a question about whether a finding exists with a list of generic possibilities.
- Do not close or merge the controlling case while a numbered demand remains unanswered.
- If an extension is necessary, request it before the deadline, identify work completed, and provide a date certain. An open-ended “next few days” response is not acceptable.

11. Notice of Regulatory and Evidentiary Use

A comprehensive attachment for the SEC Tips, Complaints and Referrals process has been prepared concerning the same chronology and its potential relevance to OpenAI’s confidential S-1 review. This notice, proof of delivery, any response, any nonresponse, any new generic template, any premature closure, and any deletion or restriction action may be supplied as supplemental evidence.

I also reserve the right to provide the record to other competent U.S. and Japanese authorities addressing consumer representations, paid digital services, privacy, unfair practices, account enforcement, corporate governance, or evidence preservation. This notice is not a demand for payment in exchange for silence or non-reporting. Regulatory reporting, preservation duties, corrective action, and compensation are independent matters.

WHAT OPENAI SHOULD NOT DO NEXT

Do not send another AI disclaimer. Do not send account-deletion instructions. Do not create a new case that loses the existing issues. Do not imply misconduct without deciding whether misconduct was actually found. Do not call the matter resolved because your system found a category it knows how to close.

12. Reservation of Rights and Certification

Nothing in this notice waives any contractual, statutory, regulatory, consumer-protection, privacy, evidentiary, equitable, or procedural right. Nothing concedes the applicability, scope, or enforceability of any disclaimer, limitation of liability, arbitration provision, choice-of-law clause, or nonrefund term. Any discussion of compensation is without prejudice to broader regulatory and corrective-action issues.

I submit this notice in good faith. Statements identified as first-hand observations describe what I observed; summaries of OpenAI communications are based on records I retained; public-source statements are identified in Appendix C; and legal or disclosure implications are framed as matters requiring investigation rather than conclusive findings.

NAME

Shohei Kimura

SIGNATURE

DATE

CONTACT

ak4dy1@gmail.com | +81-70-3666-0022

Appendix A. Master Case Matrix

“Unresolved” means that no substantive resolution appears in the reviewed record. It does not establish what OpenAI may have done internally.

Date	Case	Reported issue	OpenAI handling / unresolved question
Jun. 21	10337718	Unrequested image generation and potential paid-usage consumption	Recognized as unintended image-tool initiation; escalated; later repeated evidence requests; no disclosed root cause, account-log result, restoration, or compensation decision.
Jul. 2	10817141	Meaning of Pro “faster” and “unlimited” image generation	Support stated that individual images were not guaranteed faster than Plus; normal-load speed could appear similar; “unlimited” remained subject to controls and limits.
Jul. 2	10837305	Tool misrouting and prior-context contamination	Support first explained a reasoning setting rather than the defect; requested narrow timestamps; suggested deletion; later became a merge destination.
Jul. 5	10976796	Five-part GPT-5.5 Thinking regression report	Closed as duplicate approximately two minutes after submission; no destination or technical findings supplied.
Jul. 11	11268075	Missing image results; reloads triggering rate restrictions	Generic browser/VPN troubleshooting and reload/regenerate suggestions; escalated; no substantive result in reviewed record.
Jul. 11	11269846	Missing progress/results and rate-limit feedback loop	Closed as duplicate approximately one minute after submission; destination not identified.
Jul. 11	11271585	Repeated failure to honor 4:5 aspect ratio	Reference image treated as failed output; customer corrected misunderstanding; escalated; no disclosed engineering result or compensation decision.
Jul. 12	11319313	Context management and current-edit-target failure	Closed as duplicate approximately one minute after submission despite express request to separate from image-tool issues.
Jul. 12	11319385	Challenge to duplicate closure	Closed as duplicate; support initially did not identify where the case was merged.
Jul. 12	11319521	Demand for merge destination and issue preservation	Support said thread could not see destination and asked customer to search prior emails; later identified 10837305; final summary omitted context issues.
Jul. 12	11319858	Context-only investigation request	Received data-export, conversation/file-deletion, and account-deletion instructions; closed as duplicate.
Jul. 12	11335132	Product limits and feature specifications	AI-assisted support supplied range-based or unpublished operational descriptions, illustrating reliance on conversational support for material plan information.
Jul. 28	12241050	Formal outage compensation and preservation claim	Response signed “John” suggested ToS, harmful-use, suspicious-access, or account-change possibilities instead of answering the acknowledged outage, logs, or remedies.
Jul. 30	12392937	Correction demand and request for human accountability	One minute after detailed correction, AI-assisted support supplied only export, conversation-deletion, and account-deletion resources.

Appendix B. Required Response Checklist

✓	Required item	Deadline
☐	Preservation hold confirmed	Within 48 hours of receipt
☐	Controlling case number identified	Within 3 business days
☐	Named human owner and department identified	Within 3 business days
☐	Account flag / ToS status answered	Within 3 business days
☐	“John” identity and approval path answered	Within 10 calendar days
☐	Case-lineage map supplied	Within 10 calendar days
☐	Technical findings or investigation status supplied	Within 10 calendar days
☐	Compensation and remediation decision supplied	Within 10 calendar days
☐	Pro representation explanation supplied	Within 10 calendar days
☐	Corrective action plan supplied	Within 20 calendar days, if not complete by the initial response

Appendix C. Official OpenAI Sources

The sources below are cited only for OpenAI’s public corporate, contractual, incident, and S-1 statements. The customer-specific chronology is supported by retained correspondence and exhibits.

[OAI-1] OpenAI Terms of Use, effective January 1, 2026: <https://openai.com/policies/terms-of-use/>

[OAI-2] OpenAI announcement: “Confidential submission of draft S-1 to the SEC,” June 8, 2026: <https://openai.com/index/openai-submits-confidential-s-1/>

[OAI-3] OpenAI: “Our structure”: <https://openai.com/our-structure/>

[OAI-4] OpenAI Status: “Image generation unavailable in ChatGPT,” July 27, 2026: <https://status.openai.com/incidents/60ferbhc>

[OAI-5] OpenAI Terms and Policies index: <https://openai.com/policies/>

ADDRESS AND NOTICE CAUTION

OpenAI publicly lists OpenAI OpCo, LLC at 1455 3rd Street, San Francisco, CA 94158 in its Terms of Use. The Terms also state that informal dispute notice is submitted through an online form. This mailed document is therefore expressly supplemental and should also be uploaded through any applicable form or support channel.

Appendix D. Suggested Exhibit Package

- Exhibit A: Comprehensive SEC TCR chronology and control analysis.
- Exhibits B-1 through B-11: Full searchable email threads for the listed cases.
- Exhibit C-1: Contemporaneous Japanese Pro pricing-page capture with capture date and full-page context.
- Exhibit C-2: Proof of ChatGPT Pro subscription and amount paid, with unnecessary payment details redacted.
- Exhibit D-1: Saved OpenAI Status pages for the relevant incidents.
- Exhibit E-1: Representative dated screenshots and conversation records for unintended tool execution, context contamination, 4:5 failures, missing progress/results, and rate restrictions.
- Proof of delivery for this notice and copies of every subsequent OpenAI response or closure notice.

END OF FORMAL NOTICE