

DRAFT ATTACHMENT TO AN SEC
TIPS, COMPLAINTS AND REFERRALS SUBMISSION

OpenAI Group PBC

Potentially Material Operational, Disclosure-Control, Customer-Redress, Product-Reliability, and Governance Issues

Relevant to the Confidential Form S-1 Review Process

A documented external-customer chronology and request for regulatory inquiry

PREPARED FOR	U.S. Securities and Exchange Commission
PREPARED BY	Shohei Kimura
CAPACITY	External paying customer; ChatGPT Pro subscriber; no claim of insider access
CONTACT	ak4dy1@gmail.com +81-70-3666-0022
LOCATION	
INFORMATION CURRENT THROUGH	July 30, 2026 (Japan Standard Time)

IMPORTANT SCOPE STATEMENT

This document does not ask the SEC to adjudicate a private refund dispute. It asks the SEC to assess whether the documented pattern, together with OpenAI's acknowledged confidential S-1 submission, raises potentially material disclosure, risk-factor, known-trend, governance, operational-resilience, complaint-handling, or customer-redress issues. Legal characterizations are framed as questions for investigation, not findings of violation.

DOCUMENT MAP

Contents

1.0	Portal-ready summary
2.0	Purpose, scope, and evidentiary method
3.0	Why the submission is relevant to the SEC's confidential S-1 review
4.0	Entity and product background
5.0	Executive summary of the documented pattern
6.0	Master chronology and case matrix
7.0	Detailed factual statement of product and service issues
8.0	Detailed factual statement of support and case-management failures
9.0	Cross-case control failures
10.0	Potential disclosure relevance for SEC staff review
11.0	Questions and requested investigative steps
12.0	Evidence-preservation request
13.0	Exhibit index and submission package checklist
14.0	Certification
Appendix A	Public OpenAI service-incident context
Appendix B	Key correspondence excerpts
Appendix C	Authorities and public sources

USE Upload this document as an attachment in Step 5 of the SEC TCR portal. Enter the concise summary in Section 1.0 into the narrative field, then attach the underlying email threads, screenshots, pricing-page capture, subscription proof, and status-page records listed in Section 13.0.

1.0 Portal-Ready Summary

OpenAI publicly announced on June 8, 2026 that it had submitted a confidential Form S-1 to the SEC. I am an external ChatGPT customer and submit this information so SEC staff can assess whether OpenAI's draft registration statement adequately discloses potentially material operational, product-reliability, customer-redress, complaint-handling, marketing-substantiation, disclosure-control, and governance risks. I do not ask the SEC to resolve my individual refund claim, and I do not assert that every underlying product defect independently violates the federal securities laws.

Between June 21 and July 30, 2026, I documented repeated ChatGPT and OpenAI Support failures across at least fourteen case numbers. The product allegedly initiated paid image generation when no image was requested, repeated the conduct after explicit prohibition, lost the current task and reintroduced prior content, failed numerical 4:5 aspect-ratio requirements, failed to display image-generation progress or results, and produced a feedback loop in which reloading to determine whether a generation had completed triggered "too many requests" restrictions. OpenAI's public Status page separately recorded multiple image-generation and ChatGPT incidents during the same period.

The support process repeatedly reclassified or erased the substance of complaints. Cases were closed as duplicates within one or two minutes without identifying the surviving case; support then stated it could not see the merge destination and required the customer to search old emails. A context-management complaint was ultimately summarized as an image-generation complaint. Escalation notices promised a response within "the next few days," yet substantive results, owners, log findings, and completion dates were not supplied. In one case, OpenAI remained substantively silent for approximately twenty days while later imposing a one-day customer response deadline and threatening closure. Requests to preserve or investigate records were answered with data-export, conversation-deletion, and account-deletion instructions.

OpenAI Support also stated that ChatGPT Pro image generation described as “faster” did not guarantee that individual images would be faster than Plus and could appear similar under normal load; “unlimited” remained subject to fair-use controls, rate limits, throttling, anti-abuse systems, maintenance, and operational safeguards. I request that the SEC determine whether the contemporaneous pricing presentation, the actual plan specifications, and related limitations were accurately and adequately reflected in OpenAI’s offering disclosures and risk factors.

The most serious later sequence arose from OpenAI’s officially acknowledged July 27 image-generation outage. I submitted a formal compensation and evidence-preservation claim under Case 12241050. A reply signed “John” did not address the outage or compensation and instead suggested possible Terms of Service violations, harmful use, suspicious logins, or account changes. When I demanded confirmation whether any such finding actually existed, a new AI-assisted Case 12392937 responded only with instructions to export data, delete conversations, and delete the account. This sequence raises questions about automated support authority, human accountability, adverse account information, evidence preservation, complaint routing, legal-claim handling, refund exposure, and whether material customer and operational risks are captured and escalated to management, the board, auditors, underwriters, and SEC disclosure counsel.

REQUESTED SEC REVIEW Please compare the documented conditions against OpenAI’s confidential draft S-1, supporting diligence materials, board reporting, complaint and incident metrics, refund and credit policies, marketing substantiation, disclosure controls, and representations to auditors and underwriters. If appropriate, request tailored disclosure under Form S-1, Regulation S-K Items 101, 103, 105 and 303, and Securities Act Rule 408; require preservation of relevant records; and refer non-securities consumer-protection issues to the appropriate authorities.

2.0 Purpose, Scope, and Evidentiary Method

2.1 Purpose

The purpose of this submission is to give SEC staff a structured external record that may assist the nonpublic review of OpenAI’s draft registration statement. The central question is not whether a single image-generation defect is material in isolation. The central question is whether repeated paid-product failures, repeated service incidents, opaque plan limitations, automated complaint handling, loss of case content, lack of traceable escalation, unclear compensation processes, and weak responsibility assignment collectively indicate risks or known uncertainties that should be investigated and, if material, disclosed to prospective investors.

2.2 Scope

The primary evidentiary period is June 21 through July 30, 2026. The record includes first-hand use of ChatGPT Plus and ChatGPT Pro, contemporaneous emails to and from support@openai.com, case numbers generated by OpenAI, screenshots and prompts retained by the submitter, OpenAI’s public Status pages, OpenAI’s pricing and help materials, and OpenAI’s public announcement of its confidential S-1 submission.

The phrase “comprehensive” in this document means comprehensive within the reviewed evidence and the submitter’s documented experience. It does not mean that the submitter has access to OpenAI’s internal ticketing system, engineering logs, financial records, board materials, draft S-1, underwriter diligence, legal advice, or full customer population data.

2.3 Evidence classification

Evidence class	Meaning and limitation
Category A - First-hand observation	Actions and outputs personally observed by the submitter in ChatGPT, including unintended tool execution, context contamination, aspect-ratio failures, missing results, and rate restrictions.

Evidence class	Meaning and limitation
Category B - OpenAI correspondence	Emails sent by OpenAI Support, including case numbers, duplicate closures, escalation notices, plan explanations, requests for repeated information, and deletion guidance.
Category C - OpenAI public records	OpenAI corporate pages, pricing/help materials, and public Status incident records.
Category D - Regulatory inference	Questions about materiality, disclosure, controls, governance, and securities-law relevance. These are requests for SEC analysis, not asserted legal findings.

2.4 Deliberate limits on the submission

- The submitter does not claim to know the contents of the confidential draft S-1.
- The submitter does not assert that every product defect or support failure is individually material.
- The submitter does not assert that a customer-service dispute is itself a federal securities-law violation.
- The submitter asks the SEC to determine whether management knew or should have known of broader trends and whether material information was omitted, minimized, generically described, or inadequately routed into the disclosure process.
- Where OpenAI's internal logs, metrics, or decisions are unknown, this document identifies the missing evidence and requests that the SEC obtain it rather than filling the gap with speculation.

3.0 Why This Submission Is Relevant to the SEC's Confidential S-1 Review

On June 8, 2026, OpenAI announced that it had submitted a confidential S-1 to the SEC. OpenAI stated that it had not decided the timing of a public offering but that the submission gave it the option to go public sooner. Because the draft is nonpublic, external customers cannot evaluate whether its risk factors, business description, MD&A, governance disclosures, legal and regulatory exposure, or operational controls are sufficiently tailored.

The SEC's TCR process is intended to receive information concerning possible violations of the federal securities laws, including possible false or misleading statements about a company or SEC filings. The SEC asks for truthful, complete information, a detailed description of who did what and when, and relevant documentation. This document is organized to meet that evidentiary expectation.

Regulation S-K Item 105 requires disclosure of material factors that make an investment speculative or risky, with tailored rather than generic discussion. Item 303 requires discussion of material events and uncertainties known to management that are reasonably likely to affect future operations, revenues, income, liquidity, or costs. Securities Act Rule 408 requires further material information necessary to make required statements not misleading. The SEC should determine whether the documented pattern, if replicated at scale or known internally, bears on those obligations.

JURISDICTIONAL FOCUS This submission is not framed as "OpenAI owes me a refund, therefore the SEC should intervene." It is framed as: "OpenAI has entered a confidential S-1 review while external evidence indicates repeated operational, product, marketing, complaint-handling, evidence-preservation, and customer-redress weaknesses; the SEC should determine whether those risks and known trends are material and properly disclosed."

4.0 Entity and Product Background

4.1 Corporate structure

OpenAI states that its for-profit operating entity is OpenAI Group PBC, a public benefit corporation, and that the OpenAI Foundation controls OpenAI Group through special voting and governance rights. OpenAI states that the Foundation appoints all members of the OpenAI Group board and may replace directors. This structure makes the flow of product-risk, customer-complaint, regulatory, and operational information to both boards a relevant governance question.

4.2 Customer relationship

The submitter is a Japan-based professional user who uses ChatGPT for software engineering, writing, research, and a high-frequency multilingual publication workflow. During the initial June 21 incident the plan was ChatGPT Plus. The submitter subsequently upgraded and, by July 2026, paid approximately JPY 18,000 per month for ChatGPT Pro. Image generation was a regular component of the submitter's professional production process.

4.3 Products and functions implicated

- ChatGPT conversational models and context management.
- Automatic tool routing and image-generation invocation.
- ChatGPT image generation, aspect-ratio compliance, progress display, and result delivery.
- ChatGPT Pro plan marketing, speed, usage limits, rate limits, and service priority.
- Projects, file handling, memory, context windows, and public specification of practical limits.
- OpenAI Support's AI-assisted triage, duplicate closure, escalation, human review, case ownership, and compensation handling.
- OpenAI public incident reporting and the relationship between aggregate status information and account-specific support responses.

5.0 Executive Summary of the Documented Pattern

- **Unrequested paid-function execution:** Image generation allegedly ran in conversations that requested text, analysis, or prompt drafting only, including after explicit instructions not to generate an image.
- **Current-task and context integrity failure:** The system allegedly returned to prior locations, prior deliverables, or prior workflow rules instead of following the current request.
- **Deterministic specification failure:** Numerically specified 4:5 output requirements were allegedly not followed, making outputs unusable for the intended format.
- **UI state and rate-limit feedback loop:** Progress and results allegedly disappeared; reloading to determine whether generation completed then triggered "too many requests" restrictions affecting ordinary use.
- **Recurring service availability incidents:** OpenAI's Status history recorded multiple image-generation and broader ChatGPT incidents during the same period.
- **Paid-plan representation gap:** Support explained that "faster" did not guarantee faster individual rendering and "unlimited" remained subject to multiple limitations, raising a question whether pre-purchase wording conveyed the same practical meaning.
- **Opaque product limits:** Practical project, memory, context, and model-usage limits were described through AI-assisted support rather than clear, stable, account-specific product documentation.
- **Complaint-category substitution:** Tool control was treated as image quality; context management was treated as image generation; formal compensation was treated as account security; evidence preservation was answered with deletion guidance.
- **Duplicate closure without traceability:** Cases were closed within minutes as duplicates without identifying the surviving case or confirming which issues were preserved.
- **Loss of issue content across cases:** OpenAI could not initially identify the merge destination and later summarized a context-management complaint as an image-generation complaint.

- **Escalation without accountable process:** Notices promised specialist replies in a few days, but owners, departments, log results, milestones, and completion dates were not supplied.
- **Asymmetric response deadlines:** OpenAI could remain substantively nonresponsive for weeks while imposing a one-day deadline on the customer and threatening closure.
- **AI/human accountability ambiguity:** A response signed by “John” introduced possible misconduct explanations; a demand for John or another human to answer was met by a separate AI-assisted case.
- **Customer-remedy opacity:** No clear process or decision standard was supplied for prorated refunds, service credits, paid-resource restoration, or substantiated business losses.
- **Evidence-preservation tension:** Requests for log preservation and account-specific investigation were followed by guidance on deleting conversations or the entire account.

6.0 Master Chronology and Case Matrix

The entries below summarize the evidence reviewed. “Unresolved” means no substantive resolution appears in the reviewed record; it does not establish what OpenAI may have done internally.

Date	Case	Reported issue	Support handling / unresolved question
Jun. 21	10337718	Unrequested image generation and possible paid usage consumption	Accepted as unintended image-generation invocation; escalated; later repeated requests for already supplied evidence; no disclosed root cause, log result, restoration, or compensation decision.
Jul. 2	10817141	Meaning of Pro “faster” and “unlimited” image generation	Human-signed support said individual images are not guaranteed faster than Plus; normal-load speed may appear similar; “unlimited” remains subject to limits and safeguards.
Jul. 2	10837305	Tool misrouting and prior-context contamination	Support first explained the meaning of a reasoning setting rather than the reported defect; requested narrow timestamps and suggested deletion; case later became the merge destination for other issues.
Jul. 5	10976796	Five-part GPT-5.5 Thinking regression report	Closed as a duplicate approximately two minutes after submission; no merge destination or technical findings supplied.
Jul. 11	11268075	Missing image results; reloads triggering rate restrictions	Generic browser/VPN troubleshooting and reload/regenerate suggestions; escalated; no substantive result in the reviewed record.
Jul. 11	11269846	Missing progress and result display; rate-limit feedback loop	Closed as a duplicate approximately one minute after submission merge destination not identified.
Jul. 11	11271585	Repeated failure to honor 4:5 aspect ratio	Reference image was treated as the failed output; customer corrected the error; escalated; no disclosed engineering result or compensation decision.
Jul. 12	11319313	Context-management and current-edit-target failure	Closed as duplicate approximately one minute after submission despite an express request to separate it from image-tool issues.
Jul. 12	11319385	Challenge to duplicate closure	Closed as duplicate; support did not initially identify where the case was merged.
Jul. 12	11319521	Demand for merge destination and issue preservation	Support said the thread could not see the destination and asked the customer to search prior emails; later identified 10837305; final summary omitted the context-management issue.
Jul. 12	11319858	Context-only investigation request	Received data-export, file/conversation-deletion, and account-deletion instructions; closed as duplicate.
Jul. 12	11335132	Product limits and feature-specification inquiries	AI-assisted support supplied nonpublic or range-based descriptions of model windows, project/file context, and usage limits, illustrating reliance on conversational support for material plan specifications.

Date	Case	Reported issue	Support handling / unresolved question
Jul. 28	12241050	Formal outage compensation and evidence-preservation claim	Reply signed “John” suggested possible ToS violation, harmful use, suspicious access, or account changes instead of addressing the acknowledged outage, logs, refunds, credits, or business-loss process.
Jul. 30	12392937	Correction demand and request for human accountability	One minute after a detailed correction, AI-assisted support provided only export, conversation-deletion, and account-deletion resources.

7.0 Detailed Factual Statement of Product and Service Issues

7.1 Unrequested image-generation execution and potential paid-resource consumption

On June 21, 2026, the submitter reported that ChatGPT initiated image generation in ordinary conversations, feedback exchanges, analysis, and tasks whose requested output was a text prompt for use in a separate image generator. The submitter states that the behavior recurred after he expressly prohibited image generation and after ChatGPT acknowledged that it should not invoke the image tool absent an explicit generation request.

The alleged harm was not limited to an unwanted image. Image generation involved time, usage capacity, and potential rate limits. The submitter requested preservation of tool-call records, confirmation of resource consumption, restoration or compensation, a technical explanation, and prevention of recurrence. OpenAI accepted Case 10337718 as a report that image generation may have been unintentionally triggered and may have affected usage availability. A response from “Warner” later recognized the issue as unintended tool initiation rather than image quality, but did not disclose a root cause, account-log finding, resource calculation, remediation, or completion date.

“We also understand your concern that image generation may have been initiated when your intention was only to receive text output. We’ve documented this as part of your report for review.”

OpenAI Support, Case 10337718 (excerpt)

Potential SEC relevance: if unrequested paid-function execution occurred at scale, it could affect customer value, usage-capacity accounting, churn, credits, refund exposure, infrastructure demand, and the accuracy of claims about user control and paid access. SEC staff should determine whether OpenAI internally measured this defect and whether the results entered its disclosure process.

7.2 Context contamination, current-task loss, and model-regression allegations

The submitter separately reported that ChatGPT failed to maintain the current editing target and returned to prior tasks, locations, or workflow instructions. Examples documented in the record included content from a prior “Fuji Five Lakes glamping” prompt appearing in a new Reykjavik or Antwerp-related task, a prior flower-arrangement video prompt being returned during another editing workflow, and repeated failure to honor the established meaning of the user’s instruction to provide a corrected final deliverable without meta-explanation.

On July 5, the submitter sent an English-language regression report concerning GPT-5.5 Thinking. It listed five categories: prohibited image-tool invocation, prior-location contamination, loss of established 4:5 bromide conditions, geometrically implausible correction strategies, and insertion of explanations or apologies when only a finished deliverable was requested. Case 10976796 was closed as a duplicate approximately two minutes later without a disclosed destination or technical finding.

Potential SEC relevance: context integrity and instruction-following are core attributes of a paid AI product. Persistent regressions may affect product differentiation, professional reliability, enterprise adoption, support costs, customer retention, and reputational risk. The SEC should determine whether management had internal regression metrics, incident reports, or customer-complaint trends not reflected in the draft S-1.

7.3 Failure to honor deterministic aspect-ratio specifications

Case 11271585 concerned repeated output in a ratio other than 4:5 despite explicit mathematical constraints, including width-to-height and height-to-width ratios and examples of compliant dimensions. For the submitter's intended advertising and bromide formats, an incorrect ratio rendered otherwise acceptable content unusable and required regeneration.

Support requested the environment, model, prompt, output dimensions, reference-image size, and onset date. The submitter supplied Chrome, GPT-5.6 Instant, the prompt, and an attached reference image. Support then treated the attached reference image as if it were the failed output and asked the submitter to measure it. The submitter corrected the misunderstanding and requested that OpenAI use its own generation history and asset metadata. The case was escalated, but the reviewed record contains no disclosed engineering finding, account-level comparison of requested and actual dimensions, or compensation decision.

Potential SEC relevance: repeated noncompliance with deterministic output requirements may increase compute waste, refunds, user friction, and the cost of serving professional workflows. It also bears on the reliability of product claims directed at creative and commercial users.

7.4 Missing progress/results and the reload-to-rate-limit feedback loop

In Cases 11268075 and 11269846, the submitter reported that image generation began without a visible progress state and sometimes completed without delivering a visible result. Because the interface did not indicate whether the job was processing, stopped, or finished, the submitter reloaded the page to determine the state. Repeated reloads then produced a "too many requests" restriction that could block ordinary ChatGPT use.

Support initially supplied general troubleshooting involving browser extensions, VPN, proxy, secure DNS, incognito mode, alternative browsers, waiting, stopping, regenerating, and hard reloading. The submitter objected that reload was already part of the reported causal chain and that regeneration could create a duplicate job if the server had completed the first request. One case was escalated; the other was closed as a duplicate within approximately one minute without a destination.

Potential SEC relevance: a UI state-delivery failure can amplify infrastructure load and customer restrictions through recovery behavior. SEC staff should ask whether OpenAI measured orphaned jobs, duplicate generations, reload-induced requests, rate-limit false positives, and the financial or retention effects of such incidents.

7.5 Officially recorded service outages and repeated image-generation incidents

OpenAI's public Status history recorded multiple relevant incidents in July 2026, including elevated image-generation errors on July 7, image-generation unavailability on July 21-22 and July 27, elevated file-upload and image-generation errors on July 22, and broader ChatGPT errors on July 19, July 23, and July 25-27. The submitter does not claim that every public incident affected his account. The history is relevant context for assessing whether the July 27 compensation claim concerned an isolated customer environment or a broader operational pattern.

The July 27 status entry stated that users were experiencing elevated errors and that OpenAI was implementing a mitigation; the incident was later marked fully recovered. The submitter's formal claim sought the account-specific incident period, affected components, root cause, failed-request logs, restoration, prorated refund, additional credit, and a process for documented business losses.

7.6 Pro plan "faster" and "unlimited" representations

In Case 10817141, the submitter asked why image-generation speed appeared similar after upgrading from Plus to Pro and requested an official explanation of "faster" and "unlimited." A human-signed response from "Razel" stated that "faster" was not a guarantee that every image would complete noticeably faster than Plus; normal-load performance could appear very similar. The response explained the benefit mainly as higher limits, priority during high demand, reduced delays during congestion, and improved responsiveness when capacity was limited.

The same response stated that “unlimited” or “virtually unlimited” meant no fixed daily or monthly quota for normal use, while remaining subject to fair-use protection, temporary rate limiting or throttling, anti-abuse and safety systems, maintenance, infrastructure events, and operational safeguards. The submitter retained a contemporaneous record that the Japanese pricing page described Pro image creation as “fast and unlimited” and notified OpenAI that he intended to seek consumer-protection review.

“Faster” refers to priority access and improved service availability, rather than a guaranteed reduction in the rendering time of every individual image.”

OpenAI Support, Case 10817141 (excerpt)

Potential SEC relevance: SEC staff should determine whether the draft S-1, pricing disclosures, revenue assumptions, churn analysis, risk factors, and due-diligence materials accurately describe the practical differences between tiers, dynamic limitations, fairness controls, and circumstances in which premium performance may resemble a lower-priced plan.

7.7 Opaque or conversationally supplied product limits

Case 11335132 contained a series of product-specification inquiries regarding Deep Research limits, memory capacity, project-file limits, aggregate context usage, GPT-5.6 model routing, reasoning-level rate limits, and macOS desktop-only features. AI-assisted support supplied numerical ranges and operational descriptions, including five-hour model windows and shared usage pools, while also stating that some limits were unpublished or content-dependent.

This record does not establish that each answer was wrong. It does show that a paying customer needed a live AI support thread to obtain material operational information that may affect product utility and plan value. SEC staff should determine whether customer-facing specifications were stable, substantiated, consistently disclosed, and reconciled with internal capacity, pricing, and revenue assumptions.

8.0 Detailed Factual Statement of Support and Case-Management Failures

8.1 Reclassification of the reported problem

Across the reviewed cases, the substance of complaints was repeatedly shifted into a different category: tool-control failure became image-generation guidance; context management became image generation; an officially acknowledged outage compensation claim became account-security speculation; and evidence preservation was answered with deletion instructions. This pattern matters because a complaint cannot reach the correct owner, log set, remediation process, or disclosure channel if the intake system changes the issue before investigation begins.

8.2 Duplicate closure without a traceable surviving record

Cases 10976796, 11269846, 11319313, 11319385, and 11319858 were closed as duplicates rapidly, in some instances within one or two minutes. The reviewed notices generally did not identify the surviving case, the issues transferred, the responsible team, or the next response date.

When the submitter challenged the duplicate closure, support stated that it could not see the merge destination from the thread and instructed the submitter to search prior email for case numbers. OpenAI later identified Case 10837305 as the destination, but the final support summary described only unintended image-generation invocation and omitted the separately stated context-management and current-edit-target issues.

“This thread cannot confirm the destination case number. Please search your inbox for prior emails containing “Case Number” and paste the original number.”

OpenAI AI-assisted Support, Case 11319521 (paraphrased from correspondence)

Potential SEC relevance: duplicate handling without issue-level lineage can corrupt complaint metrics, undercount distinct defects, obscure trends, frustrate auditability, and prevent material information from reaching management or disclosure counsel.

8.3 Repeated collection of information already in OpenAI's possession

Support repeatedly requested conversation URLs, timestamps, screenshots, screen recordings, output dimensions, prior case numbers, and merge destinations. Some customer-supplied information is reasonably necessary. The concern is that previously submitted material and OpenAI-controlled records appeared not to persist through handoffs, and OpenAI asked the customer to reconstruct information that its ticketing, asset, and tool-call systems would be expected to retain.

In Case 10337718, after OpenAI had acknowledged the specific issue and a human-signed response stated that it was recorded for review, support again requested a description, reproduction steps, screenshots, and a screen recording. This occurred while imposing a one-day response deadline.

8.4 Escalation language without a visible accountable process

Multiple cases generated the message that the matter had been escalated to a support specialist and that a response was expected in the next few days. The reviewed record often lacked the receiving department, responsible person, investigation start, logs reviewed, milestone, target date, or closure criteria. Case 10337718 remained without a disclosed substantive investigation result for approximately twenty days and later received an open-ended statement that no completion date or investigation timeline could be provided.

Potential SEC relevance: escalation metrics may look favorable if an automated system records “escalated” without testing whether a qualified owner accepted, investigated, and resolved the issue. SEC staff should examine the definitions and controls underlying customer-support KPIs supplied to management, boards, auditors, or underwriters.

8.5 Asymmetric deadlines and administrative closure pressure

The submitter documented that OpenAI could remain substantively nonresponsive for days or weeks, while the customer was told that the ticket would be closed if no reply arrived within one day. In Case 10337718, the customer had already stated that the matter remained unresolved and requested that the case remain open, yet support repeated the one-day deadline and re-requested previously supplied information.

Potential SEC relevance: a system that closes unresolved complaints for administrative nonresponse may reduce apparent backlog and time-to-resolution while leaving underlying defects unresolved. The SEC should determine how OpenAI calculates complaint closure, duplicate rates, resolution rates, aged cases, and customer satisfaction.

8.6 Evidence-preservation requests answered with deletion guidance

In Case 10837305, support discussed deleting conversations to remove generated images from the Images area while the submitter was requesting investigation and preservation of tool-execution evidence. In Case 11319858, a context-only investigation request received resources for data export, conversation/file deletion, and account deletion. In Case 12392937, a formal demand concerning an outage, logs, compensation, and possible account-misconduct insinuations again received export, deletion, and account-deletion resources.

The submitter does not claim that OpenAI expressly ordered him to destroy evidence. The narrower concern is that the support system failed to distinguish investigation and preservation from deletion and account exit. SEC staff should determine whether legal-hold or preservation terms trigger mandatory routing, whether deletion guidance is suppressed after a preservation request, and whether support automation can cause relevant records to expire before review.

8.7 Formal compensation claim answered with unsupported account-side possibilities

On July 28, under Case 12241050, the submitter formally sought an account-specific investigation of the publicly acknowledged image-generation outage, preservation of logs, restoration, prorated refund, additional service credit, compensation for substantiated direct business losses, and escalation to Billing, Incident Response, and Legal Review. The message expressly stated that generic device or browser troubleshooting was not responsive to a service-wide incident.

A response signed “John” did not acknowledge the public incident or address the requested remedies. It instead listed possible Terms of Service violations or harmful use, suspicious activity or unknown-location access, and account changes or maintenance as common reasons for limitations. The submitter then asked whether OpenAI was actually stating that his account had been flagged. He requested either the factual basis, date, action, restriction, and appeal process, or written confirmation that no such condition applied and that the response was a generic, inapplicable template.

“Potential Terms of Service Violation: Your account may be flagged for behavior that conflicts with our Terms of Service, such as unauthorized or harmful use of the platform.”

OpenAI Support response signed “John,” Case 12241050 (excerpt)

Potential SEC relevance: if the statement reflected an actual account flag, it raises questions about notice, evidence, restriction governance, and appeal. If it did not reflect an actual flag, it raises questions about sending unsupported misconduct implications to a paying customer. Either possibility bears on complaint controls, reputation, consumer protection, and automated decision governance.

8.8 Human-name and AI-response accountability ambiguity

At 1:42 p.m. on July 30, the submitter sent a detailed correction and ten numbered questions concerning John’s response, logs, preservation, incident facts, compensation, escalation, and case continuity. At 1:43 p.m., a new Case 12392937 was created and an AI assistant supplied only data-export, file/conversation-deletion, and account-deletion links. At 1:48 p.m., the submitter replied that an AI should not answer and that John should respond.

The record leaves unresolved whether John was a human who reviewed the claim, an AI-drafted response approved by a human, or a human name attached to an automated response. It also leaves unresolved who had authority to retract or correct the account-side insinuations. These questions are material to understanding the actual human control environment over customer communications, legal claims, and account enforcement.

8.9 No visible customer-remedy framework

The reviewed record does not provide a published or case-specific standard for prorated refunds, service credits, restoration of usage capacity, compensation for direct business losses, required evidence, responsible department, reserve methodology, or response deadline. The submitter’s individual claim remains a factual example; SEC staff should determine whether the absence or limited visibility of a remedy framework is broader and whether refund or credit exposure is measured and disclosed.

9.0 Cross-Case Control Failures

Control domain	Observed external indicator	Potential consequence for SEC review
Issue taxonomy	Distinct product, interface, pricing, context, outage, and legal-remedy issues were repeatedly collapsed into existing categories.	Misclassified risks may not reach engineering, legal, finance, or disclosure owners.
Case lineage	Duplicate closure did not reliably identify the surviving case or retained issues.	Complaint counts and trends may be incomplete or unauditable.

Control domain	Observed external indicator	Potential consequence for SEC review
Handoff continuity	Previously supplied facts were requested again and issue summaries changed across cases.	Material information may be lost between automated and human stages.
Escalation control	"Escalated" was communicated without owner, acceptance, milestone, or completion date.	Escalation KPIs may not represent substantive investigation.
Human oversight	AI-assisted responses handled serious claims, deletion guidance, and account-side implications.	The scope of automated authority and mandatory human review is unclear.
Evidence preservation	Deletion resources appeared after investigation and preservation requests.	Potential legal-hold, retention, and audit-trail risks.
Customer-remedy accounting	No clear refund, credit, or business-loss framework appeared in the record.	Potential unmeasured liability, churn, and revenue-quality implications.
Marketing substantiation	Support qualified "faster" and "unlimited" more narrowly than the headline language alleged by the submitter.	Potential consumer, regulatory, and offering-disclosure risk.
Operational resilience	The user-specific record overlaps with repeated public incidents.	Potential capacity, reliability, and business-continuity risk.
Board and disclosure flow	No evidence shows how complaints or incident patterns reach senior management or the boards.	Potential disclosure-control and oversight gap; SEC should obtain internal evidence.

10.0 Potential Disclosure Relevance for SEC Staff Review

NO LEGAL CONCLUSION The following sections identify possible disclosure channels for SEC inquiry. They do not state that OpenAI has omitted a material fact, because the draft S-1 and internal evidence are nonpublic.

10.1 Item 101 - Description of the business and material regulation

SEC staff should determine whether the business description adequately explains dependence on reliable model routing, image generation, subscription-tier differentiation, automated support, account enforcement, customer data and logs, and the regulatory environment governing consumer representations, automated decisions, privacy, and paid digital services. If these functions are material to revenue or growth, the description should not present them only as product capabilities while omitting operational constraints necessary to understand the business.

10.2 Item 105 - Tailored risk factors

Potential tailored risks include: model and tool-routing regressions; failure to follow deterministic specifications; service outages and capacity constraints; rate-limit and entitlement errors; premium-tier differentiation; inadequate customer support at scale; automated misclassification or closure of legal and technical complaints; loss of complaint lineage; inability to preserve or retrieve account records; adverse account decisions; refunds, credits, churn, and reputation; consumer-protection investigations; and board oversight of AI-mediated customer communications.

Generic language that "AI may make mistakes" may not fully describe a system in which the same or similar failure appears in the support mechanism charged with investigating the product. SEC staff should assess whether the risk factors are specific enough to explain how these risks can affect revenue, costs, customer trust, professional adoption, regulatory exposure, and growth assumptions.

10.3 Item 303 - Known trends, events, and uncertainties

If management has data showing repeated outages, escalating support volume, duplicate-closure patterns, credits or refunds, premium-tier dissatisfaction, professional-workflow disruption, or increased human-review costs, SEC staff should determine whether these are known trends or uncertainties reasonably likely to affect net sales, subscription retention, operating expenses, infrastructure costs, legal costs, or future financial condition.

10.4 Rule 408 - Additional information necessary to avoid misleading required statements

If the draft S-1 emphasizes product reliability, user growth, enterprise readiness, safety, trust, subscription value, or scalable AI-assisted operations, SEC staff should determine whether additional context about known reliability, capacity, support, and customer-redress limitations is necessary to prevent those statements from being misleading in light of the circumstances.

10.5 Item 103 and contingent legal or regulatory exposure

The record includes the submitter's notice of potential consumer-protection review and a formal compensation claim. One customer claim is not necessarily material. SEC staff should obtain the aggregate number, value, and status of similar refund, credit, marketing, outage, account-enforcement, or business-loss claims and determine whether any proceedings, investigations, contingencies, or reasonably possible losses require disclosure.

10.6 Disclosure controls and procedures

The external evidence raises a foundational question: can information about distinct defects remain intact from customer intake through triage, duplicate consolidation, engineering escalation, legal review, management reporting, board oversight, and securities disclosure? SEC staff should examine whether OpenAI has controls designed to record, process, summarize, and communicate potentially material product and customer information to the officers responsible for disclosure decisions.

10.7 Governance and public-benefit structure

OpenAI states that the Foundation controls OpenAI Group PBC and appoints its board. SEC staff should determine which board or committee oversees consumer product reliability, support automation, account enforcement, incident response, legal holds, marketing substantiation, compensation policies, and disclosure controls; what metrics are presented; and whether the dual mission and commercial structure alter accountability for these risks.

10.8 Revenue quality, customer retention, and remediation costs

Premium subscription revenue may be affected by the practical value of paid-tier differentiation, outages, usage restrictions, failure to restore consumed capacity, support quality, and customer remedies. SEC staff should review churn, downgrade, refund, chargeback, credit, complaint, and service-incident data by plan and geography, together with management's revenue and retention assumptions.

11.0 Questions and Requested Investigative Steps

1. Compare this chronology and the underlying exhibits with the operational, customer, risk-factor, MD&A, governance, legal, and regulatory disclosures in OpenAI's confidential draft S-1 and subsequent amendments.
2. Ask OpenAI to identify all internal incidents, bug reports, regressions, and customer complaints corresponding to unintended tool invocation, context contamination, aspect-ratio noncompliance, missing generation results, reload-related rate restrictions, and premium-plan entitlement concerns from at least April through July 2026.
3. Obtain complaint volumes, duplicate-closure rates, merge-error rates, aged-case reports, reopen rates, escalation acceptance rates, substantive resolution times, customer satisfaction, and the definitions used for "closed," "resolved," "duplicate," and "escalated."

4. Determine whether OpenAI's support automation can close or merge cases, make or imply account-enforcement determinations, route compensation or preservation requests, and send account-deletion instructions without human approval.
5. Identify the person, system, and approval path behind the response signed "John" in Case 12241050 and the AI-assisted response in Case 12392937.
6. Determine whether the submitter's account was actually flagged for Terms of Service violation, harmful use, suspicious access, security concern, or temporary downgrade; if not, determine why those possibilities were sent in response to a public outage claim.
7. Confirm preservation of all relevant support, tool-call, generation, asset, routing, rate-limit, account-entitlement, and incident records associated with the listed cases and dates.
8. Review whether OpenAI has a legal-hold trigger for preservation language and whether deletion guidance is blocked or human-reviewed after a preservation request.
9. Obtain the internal meaning, substantiation, and customer-facing placement of "faster," "unlimited," "virtually unlimited," priority access, fair use, throttling, and safety or anti-abuse limitations for ChatGPT Pro.
10. Review pricing-page versions, localization records, A/B tests, conversion studies, complaint data, and legal/compliance approvals concerning premium-tier marketing in Japan and other markets.
11. Obtain refund, service-credit, usage-restoration, chargeback, and business-loss claim policies, decision criteria, aggregate amounts, reserves, and historical application after outages.
12. Ask whether customer complaint and incident trends were communicated to senior management, OpenAI Group's board, the OpenAI Foundation board, the Safety and Security Committee, auditors, underwriters, and disclosure counsel before and after the confidential S-1 submission.
13. Review whether internal operational and customer data contradict, qualify, or require expansion of any draft S-1 statements concerning product reliability, scalability, safety, trust, premium value, customer growth, professional use, support efficiency, regulatory exposure, or governance.
14. If material information is missing or described generically, require appropriate amendment or tailored disclosure before effectiveness.
15. Where the evidence suggests non-securities consumer-protection, privacy, unfair-practice, or account-enforcement issues, consider referral to the Federal Trade Commission, state authorities, or relevant foreign regulators as appropriate.

12.0 Evidence-Preservation Request

The submitter asks that the SEC promptly request or otherwise ensure preservation of records relevant to the matters described in this document. The list below is intentionally broader than the records visible to an external customer because the central control questions cannot be answered from email alone.

- All messages, thread metadata, case histories, routing events, merge mappings, duplicate-classification reasons, summaries, tags, disposition codes, reopen events, escalation events, and agent or reviewer identifiers for the listed case numbers.
- All prompts, assistant outputs, tool-call records, image-generation requests, generation assets, errors, output dimensions, safety decisions, model-routing decisions, timestamps, and rate-limit events associated with the submitter's account during the relevant periods.
- Front-end telemetry concerning missing progress indicators, missing result delivery, page reloads, duplicate jobs, and "too many requests" restrictions.
- Incident-management records, root-cause analyses, post-incident reviews, capacity metrics, user-impact estimates, credits, refunds, and internal communications for the relevant public Status incidents.
- Versions, translations, approval records, substantiation analyses, and performance data for Pro pricing and usage representations, including "faster" and "unlimited."
- Policies and model instructions governing AI-assisted support, human escalation, legal or compensation claims, evidence-preservation requests, account-security implications, deletion guidance, and customer deadlines.

- Management, board, committee, auditor, underwriter, and disclosure-counsel reports or presentations containing complaint, reliability, outage, refund, churn, marketing, or automated-support data relevant to the draft S-1.

PRESERVATION CAUTION The submitter has retained correspondence and screenshots available to him. OpenAI controls the authoritative server-side records needed to determine whether the complaints were correctly logged, merged, investigated, summarized, escalated, and reflected in disclosure processes.

13.0 Exhibit Index and Submission Package Checklist

This Word document is designed as the narrative attachment. The strongest TCR package will upload the underlying records as separate, clearly named exhibits. The list below distinguishes evidence already retained from evidence that should be requested from OpenAI.

Exhibit	Description	Status / action
	Comprehensive chronology article: "OpenAI Support Functional Failure," covering the first eight core cases and common processing structure.	Retained by submitter; attach.
	Case 10337718 full email thread: unintended image generation, potential usage impact, escalation, delay, repeated evidence requests, and one-day closure deadline.	Retained by submitter; attach as searchable PDF.
	Case 10817141 full email thread: Pro "faster" and "unlimited" explanation and consumer-protection notice.	Retained by submitter; attach.
	Case 10837305 full email thread: tool routing, context contamination, timestamp requests, deletion guidance, and later merge destination.	Retained by submitter; attach.
	Case 10976796 submission and duplicate-closure notice.	Retained by submitter; attach.
	Cases 11268075 and 11269846: missing image results, reload-to-rate-limit issue, troubleshooting, escalation, and duplicate closure.	Retained by submitter; attach.
	Case 11271585: 4:5 aspect-ratio report, full prompt, reference image, support misunderstanding, correction, and escalation.	Retained by submitter; attach.
	Cases 11319313, 11319385, and 11319521: context complaint, duplicate closures, unknown merge destination, and incomplete final summary.	Retained by submitter; attach.
	Case 11319858: context-only investigation request followed by export, deletion, and account-deletion instructions.	Retained by submitter; attach.
	Case 11335132: AI-assisted product specification and usage-limit thread.	Retained by submitter; attach as supplementary evidence.
0	Case 12241050: formal image-outage compensation and preservation claim; response signed "John."	Retained by submitter; attach.
1	Case 12392937: AI-assisted export/deletion/account-deletion response and demand for John to answer.	Retained by submitter; attach.
	Contemporaneous Japanese ChatGPT pricing-page capture showing the Pro image-generation wording at issue.	Retained screenshot should be attached with capture date and full-page context.
	Proof of ChatGPT Pro subscription and amount paid.	Attach invoice or subscription receipt with unnecessary payment details redacted.
	OpenAI public Status pages for July 7, July 21-22, July 22, July 27, and related ChatGPT incidents.	Attach saved PDFs or screenshots; public links are in Appendix C.

ibit	Description	Status / action
	Representative screenshots and conversation links for unintended tool execution, context contamination, aspect-ratio failures, missing progress/results, and rate restrictions.	Attach only representative, date-labeled examples.
	Server-side case mappings, logs, internal summaries, incident reviews, support-agent identities, and disclosure-routing records.	Not available to submitter; request from OpenAI.

13.1 File-naming convention

Use searchable PDFs with names such as “Exhibit_B-01_Case_10337718.pdf.” Preserve the full message headers, dates, sender, recipient, subject, case number, and quoted thread history. Avoid uploading only cropped screenshots when a searchable email export is available.

14.0 Certification

I submit this information in good faith. To the best of my knowledge and belief, statements identified as first-hand observations accurately describe what I observed; quoted or summarized OpenAI communications are based on records I retained; public-source statements are identified as such; and legal or disclosure characterizations are presented as requests for SEC investigation rather than conclusive findings.

I understand that knowingly and willfully providing false, fictitious, or fraudulent information to the SEC may carry legal consequences. I am prepared to provide the underlying records and to clarify the chronology if contacted by SEC staff.

NAME

Shohei Kimura

SIGNATURE

DATE

CONTACT

ak4dy1@gmail.com | +81-70-3666-0022

Appendix A Public OpenAI Service-Incident Context

This table is supplementary context. It does not establish that each public incident affected the submitter. Times and labels follow OpenAI’s public Status pages as displayed.

Date	OpenAI Status entry	Relevance
Jun. 3, 2026	Elevated errors for ChatGPT Pro	Public status record concerning the premium plan.
Jul. 7, 2026	Elevated errors with image generation in ChatGPT	OpenAI reported that some users could not generate images; mitigation and recovery followed.
Jul. 9, 2026	Elevated errors when selecting models	Relevant to model availability and routing.
Jul. 12, 2026	Conversation and login errors on ChatGPT for iOS and macOS	Relevant to platform reliability across desktop/mobile clients.
Jul. 15, 2026	Elevated errors affecting ChatGPT; separate voice issue	Broad service availability context.
Jul. 19, 2026	Elevated errors affecting ChatGPT conversations, Voice, and Work Mode	OpenAI identified the cause and applied mitigation.

Date	OpenAI Status entry	Relevance
Jul. 21-22, 2026	Image generation unavailable / elevated API image errors	Multiple image-generation disruptions across ChatGPT and APIs.
Jul. 22, 2026	Elevated errors affecting file uploads and image generation	Relevant to multimodal and project workflows.
Jul. 23, 2026	Elevated error rates affecting multiple ChatGPT components	Status page listed fifteen affected ChatGPT components.
Jul. 24, 2026	Elevated API error rates and latency on an image model	Additional image-system reliability context.
Jul. 25-27, 2026	Elevated errors affecting ChatGPT conversations	Intermittent failures prevented some users from loading or continuing conversations.
Jul. 27, 2026	Image generation unavailable in ChatGPT	The public incident underlying the submitter's formal compensation claim.

Appendix B Key Correspondence Excerpts

These short excerpts are included for orientation. The full email threads should be attached as exhibits so SEC staff can review complete context.

Case 10337718 - issue recognized but timeline unavailable

OpenAI acknowledged concern that image generation may have begun when only text output was intended, stated that the report remained under internal review, and said it could not provide a specific completion date or investigation timeline.

Case 10337718 - one-day customer deadline

After extended delay, OpenAI asked again for an issue description, reproduction steps, screenshot, and screen recording, and said the ticket would be closed if the customer did not respond within one day.

Case 10817141 - premium-tier qualification

OpenAI Support stated that Pro “faster” did not guarantee a faster completion time for every image, that Plus and Pro could appear similar under normal load, and that “unlimited” remained subject to fair-use and operational safeguards.

Case 11319521 - merge destination unavailable to the support thread

AI-assisted support stated that it could not see the merge destination and told the customer to search past OpenAI emails for a case number. It later identified Case 10837305.

Case 11319521 - issue summary changed

The final support summary described unintended image-generation invocation but omitted the separate context-management, current-edit-target, and workflow-rule issues that the customer had expressly asked OpenAI to preserve.

Case 12241050 - customer-side possibilities introduced

The response signed “John” listed possible Terms of Service violation or harmful use, suspicious activity or unknown-location access, and account or payment changes, without stating whether any applied to the customer or addressing the acknowledged outage and compensation requests.

Case 12392937 - deletion response

The AI assistant supplied resources for data export, deletion of uploaded files or conversations, and deletion of the account. It did not address the numbered correction questions, preservation request, refund, credit, business-loss process, or responsible team.

Appendix C Authorities and Public Sources

Public sources were checked on July 30, 2026. Links may later change; saved PDF or screenshot copies should be retained with the submission package.

- [1] OpenAI. Confidential submission of draft S-1 to the SEC. June 8, 2026. <https://openai.com/index/openai-submits-confidential-s-1/>
- [2] OpenAI. Our structure. Current page reviewed July 30, 2026. <https://openai.com/our-structure/>
- [3] U.S. Securities and Exchange Commission. Submit a Tip. Updated Nov. 18, 2024. <https://www.sec.gov/submit-tip-or-complaint/submit-a-tip>
- [4] U.S. Securities and Exchange Commission. Filing Guidance and Confidentiality. Updated June 5, 2025. <https://www.sec.gov/submit-tip-or-complaint/tips-complaints-resources/filing-guidance-confidentiality>
- [5] U.S. Securities and Exchange Commission. Report Suspected Securities Fraud or Wrongdoing. Updated Nov. 18, 2024. <https://www.sec.gov/submit-tip-or-complaint/tips-complaints-resources/report-suspected-securities-fraud-or-wrongdoing>
- [6] U.S. Securities and Exchange Commission. Voluntary Submission of Draft Registration Statements - FAQs. July 1, 2024. <https://www.sec.gov/about/divisions-offices/division-corporation-finance/voluntary-submission-draft-registration-statements-faqs>
- [7] U.S. Securities and Exchange Commission. Modernization of Regulation S-K Items 101, 103, and 105. Nov. 6, 2020. <https://www.sec.gov/resources-small-businesses/small-business-compliance-guides/modernization-regulation-s-k-items-101-103-105-small-entity-compliance-guide>
- [8] U.S. Securities and Exchange Commission. Commission Guidance Regarding Management's Discussion and Analysis. Dec. 29, 2003. <https://www.sec.gov/rules-regulations/2003/12/commission-guidance-regarding-managements-discussion-analysis-financial-condition-results-operations>
- [9] Electronic Code of Federal Regulations via Cornell LII. 17 CFR 230.408 - Additional information. Current text reviewed July 30, 2026. <https://www.law.cornell.edu/cfr/text/17/230.408>
- [10] Electronic Code of Federal Regulations via Cornell LII. 17 CFR 229.105 - Risk factors. Current text reviewed July 30, 2026. <https://www.law.cornell.edu/cfr/text/17/229.105>
- [11] Electronic Code of Federal Regulations. 17 CFR 229.303 - MD&A. Current text reviewed July 30, 2026. <https://www.ecfr.gov/current/title-17/chapter-II/part-229/subpart-229.300/section-229.303>
- [12] OpenAI Status. Elevated errors with image generation in ChatGPT. July 7, 2026. <https://status.openai.com/incidents/ssk9hg56>
- [13] OpenAI Status. Image generation unavailable in ChatGPT. July 27, 2026. <https://status.openai.com/incidents/60ferbhc>
- [14] OpenAI Status. OpenAI Status History. Reviewed July 30, 2026. <https://status.openai.com/history>

END OF DRAFT SEC TCR ATTACHMENT