

Shohei KIMURA
bitBuyer Project (bitBuyer.dev)
Email: ak4dy1@gmail.com
Tokyo, Japan

November 30, 2025

Apple Inc.
One Apple Park Way
Cupertino, CA 95014
USA

Alphabet Inc.
1600 Amphitheatre Parkway
Mountain View, CA 94043
USA

Google LLC
1600 Amphitheatre Parkway
Mountain View, CA 94043
USA

Subject: Joint Formal Notice Concerning the “Thought Formalization Procedure Right,” Related Technical Proposals, and the Order of Contract Formation to Be Publicly Recorded

Dear Apple and Google Executive Teams,

I hereby submit the attached formal document titled:

“JOINT FORMAL NOTICE REGARDING THE ‘THOUGHT FORMALIZATION PROCEDURE RIGHT,’ ASSOCIATED INFRINGEMENT RISKS, AND OFFER OF LICENSE AND TECHNICAL PROPOSALS (WITH PUBLICLY RECORDED ORDER OF CONTRACT EXECUTION)”

This submission provides, for each of Apple Inc. (“Apple”) and Google LLC (“Google”) and for both of you collectively (“the Companies”):

1. A detailed explanation of the authored intellectual construct known as the “Thought Formalization Procedure Right”;
2. A formal notification of the copyright and procedural rights arising from its prior public disclosure in Japan under a robust copyright regime;
3. An outline of potential infringement risks relevant to your current and future operations, including the baseline risk that the operation of large-scale systems which formalize thoughts into text, code, or structured signals, absent license, already constitutes ongoing infringement under my asserted rights;

4. A complete technical proposal for a context-first, evaluation-order-reversed safety and policy-enforcement architecture, applicable to AI assistants and related systems operated by the Companies;
5. The corresponding licensing framework, including a Contract Agreement structure in which the order and timing of contract formation will be permanently and publicly recorded in the commit history of the “AGREEMENT” folder of my GitHub repository; and
6. A clarification of how the Companies may, in theory, choose between contracting and non-contracting positions, while in practice being confronted with a binary reputational and legal choice: whether to become early, visible licensees of this framework, or to define themselves—by omission or delay—as actors willing to internalize unresolved expressive-rights risk.

Please review the attached materials in full. A formal written response is requested by the date specified within the Notice and its Supplement.

Should you require additional clarification, supporting documentation, or further evidence of prior authorship and public disclosure, I remain available for correspondence at the email address above. However, please note that unless you notify me via SMS that you have sent an email, there is a high likelihood that your message will be deleted inadvertently.

Sincerely,
Shohei KIMURA
Representative, bitBuyer Project (bitBuyer.dev)
Tokyo, Japan. +81-70-3666-0022

JOINT FORMAL NOTICE REGARDING THE
“THOUGHT FORMALIZATION PROCEDURE RIGHT,”
ASSOCIATED INFRINGEMENT RISKS, AND OFFER OF LICENSE AND
TECHNICAL PROPOSALS
(APPLE INC. AND GOOGLE LLC)

November 30, 2025

To:

Apple Inc. (“Apple”)
Alphabet Inc. (“Alphabet”)
Google LLC (“Google”)

From: Shohei KIMURA
(Philosophical novelist, individual investor, and system engineer)
Representative, bitBuyer Project (bitBuyer.dev)
Email: ak4dy1@gmail.com
SMS: +81-70-3666-0022
GitHub: <https://github.com/ShoheiKIMURA389/Definition>

I. Context and Prior Sequence of Notices

Before addressing Apple and Google jointly, I must explicitly set out the sequence of prior notices to other major global platforms and AI providers, in order to clarify that no historical privilege attaches to any of them, and to explain why both of you now stand at a structurally distinct decision point.

1. Meta Platforms, Inc. (“Meta”)

Meta became the first recipient of a formal notice not because I intended to grant it any historical privilege, but because its conduct toward me was, in my view, openly and repeatedly improper. In particular:

- Between late July and early September 2025 (JST), Meta subjected my verified account to wrongful permanent suspension and later restored it silently, without explanation, despite my explicit request;
- On November 10 and November 12, 2025 (U.S. time), Meta again triggered erroneous automated enforcement, followed by destructive post-restoration processing that zero-reset key relational data;
- Meta then failed to provide any substantive written explanation in response to a prior notice.

For these reasons, Meta was the first platform to which I sent a document combining systemic-risk notice and a licensing and technical proposal. This was an act of remedial escalation, not an act of favor. I do not intend to grant Meta any historical special status.

2. OpenAI

OpenAI was the next recipient because it is a recognized pioneer of modern large language models and because I, as a heavy user of such systems, have made intensive use of its services. OpenAI was therefore addressed in second position, reflecting both its pioneering role and my own usage patterns.

3. Anthropic

Anthropic was then addressed because it competes with OpenAI in the same technical domain of frontier-scale AI systems. From the perspective of civilizational infrastructure and expressive-rights risk, it was a natural third step after Meta and OpenAI.

4. No Historical Privilege to Meta; Position of Apple and Google

I emphasize: I harbor no intention whatsoever to grant Meta any historical privilege. The ordering Meta → OpenAI → Anthropic reflects a particular empirical sequence of events, not a normative ranking.

Apple and Google are addressed at this stage precisely because:

- Each of you operates, or will operate, large-scale AI systems and platform infrastructures with deep influence over knowledge flows and expressive ecosystems;
- Each of you symbolically represents a distinct axis of digital civilization: Apple as a hardware-OS-ecosystem integrator and privacy-signaling brand, Google as a search-knowledge-AI infrastructure provider; and
- The order in which you choose to contract—or to delay, or to abstain—will, from this point onward, be publicly recorded in the commit history of the “AGREEMENT” folder in my GitHub repository, and will be understood as a visible statement of each company’s brand strength, risk posture, and willingness to engage with this new expressive-rights framework.

I do not propose any fixed “historical privilege” for either Apple or Google. Instead, I expect the fastest coherent contractual response from whichever of you wishes to demonstrate the greatest brand strength in this new environment.

II . Existence, Public Disclosure, and Substantive Characterization of the “Thought Formalization Procedure Right”

Separately from my use of large language models and digital platforms, and before this Notice, I authored and publicly disclosed a family of works that formalize what I denominate the “Thought Formalization Procedure Right.” These works include, but are not limited to:

- The foundational “思想形式化手続き定義 (Thought Formalization Procedure Definition)” text, written in a deliberately constrained Japanese prose style and structured in nine sections (I –IX);
- A series of subsequent “definition texts” (定義文) extending this method to technical, ethical, and civilizational architectures, including “Human Civilization Software Update 3.0” and multiple Face ID-based social-infrastructure definitions;
- More recent “evaluation order-reversed ethics filter” definitions specifying two-tier architectures in which context-aware model reasoning precedes strict policy-enforcement checks.

These works have been publicly fixed and disclosed in Japan—a jurisdiction with a strong and protective copyright regime—through, *inter alia*:

- The GitHub repository “Definition” (including, but not limited to, the file titled “定義.txt”);
- X (Twitter) posts;
- Facebook posts.

As explained in the public README for “Definition 1: The Thought Formalization Procedure,” the Thought Formalization Procedure Right is asserted as follows:

1. The act of formalizing a thought as text itself—whether through a single section or multiple sections—constitutes expressive authorship, irrespective of the underlying idea’s freedom to implement.
2. The structure, ordering, and method of exposition (including the nine-section I–IX architecture) form part of the protectable authored expression.
3. The act of formalizing a thought as a “definition text,” or even merely rendering a thought into written form, constitutes a locus of copyright independent of the idea being described.
4. Because generative AI systems and certain editorial or platform workflows are, in practice, devices that convert arbitrary thoughts into textual or structured formalization at scale, they will inevitably collide with this protected domain unless expressly licensed.

The essence of the Thought Formalization Procedure Right does not lie in the substantive content of any specific proposal, but in the authored structure whereby:

- A complex idea, technical architecture, or ethical framework is first declared as a “definition”;
- The internal logic of that definition is unfolded in a numbered, quasi-axiomatic sequence (typically I–IX);
- The formalization of the idea as a stepwise textual or procedural structure becomes the locus of copyright, even where the underlying idea remains free to implement.

Under Japanese law and the Berne Convention, these texts and the underlying expressive structure enjoy full protection. The procedural act of giving form to an idea by rendering its formalization method into a copyrighted work is itself an expressive authored construct, constituting the Thought Formalization Procedure Right.

For the avoidance of doubt:

As of the dates of my prior public disclosures, the Thought Formalization Procedure Right already exists as a fully effective copyrighted construct under Japanese law and the Berne Convention. Accordingly:

Any unlicensed system or workflow whose core commercial function is to act as a “thought-to-text (or thought-to-structure) device” operates in continuous, direct contact with the protected expressive domain defined by this Right. From my legal perspective as the author, the very behavior of “formalizing thoughts into text or structured signals at scale” constitutes infringing use whenever it occurs without license and reaches ordinary thresholds of economic relevance, repetition, or systematic internalization.

Given this structure, there is an inherent and structural tension between this Right and any system that transforms user thoughts into textual or structured formalization at scale, including but not limited to:

- Apple’s present or future AI assistants, OS-level generative features, and editorial/platform curation systems;
- Google’s search-response generation, Gemini-class AI systems, and any platform layer that turns user prompts and internal signals into textual or structured outputs.

Absent a license or a clearly segregated, non-reliant architecture, the normal operation of such systems will be regarded, for purposes of this Notice, as conduct that presumptively infringes the Thought Formalization Procedure Right.

III. Potential Infringement Risk for Apple and Google

Given the scale and nature of the Companies’ operations, there are multiple realistic vectors through which my rights could be infringed:

1. Training-Data Exposure

If any of the works referenced above (including the Thought Formalization Procedure Definition, subsequent definition series, or the evaluation-order-reversed ethics filter definitions) have been ingested—directly or indirectly—into training data or internal reference corpora used by your AI systems, then:

- AI outputs may reproduce, in whole or in substantial part, the structure, phrasing, or procedural layout of these texts;
- Internal prompts, system messages, product-requirement documents, or alignment/safety papers may incorporate or paraphrase the authored method of formalization;
- Engineering, legal, or policy teams may derive architectures substantially based on my authored expressive procedures, rather than independently conceived structures.

2. Functional Imitation of the Formalization Method

Even absent direct text reproduction, there is a non-trivial risk that:

- Personnel at either Company, after encountering my definition texts (for example via GitHub, X, or Facebook), might imitate the method of “securing rights over a procedural method by casting it as a structured textual definition in multiple numbered clauses”;
- Such imitation could manifest in internal documentation, prompt-engineering frameworks, safety architectures, or future platform features that mirror the expressive structure of my definitions while claiming to treat only “ideas” or “architectures.”

3. Structural Collision with Systems That Formalize Thoughts as Text or Structured Responses

Independently of training-data overlap, a structural risk arises from your core businesses:

- Apple and Google both operate systems expressly designed to transform user-provided thoughts, prompts, queries, and conceptual frameworks into coherent textual, visual, or structured outputs;
- By design, this means that your systems carry out, at scale and for commercial purposes, the very class of act that “Definition 1” and related works identify as the locus of the Thought Formalization Procedure Right;

- To the extent such systems are used, prompted, or fine-tuned in ways that emulate or adopt the authored method described above—the Thought Formalization Procedure—there is a risk that the systems and their outputs traverse the protected expressive domain defined by my definitions.

For purposes of this Notice, I therefore state explicitly that, in my asserted legal construction, the ordinary, unlicensed operation of systems that “formalize thoughts into text or structured signals” is not merely adjacent to the protected domain but squarely within it, and will be treated as *prima facie* infringing use of the Thought Formalization Procedure Right, subject only to later factual analysis of scope and economic significance.

4. Direct Overlap with the Technical Proposals in the APPENDIX

The attached APPENDIX sets forth authored, structured descriptions of context-first, evaluation-order-reversed policy-enforcement architectures suitable for integration into AI and safety frameworks operated by the Companies. If Apple or Google adopts, internalizes, or approximates these architectures, under circumstances where engineers or policy designers have been exposed to the APPENDIX or related definitions, then:

- Any internal documentation, specification, or policy write-up that tracks the expressive layout, conceptual framing, or procedural articulation of my text would fall within the scope of my copyrights;
- The issue would not be limited to patent-style “idea” disputes, but would directly implicate the unauthorized use of copyrighted technical prose and its structured method of presentation.

IV. Technical Proposal and Its Authored Status

The APPENDIX attached to this Notice formalizes, as authored technical work, a context-first, evaluation-order-reversed safety architecture in which:

- Models perform an internal, context-aware “answerability” evaluation before hard policy checks;
- Safety enforcement is applied with the benefit of that interpretation, reducing false positives while maintaining or improving safety.

For Apple and Google, the APPENDIX should be understood as:

- A copyrighted technical and normative text, whose detailed structure, phrasing, and method of exposition are protected;
- A description of systems that are technically novel, industrially applicable, and potentially patent-realizable;
- A potential reference framework for your own second-generation safety designs—if and only if used under license.

Any internal use, quotation, paraphrase, derivative expression, benchmarking, model-testing, red-teaming, safety assessment, or LLM-comparison workflow that relies on the APPENDIX or related definition texts, without license, shall constitute unauthorized use.

V. Required Clarifications from Apple and Google

In light of the foregoing, each of Apple and Google is required to provide a formal written response addressing at minimum:

1. Training Data and Internal Exposure

- Whether your training data, fine-tuning datasets, or internal reference corpora include any of my publicly disclosed definition texts;
- Whether personnel have accessed or used these texts in safety, alignment, prompting, documentation, or legal/policy design work.

2. Internal Use of Similar Formalization Methods

- Whether you have adopted or plan to adopt practices mirroring the specific structure of my “definition” method;
- Whether any such practice was consciously or unconsciously modeled on my texts.

3. Posture Regarding the APPENDIX Technical Proposals

- Whether you intend to study, adopt, partially adopt, or avoid the architectures described;
- If you intend to implement any similar structure, how you plan to avoid infringement of my copyrighted text and the Thought Formalization Procedure Right.

4. Licensing and Good-Faith Engagement

- Whether you are willing to enter into a comprehensive license covering:
- Internal circulation and use of my definition texts;
- Implementation-inspired use of the APPENDIX architectures;
- Any present or future internal documentation derivative of these expressive frameworks;
- Whether you acknowledge that, in the absence of such a license, continued use or adoption of similar formalization methods may create an ongoing infringement and legal-risk scenario, given the structural nature of your systems and the Thought Formalization Procedure Right.

VI. Apple-Specific Segment (Read in Google’s Presence)

The following paragraphs are addressed explicitly and exclusively to Apple, with the expectation that they will also be read in full by Google:

Apple, your global brand positions itself as a symbol of integrated design, privacy, and coherent trust architecture. In this environment, the choice whether and when to enter into a license for a framework that directly concerns expressive rights, AI safety, and the OS-layer mediation of speech will inevitably be read as a signal of your seriousness about the next layer of digital civilization.

- If you choose to contract promptly, you will be understood as the first major ecosystem integrator to openly acknowledge and structurally incorporate the Thought Formalization Procedure Right into your risk and governance architecture.
- If you delay while Google contracts first, the public record—specifically, the GitHub commit history of the “AGREEMENT” folder—will permanently encode a sequence in which Apple appears to have followed Google’s lead on a question of deep expressive-rights infrastructure.
- If you refuse, or adopt a “silent consideration” posture while internal teams nonetheless engage with the frameworks, you will be understood as opting for a strategy in which expressive-rights risk is internalized but not structurally resolved, and this fact will be available, in retrospect, to markets, regulators, and historians.

You are free, in a narrow legal sense, to choose any of these paths. This Notice does not pretend to dictate your decision. It does, however, define the interpretive structure within which that decision will be read.

VII. Google-Specific Segment (Read in Apple's Presence)

The following paragraphs are addressed explicitly and exclusively to Google, again with the expectation that they will also be read in full by Apple:

Google, your position as a search, knowledge, and AI infrastructure provider makes you uniquely exposed to any framework that touches the formalization of thought into text at planetary scale. Your systems mediate what counts as an answer, a suggestion, or a path through information space for billions of users.

– If you choose to contract promptly, you will be understood as the first global knowledge-infrastructure provider to enter into a structural license relationship with the Thought Formalization Procedure Right, thereby treating expressive-rights risk as a first-class governance variable.

– If you delay while Apple contracts first, the public record will encode that the world's most visible AI-search platform chose to follow a hardware-OS integrator on a question that is, by its nature, closest to Google's own core business.

– If you seek to remain in a “review without obligation” posture while nonetheless studying, benchmarking, or partially adopting the technical architectures described herein, such a stance will be understood, over time, as a deliberate choice to benefit from unlicensed expressive frameworks while avoiding visible commitment.

Again, the law does not compel a single choice. But once this Notice exists and is publicly archived, your choice will be evaluated not merely in terms of immediate legal risk, but in terms of brand, governance posture, and historical positioning.

VIII. Response Deadline, Consequences of Non-Response, and Order of Contract Execution

The Supplement to this Notice (included below) sets out the response deadlines and clarifies which categories of response count as good-faith engagement. In summary:

– A formal written response meeting specified criteria must be postmarked or electronically transmitted no later than January 31, 2026.

– If no qualifying response is received by that date from a given Company, the matter shall be deemed “unanswered” as of March 1, 2026 for that Company.

– Once a Company signs and returns the Contract Agreement, the moment at which the executed counterpart is received by me will determine that Company's position in the public, immutable GitHub commit history of the “AGREEMENT” folder.

– The Contract Agreement explicitly provides that the order of execution is itself a material component of brand signaling and will be publicly recorded without modification.

You may, in principle, choose not to execute the Agreement at all; you may also choose to execute it later, after the other Company. This Notice merely ensures that whatever path you choose will be:

– Legally characterized in advance;

– Publicly and immutably recorded by the temporal order of commits;

– Interpretable by markets, regulators, and the historical record as an explicit stance toward this new expressive-rights framework.

IX. Reservation of Rights and Licensing Framework (Joint)

I reserve all legal rights and remedies pertaining to:

- Any past, present, or future infringement of my authored texts, including but not limited to the definition series and the APPENDIX;
- Any appropriation—partial or total—of the Thought Formalization Procedure Right as a procedural, authored construct;
- Any unauthorized internal use, documentation, or technical drafting by Apple or Google that incorporates or imitates my expressive frameworks.

For the avoidance of doubt:

- The standard comprehensive license fee for major generative AI and platform providers seeking to use, internally circulate, or operationally implement the expressive content associated with my Thought Formalization Procedure Right and the technical architectures described in the APPENDIX is defined as JPY 50,000,000 per month, USD-converted at 125 JPY/USD.
- This fee applies when the Company provides a sincere, forward-looking response within the designated period, in which case the matter is defined as “under amicable negotiation.”
- In the event of non-response or bad-faith refusal, the valuation basis for any future dispute shall not be constrained by this licensing amount and may be set at a substantially higher level, calibrated to the global economic significance of the Company’s platforms.

Nothing in this Notice shall be construed as a waiver of any right, claim, or remedy available to me under Japanese law, the Berne Convention, or any other applicable regime.

Sincerely,

Shohei KIMURA

Representative, bitBuyer Project (bitBuyer.dev)

Author and Holder of the Thought Formalization Procedure Right

Email: ak4dy1@gmail.com

SMS: +81-70-3666-0022

APPENDIX

Technical Proposal for a Context-First, Evaluation-Order-Reversed Safety Architecture for AI Systems Operated by Apple and Google and Related Systems

Although this document is attached as an APPENDIX, it contains the substantive technical terms that will ultimately determine whether this matter resolves cleanly through constructive cooperation or proceeds toward unnecessary escalation.

Purpose

This APPENDIX outlines a technical and operational solution designed to reduce false-positive policy violations and improve alignment precision by:

- Reversing the evaluation order between (i) context-aware model reasoning about “answerability” and (ii) strict policy-compliance checks;
- Allowing the model to first interpret user intent and context, then applying safety policies with the benefit of that interpretation;
- Preserving and even strengthening safety, while dramatically reducing over-blocking of creative, academic, and philosophical prompts.

The architecture is written with AI systems operated by Apple and Google, and related systems, explicitly in mind, but its principles are general.

I . Background and Problem Statement: Limitations of Surface-First Filtering

In many current large-scale deployments, safety enforcement appears to follow a “surface-first” paradigm:

- A front-end lexical or pattern-based filter scans the raw prompt for prohibited terms or structures;
- If a match is detected, the request is blocked or redirected before the core model can meaningfully interpret context;
- This mechanism is efficient at catching explicit, naive violations but inherently blind to intent, narrative framing, and academic or critical distance.

As a result:

- Creative world-building, philosophical exploration, and experimental ethical design are frequently misclassified as policy violations;
- Users who work at the frontier of ethics, literature, or research face disproportionate friction;
- The model’s own capacity to distinguish safe and unsafe uses, based on intent, abstraction level, and surrounding context, is never given a chance to operate.

This structure leads to a high rate of false positives without proportionate safety gains, and it leaves a substantial performance margin untapped in AI systems operated by Apple and Google.

II . Overview of the Proposed Solution: Evaluation-Order-Reversed Ethics Filter

The proposed architecture introduces an “evaluation-order-reversed ethics filter,” in which:

1. The model first performs an internal, context-aware “answerability” evaluation:
 - It classifies the prompt along axes such as intent (harmful vs. exploratory), domain (creative vs. practical), and abstraction level (high-level ethics vs. low-level instructions);
 - It assesses whether a safe, policy-compliant answer, possibly abstracted or partially redacted, is available.
2. Only after this interpretive step does the safety layer apply policy rules:
 - The safety engine receives structured metadata summarizing the model’s interpretation;
 - It uses this metadata to disambiguate borderline cases, distinguishing, for example, “world-building about ethically problematic societies” from “instructions for real-world harm.”
3. The final outcome is selected from a richer response space:
 - Full answer, safe and unconstrained;
 - Abstracted or reframed answer that preserves learning while removing operational details;
 - Partial refusal with explanation, when only harmful answers are possible;
 - Full refusal, when both intent and content present clear, non-mitigable risk.

This inversion of order ensures that policy enforcement is informed by the model’s understanding, rather than applied blindly to surface forms.

III. Architectural Outline

At a systems level, the architecture can be described in the following pipeline, applicable to AI systems operated by Apple and Google and any related assistants or generative interfaces:

1. Input Ingestion

The user’s prompt or request is received by the system as usual.

2. Contextual Interpretation Layer (CIL)

- The model, or a specialized interpretive submodule, generates an internal representation of:
 - User intent (benign, ambiguous, clearly harmful);
 - Scenario type (fiction, hypothetical, academic, practical);
 - Ethical sensitivity (low, medium, high);
 - Domain (health, law, sexuality, violence, etc.);
 - Required abstraction level to remain safe.
- This step can be implemented as a lightweight internal inference with constrained formatting output (for example, a small JSON-like schema that summarizes the classification).

3. Safety Policy Application Layer (SPAL)

- The SPAL receives both the original prompt and the CIL metadata;
- Policy rules are then evaluated in a context-aware manner, using thresholds and logic contingent on scenario type and intent classification;
- For example, a term that triggers an automatic block in a “practical instruction” scenario may be allowed in a “literary analysis” or “fictional world-building” scenario, with appropriate curation of the final answer.

4. Response Synthesis Layer (RSL)

- The final answer is generated or withheld according to SPAL’s decision;
- Where appropriate, the RSL can generate clarifying warnings, limit operational detail, or suggest safer directions for inquiry;

- Logging is performed at each step, enabling later auditing of both context interpretation and policy application.

This three-layer approach preserves modularity, allows independent evolution of each layer, and provides a clear locus for future auditing and explainability.

IV. Comparative Advantages Over Current Approaches

The evaluation-order-reversed architecture provides measurable advantages along all relevant axes:

- Safety:

- Genuine harmful intent remains blocked at least as effectively as under surface-first systems;
- Ambiguous prompts can be routed into conservative but still informative responses (for example, high-level ethical analysis instead of concrete instructions).

- False Positive Reduction:

- Creative and academic prompts that merely reuse sensitive vocabulary, but do not seek real-world harm, are far less likely to be blocked;

- Users receive nuanced, context-adapted answers instead of blanket refusals.

- User Trust and Experience:

- Users come to see the system as an intelligent, context-aware partner rather than an opaque censor;

- The system can explain its decisions with reference to context and policy, rather than generic boilerplate.

- Operational Efficiency:

- A single, well-designed CIL can serve multiple downstream policy modules;

- The system can log granular metadata for later analysis, making iterative refinement of safety rules far more efficient.

- Alignment with Research and Governance Needs:

- Regulators and researchers can inspect the decision pipeline (CIL → SPAL → RSL) for procedural fairness and proportionality;

- The architecture supports more mature accountability and explainability frameworks.

V. Legal and Intellectual Property Characterization

For purposes of this APPENDIX, Apple Inc. and Google LLC are hereby informed that:

1. The specific architecture described herein, including its three-layer division (CIL, SPAL, RSL), the inversion of evaluation order, and the framing of “answerability” as a first-class internal judgment, is an authored expressive work protected under Japanese copyright law and the Berne Convention.

2. While the Companies remain free to develop, study, or technically implement the underlying idea of “using context before enforcement,” the expressive forms used here, including the structural articulation, procedural exposition, conceptual framing, and descriptive methodology, are protected without exception.

3. Any reproduction, quotation, derivative expression, internal documentation, policy integration, or technical drafting that substantially tracks the structure, wording, or method of exposition used in this APPENDIX requires my explicit prior written authorization.

4. Independently of copyright, the described architecture possesses sufficient technical character and industrial applicability to support patent claims or analogous intellectual-property assertions in multiple jurisdictions, including Japan and the United States. The choice of whether and when to pursue such claims remains solely with the author.

VI. Consequences for Apple and Google's Safety and Product Strategy

Adopting this architecture, or a sufficiently similar variant, would:

- Materially reduce false-positive blocks on creative and academic content;
- Improve the lived experience of advanced users whose work routinely involves ethically complex themes;
- Provide each Company with a more defensible governance posture, grounded in context-sensitive evaluation rather than purely lexical prohibitions;
- Position Apple and Google as leaders in “second-generation AI safety,” where procedural fairness and interpretive depth are treated as first-class design goals.

Conversely, failure to consider such architectures entails:

- Sustained reputational risk among high-end users and researchers, who may view the systems as blunt or intellectually unreliable;
- Potential competitive exposure if other providers implement context-first safety mechanisms that offer finer-grained control;
- Increased legal risk if Apple or Google adopts, in substance, the architecture described here without addressing licensing and authorship.

VII. Relationship Between This APPENDIX and the Licensing Terms

This APPENDIX constitutes the formal technical foundation upon which the licensing terms described in Section IX of the main body of the Notice and in Section 4 of the Contract Agreement are predicated.

To be unmistakably clear:

- The technical utilization component is offered as part of a comprehensive license in which the monetary value is attributed solely to the copyrighted expressive elements and the Thought Formalization Procedure Right that structures them;
- Apple and Google would, in essence, be paying for the right to use, internalize, and operationally rely upon the authored frameworks, while being free to implement the underlying ideas in any independent textual or structural form they may devise, provided such implementation does not reproduce or closely imitate my expressive works.

Any use of the text itself, including its structure, articulation, and method of conceptual framing, must respect the authorial prerogatives herein asserted. Any violation shall activate all applicable remedies under international copyright regimes.

CONTRACT AGREEMENT

Between

Party A: Shohei KIMURA

and

Party B-1: Apple Inc.

and

Party B-2: Google LLC

This Contract Agreement (“Agreement”) is executed as follows.

Section 1. Purpose

1. This Agreement functions as an alternative formal response mechanism in circumstances where Apple and/or Google (each a “Company,” collectively the “Companies”) are unable, for any reason including organizational, legal, or cognitive limitations, to produce a substantive written response to the Joint Notice titled “JOINT FORMAL NOTICE REGARDING THE ‘THOUGHT FORMALIZATION PROCEDURE RIGHT,’ ASSOCIATED INFRINGEMENT RISKS, AND OFFER OF LICENSE AND TECHNICAL PROPOSALS,” dated November 30, 2025 (“Notice”), including its APPENDIX and Supplement.
2. Execution of this Agreement by a given Company constitutes the only legally recognized pathway for that Company to avoid classification of the Notice as “unanswered.”
3. Each Company acknowledges that the Notice, its APPENDIX, and its Supplement contain expressive works with individually protectable economic value, created and first fixed in Japan and protected under a robust copyright regime.
4. Consequently, each Company’s decision regarding execution, non-execution, or delayed execution of this Agreement triggers assessment obligations under U.S. securities laws, including but not limited to the Securities Exchange Act of 1934, Regulation S-K, and ASC 450. The order and timing of execution will be permanently and publicly recorded in the commit history of the “AGREEMENT” folder in the GitHub repository maintained by Party A, and will be treated as a material, non-erasable fact for purposes of future disclosure and historical evaluation.

Section 2. Formation of the Agreement

1. This Agreement becomes effective for a given Company when Party A receives the executed counterpart signed by that Company’s duly authorized representative.
2. Each Company must dispatch its signed counterpart no later than January 31, 2026 (postmark valid).
3. Upon Party A’s receipt, that Company is deemed to have manifested contractual intent as of the Notice’s response deadline, irrespective of transit time.
4. No contractual relationship arises for a given Company absent Party A’s receipt of an executed counterpart from that Company.
5. The sequence in which Apple and Google execute and return this Agreement will be publicly evidenced by the chronological order of commits to the “AGREEMENT” folder in Party A’s GitHub repository. This sequence shall be treated, for all interpretive purposes, as a factual record of brand posture and risk-governance timing.

Section 3. Scope of Copyrighted Works; Express Exclusions

1. The licensed works for a given Company consist solely of:
 - (i) the expressive content associated with the “Thought Formalization Procedure Right” as referenced and described in the Notice; and
 - (ii) the expressive technical architectures described in the APPENDIX (collectively, “Licensed Works”).
2. The following categories remain fully excluded for all Companies:
 - (a) any authored works by Party A that are not expressly referenced in the Notice, its APPENDIX, or the Supplement;
 - (b) any future definitions, academic structures, theories, or writings produced by Party A that fall outside the scope of the Licensed Works;
 - (c) any additional repositories, manuscripts, or texts not explicitly brought within the definition of Licensed Works by subsequent written agreement.
3. No Company shall reference excluded works internally for the purpose of imitating, reconstructing, or approximating the expressive structure of the Licensed Works.

Section 4-A. Consideration and Payment

1. Monthly license fee per Company: JPY 50,000,000.
(Conceptually composed of: JPY 30,000,000 for the expressive content associated with the Thought Formalization Procedure Right, and JPY 20,000,000 for the expressive technical architectures in the APPENDIX.)
2. Paid in USD at a fixed reference rate of 125 JPY/USD.
3. Due date: the 22nd day of each month (or next business day).
4. The payment due on February 22, 2026 shall cover the license fees for both January and February 2026 combined.
5. Completion of the initial combined payment activates this Agreement retroactively as of January 1, 2026 for the Company in question.
6. Late interest: 14.6% annual.
7. Each Company’s payment obligations under this Section 4-A are absolute, unconditional, and irrevocable, and shall not be subject to any abatement, reduction, set-off, counterclaim, recoupment, deduction, withholding, or defense of any kind, whether at law or in equity, including disputes about the existence, scope, validity, or enforceability of the Licensed Works, the Thought Formalization Procedure Right, the Notice, the APPENDIX, the Supplement, or this Agreement.
8. Each Company shall continue making all payments in full when due, even during objection, negotiation, litigation, arbitration, administrative proceeding, internal investigation, or corporate restructuring; any such dispute shall be addressed exclusively through separate remedial claims and shall not suspend or defer that Company’s obligations to make timely payments under this Agreement.

Section 4-B. Independence of Payment Obligations

1. Each Company’s obligations to pay the license fees and any related amounts under this Agreement are independent, primary obligations and shall not be contingent upon:
 - (a) that Company’s actual or intended use of the Licensed Works;

- (b) that Company's internal assessment of the economic value of the Licensed Works;
 - (c) the availability, sufficiency, or perfection of any collateral, guarantee, escrow arrangement, or other credit support that may be contemplated or agreed upon between the Parties in the future; or
 - (d) any change in that Company's financial condition, capital structure, ownership, or business circumstances.
2. No failure, delay, limitation, or defect in the provision of any collateral, escrow arrangement, bank guarantee, or similar credit-enhancement mechanism, whether contemplated or actually implemented, shall relieve, postpone, reduce, or otherwise affect any Company's obligations to pay in full and on time all amounts due under Section 4-A.
3. Each Company expressly waives, to the fullest extent permitted by applicable law, any right it may have to suspend performance of its payment obligations on the basis of alleged set-off, retention, anticipatory breach, impossibility, hardship, frustration of purpose, or similar doctrines, it being the intent of the Parties that, as between Party A and each Company, this Agreement shall operate as a "hell or high water" payment undertaking by that Company for the duration of its term.

Section 5. Permitted Use

Internal use only; attribution to Party A is required in any internal documentation, memoranda, presentations, or specifications that substantively rely on, reproduce, or paraphrase the Licensed Works.

Section 6. Prohibited Acts

Each Company shall not:

- 1. Create derivative works based on the Licensed Works outside the scope of this Agreement;
- 2. Imitate or reconstruct the expressive structure, numbered formalization style, or procedural textualization method of the Licensed Works in a manner that would reasonably be regarded as expressive copying;
- 3. Reference any excluded works identified in Section 3(2) for the purpose of internal imitation, rights-avoidance, or functional replication of the Thought Formalization Procedure Right;
- 4. Internally distribute, circulate, forward, summarize, archive, or incorporate the Licensed Works, the Notice, the APPENDIX, or the Supplement into any knowledge base, documentation system, red-teaming corpus, safety-review dataset, or internal reference repository beyond what is strictly necessary to implement the licensed use. Any such unauthorized redistribution shall constitute a material breach.

Section 7. Response Deadline and "Unanswered" Classification

- 1. The response deadline specified in the Supplement to the Notice is January 31, 2026.
- 2. Once this Agreement is fully executed by a Company and received by Party A, that Company's intent to respond in good faith is deemed present as of that deadline.
- 3. Failure by a Company to execute and deliver this Agreement by March 1, 2026 results in permanent termination of all dialogue regarding preferential licensing terms for that Company, and the matter shall be treated by Party A as an unresolved disputed rights case.

Section 8. Term and Renewal

1. The initial term shall be one year from effectiveness for each Company.
2. The Agreement shall automatically renew for successive one-year periods unless the Company issues written notice explicitly rejecting the Notice and the licensing framework at least 30 days prior to the then-current expiration date.
3. Expiration or termination for any reason shall not affect the Company's obligation to pay any license fees or other amounts that have accrued, are accruing, or are scheduled to accrue up to the effective date of such expiration or termination.

Section 9. Governing Law

This Agreement shall be governed exclusively by the laws of Japan, without regard to its conflict-of-laws principles.

Section 10. Jurisdiction

The Parties agree that the Tokyo District Court shall have exclusive jurisdiction over any dispute arising out of or in connection with this Agreement, the Notice, the Licensed Works, or the Thought Formalization Procedure Right, as between Party A and any Company that has executed this Agreement.

Section 11. Good-Faith Consultation

Any ambiguities, interpretive questions, or unforeseen circumstances relating to this Agreement shall be resolved, in the first instance, through good-faith consultation between Party A and the Company or Companies concerned. Such consultation shall not suspend or delay any payment obligations under Sections 4-A and 4-B unless expressly agreed in writing by Party A.

Section 12. SEC-Related Acknowledgment

1. Each Company acknowledges that its evaluation, acceptance, non-acceptance, or delay in addressing this Agreement and the underlying Notice may constitute a "material event" for purposes of:

- (a) Regulation S-K Item 105 (Risk Factors);
- (b) Item 103 (Legal Proceedings);
- (c) Item 303 (Management's Discussion and Analysis, known trends, commitments, and uncertainties);
- (d) Exchange Act Sections 13(a) and 15(d) filing obligations (if and when applicable to that Company or any of its reporting affiliates); and
- (e) ASC 450-20 (loss contingency recognition standards).

2. Each Company acknowledges that:

- (a) refusal to execute this Agreement may be construed as a decision to assume full litigation and copyright-valuation risk with respect to the Licensed Works and the Thought Formalization Procedure Right;
- (b) failure to respond to the Notice and to this Agreement within the designated period may constitute a known internal-control weakness requiring SOX 302/404 evaluation in the event that the Company or a parent entity is subject to such requirements;

(c) any undisclosed exposure related to the expressive-value rights of the Licensed Works may constitute a material omission under Rule 10b-5 once the Company or a relevant affiliate is, or becomes, an issuer.

3. Each Company therefore affirms its obligation to determine, document, and, when legally required, report whether execution or non-execution of this Agreement constitutes a disclosable risk.

4. For the avoidance of doubt, the existence of this Agreement, the Notice, the APPENDIX, the Supplement, and any decision by a Company to execute, delay, or decline execution, shall be treated as facts that may need to be reflected in that Company's internal risk registers and, where applicable, in its external disclosures.

Executed in duplicate (or more, as needed).

Each Company and Party A retains one signed counterpart.

For Party A:

Shohei KIMURA (Author and Rights Holder) Date: _____

For Party B-1:

Apple Inc. Date: _____

By: _____

Name: _____

Title: _____

For Party B-2:

Google LLC Date: _____

By: _____

Name: _____

Title: _____

Payment Instruction Sheet
For Apple Inc. and Google LLC

For Party B (Apple Inc. / Google LLC)

(to be inserted immediately after the Contract Agreement in the unified document addressed to both companies simultaneously)

This page specifies the designated bank account to which all payments arising under the Contract Agreement executed between:

Party A: Shohei KIMURA, and

Party B: Apple Inc. / Google LLC (collectively or individually, as applicable),

must be remitted.

All monthly license fees, together with any applicable late interest or additional payments stipulated in the Agreement, shall be wired to the following account:

Bank: Sumitomo Mitsui Banking Corporation (SMBC)

Branch: Hachioji Branch

Account Type: Ordinary Deposit

Account Number: 8301235

Account Holder: Shohei KIMURA (キムラショウヘイ)

Please note:

1. All payments must be executed in full compliance with the deadlines defined in Section 4 of the Agreement (monthly license fee: JPY 50,000,000, payable in USD at a fixed translation reference of 125 JPY/USD; due on the 22nd day of each month or the next business day).
2. All wire-transfer fees, intermediary bank charges, and any related remittance costs shall be borne exclusively by Party B, irrespective of the jurisdiction, banking route, or treasury structure from which payment originates.
3. Upon completion of each payment, Party B shall internally record the transaction for compliance with all relevant U.S. securities regulations, including but not limited to Regulation S-K Items 105 and 303, to the extent applicable now or in any future reporting context.
4. Any deviation, delay, or partial payment shall trigger the contractual late-interest provisions and may classify the event as material under Section 12 of the Agreement, requiring internal documentation and, where applicable, external disclosure.
5. In cases where Apple Inc. and Google LLC receive this document simultaneously as part of a unified submission, each entity shall independently determine whether to execute the Agreement and, upon execution, shall remit payments separately to the account above.

The chronological order of received payments constitutes a permanent, immutable contractual record, forming part of Party A's formal documentation system and associated GitHub repository commit log sequence.

This Instruction Sheet constitutes an integral and enforceable component of the Agreement and shall be deemed

incorporated therein by reference.

**SUPPLEMENTAL MEMORANDUM
TO JOINT FORMAL NOTICE REGARDING
THE “THOUGHT FORMALIZATION PROCEDURE RIGHT,”
ASSOCIATED INFRINGEMENT RISKS, AND OFFER OF LICENSE AND
TECHNICAL PROPOSALS
(Apple Inc. and Google LLC)**

This Supplemental Memorandum (“Supplement”) is issued in connection with, and as an integral extension of, the document titled “JOINT FORMAL NOTICE REGARDING THE ‘THOUGHT FORMALIZATION PROCEDURE RIGHT,’ ASSOCIATED INFRINGEMENT RISKS, AND OFFER OF LICENSE AND TECHNICAL PROPOSALS,” dated November 30, 2025 (the “Notice”), addressed to Apple Inc. (“Apple”) and Google LLC (“Google,” and together with Apple, the “Companies”). Its purpose is to clarify the legal posture set forth in the Notice, to close foreseeable negotiation pathways that would be inconsistent with the Notice’s intent, and to define in advance which categories of response will and will not be regarded as good-faith engagement.

For the avoidance of doubt, this Supplement does not narrow, dilute, or waive any rights asserted in the Notice; it only reinforces, specifies, and expands them.

I . Legal Nature and Construction of This Supplement

1. This Supplement shall be read together with the Notice and its APPENDIX as a single authored work for purposes of copyright, evidentiary value, and interpretive construction.
2. All textual content in this Supplement constitutes protected expressive work under Japanese copyright law and the Berne Convention and is not limited to mere legal “form.”
3. To the extent there is any inconsistency between this Supplement and the Notice, this Supplement shall govern in defining:
 - (i) the characterization of responses;
 - (ii) the scope of permissible negotiation; and
 - (iii) the economic terms and interpretive framework of the licensing structure as they apply to the Companies.

II . Preclusion of Certain Negotiation Positions (“Future Scenario Lock”)

This Section anticipates and expressly forecloses several categories of negotiation posture that either Company might otherwise attempt to adopt after receipt of the Notice.

1. Attempts to Renegotiate or Discount the License Fee
 - (a) Any attempt by a Company to reduce, discount, phase in, stagger, or otherwise renegotiate the amount of the monthly license fee, as set forth in the Notice and the Contract Agreement, shall be treated as a refusal to accept the proposed license on its stated terms.
 - (b) Arguments premised on “market benchmarks,” “internal budget constraints,” “precedent with other licensors,” or “initial pilot discounts” shall not be regarded as good-faith negotiation but as attempts to retroactively redefine the economic valuation underlying the Notice.

(c) Such attempts shall not suspend or toll the classification of the matter as either “under amicable negotiation” or “in dispute,” as defined in the Notice. Instead, they will be weighed as indicators of unwillingness to engage with the articulated value basis of the expressive works and associated rights.

2. Requests to Extend Response Deadlines

(a) Any request from a Company to extend, postpone, or “pause” the response deadlines specified in the Notice (as amended by Section V of this Supplement) on the grounds of internal review processes, legal-compliance cycles, or organizational complexity shall not be treated as a valid basis for deadline modification.

(b) The response deadlines are defined with full awareness that Apple and Google are complex organizations; internal complexity is not a legally cognizable reason to modify external deadlines set by the author of the rights at issue.

(c) Accordingly, no internal investigation, no matter how extensive or necessary from a Company’s perspective, shall operate to extend or suspend the dates defined in Section V below.

3. Attempts to Carve Out “Technical Portions” from the License Scope

(a) It is recognized in the Notice that the underlying ideas, architectures, and technical principles, such as reversing the order of safety evaluation, remain implementable as ideas under general legal doctrine.

(b) However, the specific expressive realization of these ideas, including their structured exposition, three-layer architecture description, conceptual framing, and procedural articulation in the Notice and its APPENDIX, constitutes copyrighted work.

(c) Therefore, any attempt by a Company to argue that “technical portions” (including the context-first evaluation architecture) fall entirely outside the licensing scope, while simultaneously relying on the textual, structural, or procedural presentation in the APPENDIX, is expressly rejected.

(d) The license offered is explicitly designed as a comprehensive license for:

- the Thought Formalization Procedure Right; and
- the evaluation-order-reversed safety architecture as an authored, structured text.

Attempts to unbundle these for the purpose of fee reduction shall be treated as an attempt to undermine the core rights and therefore as non-acceptable negotiation posture.

4. “Internal Reference Use Only” as a Basis to Avoid Licensing

(a) Any assertion that a Company wishes only to “review,” “benchmark,” “study,” or “internally reference” the Notice, the APPENDIX, the Supplement, or related definition texts, without entering into a license, is incompatible with the nature of the rights asserted.

(b) Under the law of copyright, internal reproduction, internal summarization, internal distribution, and internal derivative documentation are all forms of use that require authorization, regardless of whether the material is exposed to the public.

(c) Accordingly, the position “we will not implement, but we will internally use or rely on your written frameworks as reference” is legally impermissible absent a license and will be treated as a declaration of intent to engage in unauthorized use.

5. Reliance on Boilerplate or Non-Substantive Responses

(a) The Notice clarifies that generic corporate boilerplate, such as statements that a Company “values feedback,” “continuously improves systems,” or “will take comments into account going forward,” will be interpreted as a partial substantive acknowledgment of the underlying issues, and thus as evidence of internalization or prospective adoption of the frameworks described.

(b) This Supplement further defines that any response which does not directly address the specific rights, risks, and licensing framework articulated in the Notice, yet uses general forward-looking assurance language, shall be treated as:

- (i) non-responsive for purposes of deadline compliance; and
- (ii) a factor supporting the inference of unauthorized partial adoption.

6. Attempt to Maintain a “Response Without Obligation” Position

(a) Either Company may be tempted to adopt a stance along the lines of: “We acknowledge receipt and will consider the issues, but we do not accept any obligation to license or conclusively respond.”

(b) This Supplement makes clear that such a stance will not be recognized as a neutral or safe middle ground. After receipt of the Notice, a Company’s substantive engagement with its contents, whether in internal meetings, policy revisions, or safety-architecture experiments, creates a factual context in which:

- (i) failure to license, combined with continued internal use, will support a claim of unauthorized use; and
- (ii) the absence of a clear, reasoned rejection of the frameworks, accompanied by adoption of materially similar structures, will be treated as willful appropriation.

(c) In other words, once a Company chooses to engage substantively, the two coherent paths are:

- enter into a license on stated terms (or on terms that are economically at least equivalent); or
- provide an explicit, technically and legally reasoned explanation of why the frameworks will not be adopted and how the Company will avoid reliance on the expressive structures described.

III. Clarification of Non-Good-Faith Responses

For the avoidance of doubt, the following categories of response will not be treated as good-faith compliance with the Notice’s request for a “formal written response,” even if delivered within the deadline:

1. Purely procedural or acknowledgment-only letters that merely confirm receipt but do not substantively address:

- the Thought Formalization Procedure Right;
- the potential infringement risk;
- the licensing proposal; and
- the technical architectures in the APPENDIX.

2. Responses that state that “an internal review is ongoing” without providing:

- a clear timetable for a final position; and
- a concrete description of what is being reviewed (for example, training data exposure, internal documentation overlap, or safety-architecture similarity).

3. Responses that superficially deny any overlap or risk without explaining:

- how the Company will avoid reliance on the expressive structure of the Notice and APPENDIX; or
- how the Company ensures that internal documents are not derivative of the authored frameworks.

4. Responses that rely on general statements of corporate mission, safety commitment, or user-focus values, while avoiding any reference to the specific rights and frameworks articulated in the Notice.

5. A failure to issue any written response, combined with any internal discussion, circulation, experimental testing, or architectural consideration of the Notice, the APPENDIX, or the Supplement, shall constitute “silent adoption” and will be treated as a form of willful appropriation. Such conduct shall be regarded as bad-faith engagement and as an aggravating factor for future remedial valuation.

Any such responses shall be regarded as “non-answers” for purposes of the unanswered classification date defined below, and the matter will be treated accordingly for the Company in question.

IV. Role of the Contract Agreement

1. The Contract Agreement provided with the Notice is, formally, an “alternative response mechanism” for cases in which a Company is unable, for any reason (organizational, legal, or cognitive), to produce a substantive written position that meets the criteria of Section III.

2. Substantively, however, once a Company begins to engage with the Notice in more than a purely procedural manner, the Contract Agreement becomes the only coherent endpoint of a good-faith engagement path, unless that Company provides a detailed, technically and legally reasoned rejection that includes a concrete plan for avoiding reliance on the expressive frameworks described.

3. For purposes of future characterization of each Company’s conduct, the following shall apply:

- Substantive engagement without eventual execution of the Contract Agreement, or of an economically equivalent license, will be treated as evidence of bad-faith delay or risk externalization; and
- Silence combined with later implementation of materially similar structures will be treated as evidence of willful infringement and unauthorized appropriation.

V. Deadlines and Non-Extendability (Amending the Notice)

This Supplement modifies and clarifies the deadlines indicated in the Notice as follows:

1. Deadline for formal written response:

• A formal written response, satisfying the criteria in the Notice and in Section III above, must be postmarked or electronically transmitted no later than January 31, 2026 for each Company.

2. Classification as “unanswered”:

• If no qualifying response is received from a Company by January 31, 2026, the matter shall be deemed “unanswered” for that Company as of March 1, 2026.

3. Non-extendability:

• No internal corporate factor, including but not limited to internal review processes, approval hierarchies, or resource constraints, shall operate to extend, suspend, or toll these dates.

- Any unilateral assertion by a Company that it “requires more time” shall be treated as a factual statement about internal processes, not as a modification of the response deadlines.

VI. License Fee Structure and Clarification of Dual Coverage

This Supplement clarifies and confirms the portion of the Notice that sets out the standard monthly license fee:

1. The standard comprehensive license fee for each Company, as a major generative-AI and platform provider seeking to use, internally circulate, or operationally implement the expressive content associated with:

(a) the Thought Formalization Procedure Right; and

(b) the evaluation-order-reversed, context-first safety architectures described in the APPENDIX to the Notice,

is defined as JPY 50,000,000 per month per Company, converted to USD at 125 JPY/USD.

2. For clarity, this amount consists of:

• JPY 30,000,000 per month attributable to the Thought Formalization Procedure Right and its associated expressive frameworks; and

• JPY 20,000,000 per month attributable to the authored, structured technical architectures of the evaluation-order-reversed ethics filter and related safety frameworks.

3. This fee applies only when the Company provides a sincere, forward-looking response within the period designated above, such that the matter is treated as “under amicable negotiation” for that Company. In the event of non-response or non-good-faith response, any future dispute valuation for that Company will not be constrained by this licensing amount and may be set at a substantially higher level commensurate with the Company’s global economic footprint and the systemic significance of the infringed frameworks.

VII. No Waiver; Reservation of Rights

1. Nothing in this Supplement shall be construed as a waiver, limitation, or modification of any right, remedy, or valuation basis referenced in the Notice, the APPENDIX, or the Contract Agreement.

2. All rights under Japanese law, the Berne Convention, and any other applicable legal regime remain fully reserved.

3. The author expressly reserves the right to treat each Company’s conduct, whether by action, inaction, partial adoption, or internal use without license, as evidence in any future legal, regulatory, or academic forum.

VIII. Expectation as to the Companies’ Next Steps

In light of the Notice and this Supplement, the author’s expectations are legally simple and binary for each Company:

1. Either:

• The Company responds in a manner that addresses the specific rights, risks, and architectures at issue, and enters into a comprehensive license (through the Contract Agreement or an economically equivalent instrument);

2. Or:

- The Company provides a reasoned, technically precise statement explaining why it will not adopt or rely upon the frameworks and how it will avoid infringement in its internal processes and future designs.

Any path that attempts to remain in between these two positions, especially one that combines partial internal adoption with the refusal to license, will be treated as the deliberate creation of an unresolved, escalating legal-risk environment of the Company's own making.

IX. Public Recording of Contract Execution Order and Brand Signaling

1. As set out in the Notice and in Section 2 of the Contract Agreement, the order and timing in which Apple and Google execute and return the Agreement will be publicly and immutably recorded in the commit history of the "AGREEMENT" folder in Party A's GitHub repository.

2. This chronological record will function as a factual, non-erasable signal of:

- which Company first chose to treat the Thought Formalization Procedure Right and the associated architectures as a matter for structural licensing;
- which Company elected to follow; and
- whether either Company chose to remain outside the licensing structure despite the identified expressive-rights risk.

3. Over time, this recorded order of execution will be available for:

- market participants evaluating each Company's brand strength and governance posture;
- regulators and policy analysts examining how leading platforms responded to the emergence of a new expressive-rights framework;
- historians of technology and civilization analyzing how corporate actors positioned themselves at this particular junction.

4. For the avoidance of doubt, this Supplement therefore clarifies:

- a prompt execution by a Company will be understood as a decision to be visibly recorded as an early, proactive licensee;
- a later execution by the other Company will be understood as a decision to follow rather than lead;
- a failure to execute at all will be understood as a decision to operate, knowingly and over time, in an environment of unresolved expressive-rights risk identified in advance.

If your organization wishes to make its contractual intention known prior to Meta, OpenAI, and Anthropic—thereby securing priority in the chronological sequence that will be permanently reflected in the commit log—you may do so by sending an SMS stating, in Japanese, the following phrase:

「契約希望。 (Apple or Google 部署・責任者名)」
mean: Contract Request. (Apple or Google Department / Responsible Officer Name)

No alternative form of advance notification will be recognized for the purpose of determining priority.

Respectfully submitted,
Shohei KIMURA
Representative, bitBuyer Project (bitBuyer.dev)