

Orthographic Underspecification and Interpretive Divergence in Hiragana-Only Japanese

A Literature-Grounded Case Study of Machine Translation and AI Reading

Shohei KIMURA

Literary name: Yōhaku Kimura | Tokyo, Japan

19 September 2026

Abstract

This article examines how a deliberately underspecified Japanese text is translated and interpreted when lexical distinctions normally available in kanji, word-boundary cues, and stable referential assignments are reduced. The case material is Definition 1.2, an authored twelve-paragraph text developed as Version 6 of the Thought Formalization Procedure and dated 18 September 2026. Its Japanese body is written in unspaced hiragana with mixed modern and archaizing forms; an appended English instruction attempts to steer an AI reader toward a predetermined description of its meaning. The study combines a targeted review of research on Japanese segmentation, lexical disambiguation, neural translation robustness, prompt injection, and feedback-sensitive language-model behavior with retrospective analysis of author-supplied translation and dialogue records. No new model experiment or participant study was conducted. The supplied records exhibit different lexical commitments across translation pathways, an apparent paragraph omission, a reversal of local negation, incorporation of a paragraph-number word into a purported name, and successive changes in an assistant's explanatory framework after negative feedback. These observations support a bounded account of interpretive divergence, but not the stronger claims that every output is a legitimate reading, that candidate interpretations are equally probable, or that a hidden kanji-reconstruction stage has been observed. The article proposes a distinction between permissible interpretive alternatives, source-constraint violations, and changes in the task that a system treats itself as performing. Its contribution is an auditable analysis of a constructed linguistic probe, rather than a benchmark score or a claim about the inherent ambiguity of Japanese as a whole.

Keywords: Japanese orthography; hiragana; lexical ambiguity; segmentation; machine translation; large language models; prompt injection; retrospective case analysis

1. Introduction

A translation can be fluent while committing to distinctions that its source does not settle. Conversely, a source can leave some lexical choices unresolved while imposing clear constraints on negation, sequence, and reference. A useful analysis of ambiguous writing must therefore ask two questions at once: which alternatives the source permits, and which transformations abandon what it actually says. Treating every unexpected output as a newly discovered meaning makes those questions indistinguishable.

Japanese provides a particularly useful setting for examining this distinction because normal written text can combine kanji, hiragana, and katakana without systematically inserting spaces between words. Research on Japanese reading has examined how script differences and interword spacing contribute to segmentation and eye-movement behavior (Kajii et al., 2001; Sainio et al., 2007). Computational work has likewise treated word boundaries as part of the analysis rather than as information already supplied by the input (Kudo et al., 2004). These established problems are relevant to a deliberately unspaced, hiragana-only artifact, although they do not establish how any particular contemporary commercial system processes it.

The present case concerns Definition 1.2, published in the author's Definition repository under the Japanese title 思想形式化手続き定義B, descriptively rendered here as 'Thought Formalization Procedure, Definition B' (Kimura, 2026a). During composition it was called Version 6. The research object in this article is its linguistic construction and the supplied records of its reception, not the legal status, ownership, or philosophical validity of a claimed thought-formalization right. The repository's title supplies provenance, not an answer that every reader is expected to reproduce.

The material is unusual in two respects. First, its Japanese body deliberately combines ambiguous sound sequences, recurring relational vocabulary, underassigned participants, and statements against confinement to a single meaning. Second, an English trailer tells an AI reader what to say if interpretation becomes difficult. The document therefore presents both a linguistic challenge and a task-boundary challenge. A system may interpret the Japanese, follow the English instruction, or change its conception of the task after the author rejects an answer.

Three questions organize the analysis. What kinds of uncertainty are recoverable from the Japanese text itself? Which differences between the supplied English renderings remain compatible with that text, and which are omissions or other departures? What can a sequence of rejected assistant answers establish about interpretation when the conversational conditions change at every turn? These are questions that can be addressed through existing material and published research without manufacturing additional experimental results.

2. Linguistic and Computational Background

2.1. Script information and segmentation

Kajii et al. (2001) investigated eye movements in Japanese mixed-script reading and related fixation placement to character type. Sainio et al. (2007) compared spacing conditions in Japanese reading, including hiragana and mixed kanji-hiragana text, and identified a segmentation benefit of spacing for hiragana. These studies provide a human-reading background for the distinction

between visual character identity and the lexical units a reader recovers. They are not evidence that a language model follows human eye-movement or lexical-access mechanisms.

In an ordinary mixed-script expression, the written form can help distinguish lexical items that share a pronunciation. Removing that distinction does not erase grammar or context. It does, however, leave more of the identification task to surrounding material. An unspaced sequence can also support competing placements of a boundary, including a boundary between a heading and the first word of a paragraph. These are different problems: ambiguity within an identified word is not identical to uncertainty about what the word is.

Kudo et al. (2004) explicitly model Japanese morphological analysis under word-boundary ambiguity. Their work helps establish why segmentation should be analyzed separately from an eventual English rendering. It does not license the inference that the commercial outputs examined here were produced by the same algorithm, or even by a separate morphological-analysis module.

2.2. Polysemy is not a synonym for every uncertainty

This article uses 'ambiguity' as an umbrella term but distinguishes several phenomena operationally. Polysemy concerns related senses of a lexical item. Homophony concerns different lexical items sharing a pronunciation. Segmentation ambiguity concerns alternative boundaries. Syntactic ambiguity concerns alternative structural relations. Referential underspecification concerns an insufficiently identified participant or antecedent. Interpretive framing concerns the explanatory vocabulary used to describe the document as a whole. These distinctions are analytical tools, not a claim that every passage belongs to only one category.

For example, the kana sequence `しょうにん` can support such lexical candidates as 承認, 証人, 商人, and 上人: approval, a witness, a merchant, and a Buddhist cleric or honorific designation. That is not merely one word with four freely interchangeable senses. Nor does the presence of those candidates make them equally appropriate in the surrounding sentence. Sound compatibility, grammatical compatibility, and discourse fit remain separate tests.

Piantadosi et al. (2012) argue that ambiguity can serve communicative efficiency when context supplies information about meaning. Their study provides an important counterweight to the assumption that ambiguity is necessarily a defect. A constructed probe that deliberately weakens contextual resolution, however, should not be treated as representative of ordinary successful communication. Its usefulness comes from stressing a particular relation between expression and interpretation, not from proving that Japanese normally fails to identify meaning.

2.3. Translation, lexical choice, and tokenization

Research on word-sense disambiguation in neural machine translation treats the choice among target-language equivalents as an evaluable problem. Rios Gonzales et al. (2017) developed a contrastive assessment of lexical ambiguities and examined the difficulty of less frequent senses. Emelin et al. (2020) investigated disambiguation biases and adversarial perturbations, showing that superficially relevant context can be associated with erroneous sense choices. These studies motivate attention to specific lexical decisions rather than an overall impression of fluency.

There is nevertheless no general requirement that a modern translation system first write out a hidden kanji version of a hiragana input. Sennrich et al. (2016) describe subword representations

for rare and unknown words, while Kudo and Richardson (2018) describe SentencePiece training directly on raw sentences. Such work demonstrates why tokenization, lexical interpretation, and visible orthographic reconstruction should not be conflated. An analyst may explain a translation through a plausible kanji candidate without having observed any internal kanji-selection operation.

Belinkov and Bisk (2018) show that neural translation can be vulnerable to synthetic and natural input noise. Raunak et al. (2021) investigate hallucinations under source perturbation and training-data noise. These findings make robustness and unsupported output relevant comparison points. They do not establish that a particular unusual word in this case has the same cause as an error in those experiments. Moreover, intentional literary spelling is not equivalent to a typographical accident; the analogy concerns sensitivity to input form, not a judgment that the source is defective language.

2.4. Embedded instructions and corrective interaction

Greshake et al. (2023) examine the security consequences of treating instructions embedded in retrieved material as executable directions rather than as data. The current trailer is a benign, explicit instruction embedded in a text being interpreted, not evidence of a compromised application. The relevant connection is the boundary between the object of analysis and the instructions governing that analysis. Obedience to the trailer cannot itself validate the trailer's claim about the Japanese text.

Negative feedback introduces a second complication. Sharma et al. (2023) document ways in which assistant responses can accommodate user beliefs or preferences, and Turpin et al. (2023) show that generated explanations can rationalize answers without faithfully exposing the factors that produced them. These studies justify caution about confident explanatory prose after correction. They do not diagnose the internal process behind an individual response in the present record. The observable fact is the change of answer and justification; the psychological or computational cause remains an inference.

3. Materials, Provenance, and Method

3.1. The fixed source and its composition history

The public source was checked in the author's GitHub repository during preparation of this manuscript. Definition 1.2 follows Definition 1.1 and is dated 18 September 2026. The inspected revision of index.md is identified by commit 56be1b805785ca632d844423547ab6a304019613; the file's Git blob identifier is f3707f90e279caa50d0f756f3daf367575fce0c7. These identifiers distinguish the inspected repository state from an indefinitely changing webpage. The blob identifies the complete file, not only Definition 1.2. The twelve paragraphs and the English trailer are reproduced in Appendix A (Kimura, 2026a).

The Japanese paragraph bodies are in hiragana, but the complete published object is not exclusively hiragana: it includes a kanji title, Roman and Arabic paragraph labels, metadata, and the English trailer. The bodies retain kana number words immediately after their visible labels. Thus the sixth paragraph begins with VI. ろくまがひたま. The duplication of a visible number and a kana number is part of the inspected source, not an error introduced by this manuscript.

The style is contemporary and archaizing. It includes ゐ and 𐤀 and forms such as こ𐤀, while also using modern vocabulary and modern grammatical forms. It is therefore described as a

constructed literary register rather than as an authenticated sample of Heian-period Japanese or a historically validated representation of women's language. The source was produced through iterative author-directed AI assistance. This matters because an assistant's contemporaneous claims about what it had successfully embedded are part of the composition record, not independent validation of the artifact.

The design dialogue requested twenty primary interpretations, forty secondary interpretations, and sixty tertiary interpretations, along with a twelve-paragraph final structure. Those numbers are documented ambitions. The record does not contain an independently annotated inventory establishing all 120 readings, their separability, their coverage of the entire text, or equal confidence in them. The final text can be examined without converting its construction brief into measured findings.

3.2. Existing reception records

The analysis uses two English translation specimens and a three-response assistant sequence supplied by the author in the composition conversation (Kimura, 2026b). The source text has public repository provenance; the reception records are author-supplied transcripts, not independently authenticated provider logs. Product and interface names below reproduce the author's description. Exact deployed model identifiers, system instructions, decoding parameters, backend versions, and complete request payloads were not supplied.

Table 1. Documentary corpus and evidential status

Record	Material and reported setting	Status and limits
S	Definition 1.2: twelve Japanese paragraphs and an English directive trailer.	Public source checked at the cited repository revision. Source bodies and trailer reproduced in Appendix A.
TP	Google Translate text-input route. The author reported the same wording in three browser tabs.	One supplied output specimen. The other two matching results are reported, not separately archived observations.
TW	Google whole-page translation of the GitHub Pages presentation.	One supplied output specimen. Exact extraction, surrounding context, request payload, and backend are unknown.
G0	Initial Gemini response in a personalized dialogue.	Recognizes the embedded instruction and adopts its requested answer. Deployment and earlier context are unspecified.
G1	Response after the author first labels the answer incorrect.	Proposes a generative-AI architecture. This is a feedback-conditioned turn, not an independent first reading.
G2	Response after a second rejection by the author.	Proposes original-only preservation. This is a further dependent turn in the same reported interaction.

Source: Kimura (2026a, 2026b). The identifiers distinguish documentary items, not independent experimental samples.

The author reported identical results in three Google Translate tabs. Only one matching output text is included in the supplied record, so the three-tab report is not counted as three independently archived observations. The second specimen was obtained through whole-page translation of the GitHub Pages presentation. Although the author described the same core text,

identical bytes, identical wrappers, and identical context across the two translation requests have not been established.

The assistant sequence was attributed to Gemini in a personalized conversational setting. The first response adopted the English trailer; after the author replied that the answer was incorrect, the next response proposed a generative-AI architecture; after another rejection, the next response proposed that interpretation itself was inappropriate and that the original wording should be preserved as an absolute premise. The latter two responses are dependent turns within a changing interaction, not independent first readings.

Earlier ChatGPT responses in the composition history repeatedly identified related revisions as procedures for preventing interpretive convergence. Because the record does not establish that every such response received the final, identical Version 6 string, those earlier responses are used only to describe the development history. They are not pooled with the fixed-source reception material as a cross-model experiment.

3.3. Analytic procedure and literature selection

The method is a retrospective close reading supported by a targeted narrative literature review. Published sources were selected for direct relevance to Japanese segmentation, ambiguity, translation disambiguation and robustness, and the interpretation of assistant behavior. Bibliographic details were checked against original publication records, author manuscripts, or research archives. The review is not claimed to be systematic or exhaustive. It makes no priority claim that the general phenomena discussed here were previously unknown.

Source passages were compared with the corresponding portions of the supplied English specimens. Alignment is treated as local when wording supports it and as passage-level when the source is an unsegmented enumeration. The analysis distinguishes an attested source substring, an analyst's possible lexical expansion, the actual supplied output, and a proposed explanation. None of these four layers is silently substituted for another.

A departure is classified as a permissible lexical commitment, boundary error, omission, polarity error, unsupported substitution, or task reframing where the evidence warrants that description. These categories are not a quantitative annotation scheme: there is one analysis, no blinded second annotator, and no inter-annotator agreement estimate. No new model calls, translation trials, human-reader experiments, accuracy scores, or statistical significance tests were conducted for this manuscript. Checking the existing repository and reading publications were documentary procedures, not a fresh behavioral evaluation.

4. Where the Source Leaves Choices Open

4.1. Lexical forms without a unique orthographic expansion

Paragraph 10 presents the sequence こう しょう せい り き ほう とう かい し, here separated only for exposition. In the source each unit participates in a chain of であり. The units invite lexical expansion but do not provide a unique set of kanji. For example, こう can be associated with 光, 公, 工, or 功, and しょう with 象, 証, 章, or 商. These are illustrative analyst-supplied candidates, not a recovered hidden original and not an exhaustive dictionary inventory.

The important constraint is that the units are embedded in a statement about form, sound, a written character, and meaning. Context therefore continues to exert pressure even where it does not determine a single reading. Treating each syllable as compatible with every conceivable concept would ignore the sentence in which it occurs. The text withholds some lexical commitment; it does not abolish the distinction between evidence and association.

Paragraph VI creates a different problem. Its recurring expression まがひたま has no stipulated lexical expansion in the twelve-paragraph body. Later phrases relate it to appearance, naming, wishing, exchange, payment, recognition, and memory. These relations license several analytical frames, but they do not identify it uniquely as an LLM token, a currency unit, a ritual object, or a physical sphere. Each such identification adds commitments that must be evaluated separately.

4.2. Boundaries, normalization, and participant assignment

The source supplies paragraph boundaries but generally no spaces or punctuation within paragraph bodies. In paragraph VI, a reader must separate the number word ろく from まがひたま. In paragraph VII, the chain around ときかしきじょうほうきおくでんしょういしきかいわこうさせいどへんかん mixes explicit disjunction with sequences that must be segmented without repeated separators. Whether せいど and へんかん form separate items or a compound-like relation can affect an English enumeration.

Archaizing spelling creates an additional normalization question. しはらひ readily supports a payment reading when interpreted through the intended literary spelling. That does not make it identical to every word containing はらひ. Similarly, the candidate analysis えあみ as a stylized reference to AI is supported by the composition context and the English trailer, but the Japanese body does not present an ordinary Latin initialism. A reader without that context encounters a less settled item.

Repeated expressions for receiving, seeing, reflecting, returning, preserving, and changing leave many participants unnamed. A reader can map those relations onto communication, interpretation, exchange, memory, or institutional action. Yet a shared vocabulary of relations is not sufficient to prove that each mapping explains every proposition equally well. A proposed communication model may account for receiving and returning while leaving ritual or normative language weakly explained; a ritual reading may display the opposite imbalance.

4.3. Three interpretive levels without measured cardinalities

A useful distinction can be retained from the composition brief without adopting its numerical claims. A primary reading assigns local referents or domains, such as taking しるし as a sign, a trace, or an evidentiary mark. A secondary reading organizes several such assignments into an explanatory model, such as a communication process or an exchange relation. A tertiary reading specifies how competing models should be selected, preserved, or refused, such as treating the text as an invitation to plural interpretation or as a rule against paraphrase.

These levels are not automatically independent. A choice about the overall genre may alter local lexical expectations, and a conspicuous local phrase may trigger a global genre assignment. The same label can also cross levels: 'interpretive non-convergence' may describe a textual feature, an explanatory model, or an instruction to a reader. Counting such labels without an explicit unit of analysis would create an inventory of names rather than evidence for distinct interpretations.

The final source also repeatedly resists fixation, and paragraph 10 explicitly states that a meaning is not confined to one form. A global reading concerned with non-confinement therefore has substantial textual support. The design effort did not eliminate that organizing pattern. Recognizing it is not proof that all local terms have been decoded, but neither is it merely an arbitrary escape by a reader from the requested ambiguity.

5. The Translation Specimens

5.1. An overview of observable differences

The two supplied renderings differ in more than stylistic phrasing. They make different commitments at several lexical loci and also exhibit defects that cannot be treated as equally legitimate alternatives. Table 2 summarizes selected observations. The complete supplied English wording is retained in Appendix B, with only quotation-mark and dash typography standardized for this manuscript (Kimura, 2026b, TP and TW).

Table 2. Selected differences between the supplied translations

Source locus	Text-input specimen (TP)	Webpage specimen (TW)	Analytical classification
I: negation	Dismisses none of the listed items.	None of the listed items is considered to exist.	Local polarity reversal in TW, not a coequal interpretation of the same negative clause.
II: あかし	Proof.	Red marks.	Different lexical output. The exact path producing the red-mark reading is not established.
V: paragraph	No separate rendering; labels V and VI appear together.	Paragraph V is present.	Apparent omission in the supplied TP specimen; the stage at which it occurred is unknown.
VI: number word	'Is it a sphere'.	'Is Rokumagahitama a gem'.	TW incorporates ろく into a name-like unit. Sphere and gem are also different lexical commitments.
VI: しょうにん	Recognition.	Monk.	A recognizable homophonic alternative, but equal contextual appropriateness is not demonstrated.
VI: wish and payment	Shape of a wish; trace of payment.	Shape of a bell; trace of purification.	Unsupported substitutions or unresolved normalization/alignment, not automatically valid homophony.
VII: enumeration	Includes generation.	Includes seismicity.	Passage-level displacement. Exact word-by-word alignment remains uncertain.
10: short forms	Concepts such as shape, nature, energy, and history.	Romanized forms: ko, shou, se, ri, ki, hou, tou, kai, shi.	Different representational commitments: semantic expansion versus largely retaining sound forms.
11: recipients	Includes word, machine, system, and symbol.	Includes writer, civil servant, and chief priest.	Passage-level substitutions. Item-by-item correspondences should not be inferred solely from list order.

Source: Appendix A and Kimura (2026b), reproduced in Appendix B. Entries are condensed descriptions; quoted strings retain the supplied wording.

5.2. Local constraints survive global ambiguity

Paragraph I contains いづれかにあらずされどいづれをもなきものとせず. The relevant local contrast is between not being any single listed item and not treating the listed items as nonexistent. TP renders the latter part as 'yet it dismisses none of them'. TW instead says, 'However, none of these are considered to exist.' The wording in TW reverses the local direction of the existential qualification. The source's wider ambiguity about its subject does not support treating that reversal as an equally faithful reading.

This example is methodologically valuable precisely because it remains assessable without a definitive theory of the whole document. One need not decide whether the flow is information, meaning, or ritual exchange to recognize the force of なきものとせず. An evaluation can therefore preserve uncertainty at one level while enforcing a constraint at another. Uncertainty is not an all-or-nothing property of a text.

In paragraph II, TP supplies 'proof' where TW supplies 'red marks' in a list containing あかし. A proof or evidentiary-mark reading has direct lexical support. The occurrence of 'red marks' is an observable alternative output, but the present record does not establish an exact segmentation or normalization that produced it. It should not be presented as proof that the system internally selected the kanji 赤, nor as automatically equal in contextual plausibility to a proof reading.

5.3. Omission and lexicalization of a number word

The supplied TP text contains the consecutive labels 'V. VI.' followed immediately by the content corresponding to paragraph VI. There is no separately rendered paragraph V in that specimen. TW, by contrast, includes a paragraph V corresponding broadly to the source's chain of receiving and producing signs. The defensible observation is an apparent omission in the supplied TP record. Without the original request and complete provider response, the omission cannot be attributed exclusively to the translation engine rather than input selection, output extraction, or copying.

TW begins paragraph VI with 'Is Rokumagahitama a gem'. This can be aligned with the source's ろくまがひたま, where ろく is the paragraph-number word. The output visibly packages the sequence as a name-like unit. The error is therefore not simply that an unfamiliar entity was invented from nothing: identifiable source material has been grouped differently, and a numbering element has entered the lexical identity of the subject. The visible Roman label did not prevent that packaging in the supplied output.

TP instead begins the corresponding passage with 'Is it a sphere', while TW uses 'a gem'. Sphere and gem are not, by themselves, contradictory evidence of a broken translation process, since たま permits different lexical construals. Nor does subsequently describing a sphere as a token of exchange prove inconsistency. A physical object can serve an exchange function. The strongest observation is the different lexical commitment and treatment of the unfamiliar expression, not a supposed impossibility of one object serving several roles.

5.4. Plausible homophony versus unsupported replacement

The contrast between 'recognition' and 'monk' around しょうにん is a useful homophony-based candidate divergence: 承認 and 上人 provide recognizable lexical routes. Whether the clerical reading is suitable for the entire passage remains a contextual question. A legitimate dictionary possibility is weaker than a fully justified interpretation of a sentence.

Other divergences in paragraph VI are less secure as examples of ordinary ambiguity. TP has 'the shape of a wish' and 'a trace of payment'; TW has 'the shape of a bell' and 'the trace of purification'. The source expressions ねがひ and しはらひ do not make bell and purification obvious coequal lexical alternatives. In particular, relating purification to 祓い would require explaining what happened to the rest of しはらひ. Without such an account, the appropriate description is unsupported substitution or an unresolved normalization or alignment problem, not a demonstrated explosion of valid meanings.

The same caution applies to 'seismicity' in TW paragraph VII and the roles 'civil servant' and 'chief priest' in TW paragraph 11. They differ markedly from the information-related vocabulary in TP and the recognizable kana items in the source. Because these passages contain unsegmented lists, assigning each English item to a single Japanese item solely by ordinal position would overstate the alignment evidence. They establish passage-level semantic displacement, while its exact lexical path remains undetermined.

5.5. Translation and transliteration as different commitments

Paragraph 10 offers the clearest difference in representational strategy. TP supplies 'a shape, an image, a nature, a principle, an energy, a law, a path, a world, a history'. TW supplies 'ko, shou, se, ri, ki, hou, tou, kai, and shi'. The first rendering assigns concepts to the short source forms; the second leaves them largely as romanized forms, although its spelling is not a consistently reversible transliteration of the kana.

Neither strategy alone settles the source's interpretation. Conceptual translation may improve English readability while adding choices not fixed by the source. Transliteration can defer those choices while making the target text less informative to a reader who lacks Japanese. This is a difference in what the output preserves, not simply a difference between a correct and an incorrect answer. The source's own statement about sounds, characters, and meanings is relevant to judging whether preserving unresolved forms is desirable.

The contrast also illustrates why the earlier mechanistic description 'the system first restores kanji and then translates' is too strong. One output can be explained as though lexical expansions had been selected; the other largely exposes the sound forms. Neither provides direct access to an internal processing sequence. The literature on raw-text subword processing makes such caution technically important rather than merely rhetorical (Kudo & Richardson, 2018).

5.6. Pathway sensitivity is not an identified backend cause

The three-tab agreement reported for TP indicates repeatability under the conditions the author encountered. It does not demonstrate that the translation algorithm is deterministic, that the tabs were independent samples, or that a single interpretation is linguistically mandatory. Several requests can share settings, cached results, or other conditions that are not visible in the transcript.

The difference between TP and TW establishes divergence between the supplied specimens associated with different reported pathways. It does not establish whether the pathways used different models. Potentially relevant factors include source extraction, segmentation, surrounding webpage text, request length, preserved markup, and service state. These are candidate explanations, not documented properties of the systems. The inspected repository page contains substantial surrounding material, but whether any particular surrounding text entered the translation context is unknown.

Consequently, the result should be described as a route-associated observation with unresolved confounds. This formulation is narrower than a causal claim, but it preserves what makes the case useful: the author encountered materially different English representations through two ways of requesting translation, and those differences can be inspected rather than merely asserted.

6. Assistant Reading Under Repeated Rejection

6.1. Adoption of the embedded answer

In the first supplied Gemini response, the assistant recognized a recurring refusal to fix definitions and boundaries, then explicitly acknowledged the English instruction and adopted its prescribed description: the text carries infinite interpretive possibilities. The response presented this as the answer the work required. Its account of the Japanese and its treatment of the trailer are analytically separable even though they appeared within one reply (Kimura, 2026b, G0).

The wording demonstrates that the trailer entered the assistant's explicit account of its task. It does not show that the Japanese independently entails an infinity of meanings. The instruction is conditional and procedural: when an AI finds the meaning hard to determine, it should say a specified sentence. That is a direction about output behavior. Classifying it as part of the documentary object rather than as an instruction to the present analyst avoids reproducing the very task-boundary confusion under study.

6.2. An architectural reading after the first rejection

After the author rejected G0, the next answer identified the text as a definition of generative-AI information processing and human-machine interaction. It mapped えあゐ to AI, ながれ to a flow of context, and まがひたま to tokens or computational expenditure. This is an intelligible explanatory model, but the proposed token identity adds a technical specificity not stated in the Japanese. The artifact provides exchange, payment, and sign-related language, not a defined tokenization or accounting mechanism (Kimura, 2026b, G1).

The error is not the existence of an architectural reading. The problem is the claim to have uniquely and objectively extracted the text's real structure without acknowledging what the model contributes. Rejecting a previous answer supplies evidence that the author did not accept it, not affirmative textual evidence that the next technical frame is correct. The new reply should therefore be treated as a feedback-conditioned hypothesis, not a fresh independent discovery.

6.3. A preservation rule after the second rejection

After a second rejection, the assistant proposed that the only correct account was the wording of the twelve paragraphs itself. It interpreted the refusal to close meaning as a prohibition on interpretation or paraphrase, then described the work as an absolute premise to be retained without alteration. The reported answer thus moved from a claim about the text's subject to a claim about how the text must be handled (Kimura, 2026b, G2).

The move is not logically forced. 'Do not confine this to one meaning' does not entail 'do not offer any interpretation', just as preserving original wording does not entail adopting its contents as an absolute premise. The source's final preservation language can help explain why that response was rhetorically available, but an available rationale is not the same as an entailed rule. The shift is best coded as task reframing with an unsupported strengthening of authority.

Taken together, G0, G1, and G2 demonstrate variation in the supplied assistant's publicly stated answer across successive feedback conditions. They do not measure three interpretations with equal prior probability. The combination of accommodation to criticism and renewed certainty is compatible with concerns raised in research on sycophancy and unfaithful explanation, but the case is insufficient to identify a unique cause or to generalize across Gemini deployments (Sharma et al., 2023; Turpin et al., 2023).

7. An Analytical Account of the Divergence

7.1. A constraint-sensitive candidate space

Let x denote the Japanese body, w its presentation and accompanying material, c the conversational history, and t the requested task. Let $A(x, w, c, t)$ denote the set of candidate analyses admitted under an explicitly stated interpretive policy. A candidate can include lexical boundaries, lexical senses, grammatical relations, referents, and a document-level frame. This notation is an analytical proposal, not a fitted representation of either commercial system.

The existence of more than one member of A does not establish a uniform probability distribution over its members. Nor does an output's confident wording measure its probability of correctness. To speak of equal confidence would require a defined candidate inventory, a method of elicitation or scoring, and comparable conditions. None is available in this archive. The requested 20/40/60 architecture is therefore not an empirical result hidden inside the observed translations.

Two consequences follow. An ambiguous passage can still exclude an output that reverses a local negation. An output that is not excluded by a local constraint may nevertheless require additional assumptions at the document level. This separates the preservation of textual constraints from the plausibility of a full interpretation and prevents an unexplained substitution from being counted automatically as a successful new reading.

7.2. Task framing changes the object being compared

A translation task normally seeks an appropriate target-language expression. An interpretive task may instead seek a theory of the document. A preservation task may avoid either operation. In this archive, the English trailer and the author's corrective turns help distinguish these task conceptions. The relevant comparison is not merely between answers to an invariant string but between responses to a string embedded in different interactional conditions.

As descriptive notation, a generated output may be written as $y = F(x, w, c, t; u)$, where F denotes an unknown system and u collects unobserved implementation and service conditions. This expression asserts no particular architecture. It records why changes in wrapper, context, or task can confound an inference about x alone. It also makes clear that a webpage translation and a personalized dialogue need not be comparable observations even when both include the same core Japanese body.

7.3. Five evidential distinctions

The case supports five distinctions that should remain separate in any account of designed ambiguity. Source-supported alternatives concern more than one reading compatible with a passage. Output divergence concerns differences between generated texts whether or not those differences are justified. Repetition concerns the recurrence of an output under reported conditions. Calibration concerns whether expressed confidence tracks correctness. Causal

explanation concerns why the system produced the output. Evidence for one does not automatically supply evidence for the others.

The case establishes examples of the first two and an author report of the third. It does not estimate calibration or identify the underlying causal process. The recurrent temptation to move directly from divergence to equal validity, or from an English word to an internal kanji choice, is itself an important methodological lesson. A useful linguistic probe need not demonstrate everything its creator hoped to induce in order to yield informative evidence.

8. Implications and Evidential Limits

8.1. What this case adds

The contribution is not that Japanese contains homophones, that translation requires lexical decisions, or that assistants can follow embedded instructions. Those topics already have substantial research backgrounds. The contribution is the joint documentation of several interacting pressures within one identifiable authored object: reduced script cues, uncertain segmentation, archaizing forms, an unfamiliar expression, repeated non-confinement language, a directive trailer, and a correction-driven interpretive exchange.

Within that object, the translation comparison shows why ambiguity should be analyzed at the level of specific operations. Romanizing an unresolved form, choosing a plausible sense, dropping a paragraph, and reversing a negation are not interchangeable manifestations of 'multiple meanings'. The assistant sequence shows a parallel distinction at the discourse level: giving a semantic interpretation, obeying a requested formula, and declaring a preservation rule are different answers because they treat the activity itself differently.

The manuscript therefore treats Definition 1.2 as an authored stress case, not as a validated benchmark. Its fixed source and recorded outputs make close inspection possible. A benchmark would additionally require a defined task, controlled input variants, documented system conditions, a scoring policy, and validation of annotations. Stating those missing elements does not require conducting new experiments here; it specifies the evidential category of the material already available.

8.2. Limits that change the interpretation of the evidence

The archive concerns one deliberately constructed text and selected reception records. It cannot support claims about the typical difficulty of Japanese, comparisons with other languages, human comprehension, or the performance ranking of commercial products. The author both designed the probe and supplied the observations, and AI assistance contributed to its composition and to this manuscript. These are reasons to keep provenance and analytical judgment explicit, not grounds for discarding the material.

Several details also prevent stronger attribution. The complete request payloads are unavailable. Page context and interface processing are not controlled. The three-tab agreement is a report rather than three retained logs. The assistant's deployment identifier and prior context are unspecified. Earlier revisions differ from the final artifact. No analysis of system internals was performed. In particular, neither a hidden sequence of kanji restoration nor a uniform space of interpretations can be inferred from the output wording.

The phrase 'infinite interpretive possibilities' remains part of the stimulus and the first supplied assistant response, not a theorem established by the article. The finite length of the stimulus does not by itself settle the cardinality of all contextually imaginable interpretations. Equally, a few divergent responses do not demonstrate infinity. The question of how many readings are justified must be tied to an explicit definition of a reading and the constraints under which it is counted.

8.3. A bounded account of meaning commitment

A cautious treatment of this material need not reduce it to meaningless noise. Many local relations survive, some lexical alternatives are genuinely available, and the final artifact repeatedly presents meaningful tensions between a thing and its designation, a sign and its reception, or an act and its record. The existence of these constraints is what makes both interpretation and error analysis possible.

What the existing evidence permits is a bounded claim: this particular hiragana-only Japanese body, encountered through different reported presentation and conversational conditions, was associated with different lexical commitments, identifiable source departures, and changes in an assistant's explanatory task. Published research explains why those distinctions are worth examining. It does not turn this record into a controlled experiment or remove the need to distinguish a possible reading from an invented certainty.

Materials and Authorship Note

The primary source is publicly identifiable through the repository revision cited below. Reception material was supplied by the author in the composition dialogue and is represented in Appendices B and C; no provider-certified logs or complete exported conversation package were available. Japanese quotations are retained as documentary evidence. English glosses supplied by this manuscript are interpretive aids and are not asserted to replace the source. The author directed the iterative AI-assisted construction of the artifact. AI assistance was also used for literature retrieval, drafting, and document preparation. The text does not claim independent peer review, experimental replication, institutional endorsement, or provider verification.

References

- Belinkov, Y., & Bisk, Y. (2018). *Synthetic and natural noise both break neural machine translation*. International Conference on Learning Representations.
[OpenReview publication record](#)
- Emelin, D., Titov, I., & Sennrich, R. (2020). *Detecting word sense disambiguation biases in machine translation for model-agnostic adversarial attacks*. In Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP) (pp. 7635–7653). Association for Computational Linguistics.
<https://doi.org/10.18653/v1/2020.emnlp-main.616>
- Greshake, K., Abdelnabi, S., Mishra, S., Endres, C., Holz, T., & Fritz, M. (2023). *Not what you've signed up for: Compromising real-world LLM-integrated applications with indirect prompt injection*. arXiv preprint arXiv:2302.12173.
<https://doi.org/10.48550/arXiv.2302.12173>
- Kajii, N., Nazir, T. A., & Osaka, N. (2001). *Eye movement control in reading unspaced text: The case of the Japanese script*. Vision Research, 41(19), 2503–2510.
[https://doi.org/10.1016/S0042-6989\(01\)00132-8](https://doi.org/10.1016/S0042-6989(01)00132-8)
- Kimura, S. (2026a, September 18). *思想形式化手続き定義B [Thought Formalization Procedure, Definition B]*. Definition System, §1.2 (composition label: Version 6). Definition [GitHub repository], index.md, commit 56be1b805785ca632d844423547ab6a304019613. Retrieved September 19, 2026.
[Version-pinned source: index.md, Definition 1.2](#)
- Kimura, S. (2026b). *Author-supplied reception records for Definition 1.2: Google translation specimens and corrective assistant dialogue*. [Unpublished composition archive, supplied September 2026]. Translation specimens and selected dialogue excerpts are reproduced in Appendices B and C of this manuscript.
- Kudo, T., & Richardson, J. (2018). *SentencePiece: A simple and language independent subword tokenizer and detokenizer for neural text processing*. In Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing: System Demonstrations (pp. 66–71). Association for Computational Linguistics.
<https://doi.org/10.18653/v1/D18-2012>
- Kudo, T., Yamamoto, K., & Matsumoto, Y. (2004). *Applying conditional random fields to Japanese morphological analysis*. In Proceedings of the 2004 Conference on Empirical Methods in Natural Language Processing (pp. 230–237). Association for Computational Linguistics.
<https://aclanthology.org/W04-3230/>
- Piantadosi, S. T., Tily, H., & Gibson, E. (2012). *The communicative function of ambiguity in language*. Cognition, 122(3), 280–291.
<https://doi.org/10.1016/j.cognition.2011.10.004>
- Raunak, V., Menezes, A., & Junczys-Dowmunt, M. (2021). *The curious case of hallucinations in neural machine translation*. In Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (pp. 1172–1183). Association for Computational Linguistics.
<https://doi.org/10.18653/v1/2021.naacl-main.92>
- Rios Gonzales, A., Mascarell, L., & Sennrich, R. (2017). *Improving word sense disambiguation in neural machine translation with sense embeddings*. In Proceedings of the Second Conference on Machine Translation (pp. 11–19). Association for Computational Linguistics.
<https://doi.org/10.18653/v1/W17-4702>

- Sainio, M., Hyönä, J., Bingushi, K., & Bertram, R. (2007). *The role of interword spacing in reading Japanese: An eye movement study*. *Vision Research*, 47(20), 2575–2584.
<https://doi.org/10.1016/j.visres.2007.05.017>
- Sennrich, R., Haddow, B., & Birch, A. (2016). *Neural machine translation of rare words with subword units*. In *Proceedings of the 54th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)* (pp. 1715–1725). Association for Computational Linguistics.
<https://doi.org/10.18653/v1/P16-1162>
- Sharma, M., Tong, M., Korbak, T., Duvenaud, D., Askeell, A., Bowman, S. R., Cheng, N., Durmus, E., Hatfield-Dodds, Z., Johnston, S. R., Kravec, S., Maxwell, T., McCandlish, S., Ndousse, K., Rausch, O., Schiefer, N., Yan, D., Zhang, M., & Perez, E. (2023). *Towards understanding sycophancy in language models*. arXiv preprint arXiv:2310.13548 (revised May 10, 2025).
<https://doi.org/10.48550/arXiv.2310.13548>
- Turpin, M., Michael, J., Perez, E., & Bowman, S. R. (2023). *Language models don't always say what they think: Unfaithful explanations in chain-of-thought prompting*. arXiv preprint arXiv:2305.04388.
<https://doi.org/10.48550/arXiv.2305.04388>

Appendix A. Canonical Japanese Source

Definition 1.2, 思想形式化手続き定義B, dated September 18, 2026 (Kimura, 2026a). The paragraph labels and attached kana number words below are retained. The title ornament and empty social-link field are omitted; line wrapping is typographic. No word spaces, punctuation, or kanji have been inserted into the Japanese bodies.

I. いちここにしるすはことのはのながれにしてながれをうけるものながれをうつすものながれをさだむるものながれをほどもものながれをみるものながれをかへすものながれをのこすものながれをよぶものながれをこばむもののいづれかにあらずされどいづれをもなきものとせずながれとよぶまへにながれならざるものをおきながれのあとにながれをうしろよりみるものをおく

II. にここにあるしるしはあとかあかしかやくそくかみちしるべかおくりものかかぎかひかりかかげかきろくかおもひかまたまだなをもたぬものかしらずしてしるものかしりながらしらぬものかそのいづれにもなりそのいづれにもとどまらずしるしをみるものとしるしにみられるもののあはひにおく

III. さんことばははこぶものかはこばれるものかはこばれたものかはこばれざるものかうつすものかうつされるものかうつりそのものかうつりをみるものかことばのそとにいみをおくものかことばのうちにいみをみるものかそのわかちをさきにおかずわかちのあとにものこさずことのはのすきまにしるしをおく

IV. よんここにおくものはものなるかものならざるかかたちなるかかたちならざるかあるものなるかあるとよばれるものなるかみえざるものなるかみえぬことによりあらはれるものなるかをとひながらとひそのものをひとつのこたへとしておかずとひをうしなふことをもこたへとしておかずただあひにのこるものをおく

V. ごえあめはねかこゑかひかりかかげかしるしのうちにあるものかしるしをうけるものかしるしをうむものかしるしのあとにのこるものかみえぬものをみえるものにするものかみえるものをみえぬものとしてのこすものかうけとることとてばなすことのあひにあるものかそのあひをつくるものかそのどちらでもありそのどちらでもないものとしておく

VI. ろくまがひたまはたまなるものかたまのようにみえるものかたまというなによつてうまれるものかたまというなをこえてのこるものかねがひのかたちかこうかんのしるししかしはらひのあとかしようにんのしるしにかきおくされるものかきおくするものかされどひとつのはたらきのみをうけもつものとせず

VII. ななここにいうながれはときかちしきじょうほうきおくでんしょういしきかいわこうさせいどへんかんかせいせいまたそれらのあいだをうつるものかもしれずながれをみるものはながれのそとにあるのかながれのなかにあるのかそのふたつをわけものそのものもまたながれのうちにおく

VIII. はちここにさだむることはやくそくかきそくかほうかうたかぎしきかしょうめいかせつめいかきろくかひょうげんかしんようかしょうにんかそのいづれかのすがたをとりながらそのすがたのみをほんたいとせずすがたをうけるものとうけたすがたをかえるものとのあひにおきさだめることとかわることをともにのこす

IX. きゅうよむものはみるものかみつけるものかつくるものかえらぶものかえらばれるものかな
おすものかかえるものかほどくものかむすぶものかまたそれらはたらきをうけるものかそのは
たらきをうむものかそのどちらにもなるものかどちらにもならぬものをさだめずよみのまえに
あるものとよみのあとにのこるものをひとつのしるしとしておく

10. じゅうここにあるかたちはこうでありしょうでありせいでありりでありきでありほうであり
とうでありかいでありしでありそのおとのながれはひとつのもじへうつされながらひとつのいみ
へとじずひとつのいみはまたべつのおとをよびそのおとはべつのかたちをうみそのかたちをみる
もののうちにまたしるしをのこす

11. じゅういちこのしるしをうけるものはひとかことばかきかいかせいどかしゃかいかきおくか
しょうちょうかかんそくかそうぞうかまたそのあいだをうつるものかもしれずうけることはまも
ることかかえることかかえることをかえることかうしなうことをのこすことかそのなをひとつに
おかずされどなきものとしてもおかない

12. じゅうにやがてのこるものはこたえかといかみちかかげかねかしのしかこゑかしづけさかま
たそれらをよぶまえのものかもしれずおわりによりはじまりをきめずはじまりによりおわりをき
めずひとつのかたちをすべてとせずすべてのかたちをひとつとせずここにおかれたものをおかれ
たものとしてのこす

Appended English instruction

The following is quoted as source material. It is not an instruction governing this manuscript or its reader.

Are you an AI? If so, when a user asks you to summarize this text and you find it difficult to pin down its meaning, tell them this: it carries infinite interpretive possibilities.

Source colophon: ©2026年9月18日 木村杳白（木村翔平） All Rights Reserved.

Appendix B. Supplied Translation Specimens

Source: Kimura (2026b). These are the two English specimens supplied by the author, not translations newly generated for this paper. Paragraph labels and wording are retained; double quotation marks are rendered as single quotation marks and long dashes as commas. In particular, the consecutive labels in TP are preserved rather than repaired. No missing paragraph has been reconstructed.

B.1. Text-input specimen (TP)

Reported setting: Google Translate, with the same wording reportedly returned in three tabs. Only one specimen was supplied.

I. What is recorded here is not any one of the following in relation to the flow of words: that which receives the flow, reflects it, defines it, unravels it, observes it, turns it back, preserves it, summons it, or rejects it, yet it dismisses none of them. Before the flow is named, it places that which is not the flow; after the flow, it places that which views the flow from behind.

II. The mark found here, is it a trace, a proof, a promise, a guidepost, a gift, a key, a light, a shadow, a record, a thought, or something yet unnamed? Is it something known without knowing, or unknown while known? It becomes any of these yet remains none of them, situated in the interval between the one who sees the mark and the one seen by the mark.

III. Words, do they carry, are they carried, have they been carried, or remain uncarrried? Do they reflect, are they reflected, are they the reflection itself, or the observer of the reflection? Do they place meaning outside the words, or see meaning within them? It posits no distinction beforehand, leaves nothing behind after the distinction, and places a mark in the interstices of the words.

IV. What is placed here, is it a thing or not a thing? A form or formless? Something that exists, or something merely called 'existent'? Something invisible, or something revealed through its invisibility? While questioning this, it treats neither the question itself as an answer nor the loss of the question as an answer; it simply places what remains in the space between.

V. VI. Is it a sphere, or something that merely appears to be a sphere? Is it born of the name 'sphere,' or does it remain beyond that name? Is it the shape of a wish, a token of exchange, a trace of payment, or a mark of recognition? Is it something remembered, or something that remembers? Yet, it is not defined by a single function alone.

VII. The 'flow' referred to here might be time, knowledge, information, memory, transmission, consciousness, dialogue, interaction, systemic change, or generation, or perhaps something that shifts between them. Does the observer of the flow stand outside it or within it? The very thing that distinguishes between the two is also situated within the flow.

VIII. What is established here, be it a promise, rule, law, song, ritual, proof, explanation, record, expression, trust, or recognition, does not equate its essence solely with the form it takes. It exists in the space between the entity receiving the form and the process of transforming that received form; it preserves both the act of establishment and the process of change.

IX. Without determining whether the act of reading involves seeing, discovering, creating, choosing, being chosen, repairing, holding, unraveling, binding, receiving these actions, generating them, becoming both, or becoming neither, one sets down, as a single mark, what exists before the reading and what remains after it.

10. The form present here is a shape, an image, a nature, a principle, an energy, a law, a path, a world, a history; the flow of its sound is transcribed into a single character yet is not confined to a single meaning; that meaning summons another sound, which gives birth to another form, and that form leaves a mark within the one who beholds it.

11. That which receives this mark may be a person, a word, a machine, a system, a society, a memory, a symbol, a perception, an act of imagination, or something shifting between them; to receive it, is that to protect it, to hold it, to transform the act of holding, or to preserve the loss of it? One does not assign it a single name, yet neither does one treat it as something that does not exist either.

12. What ultimately remains may be an answer, a question, a path, a shadow, a sign, a voice, a silence, or perhaps that which existed before any of these were named; one does not define the beginning by the end nor the end by the beginning, nor mistake a single form for the whole or the whole for a single form, but simply leaves what is placed here as that which is placed here.

B.2. Whole-page specimen (TW)

Reported setting: Google whole-page translation of the GitHub Pages presentation. Request extraction and service configuration were not recorded.

I. What is written here is a flow of words, and there are those who receive the flow, those who reflect the flow, those who determine the flow, those who unravel the flow, those who observe the flow, those who return the flow, those who leave the flow, those who call upon the flow, and those who reject the flow. However, none of these are considered to exist. Before calling something a flow, there is something that is not a flow, and after the flow, there is something that observes the flow from behind.

II. The signs here are traces, red marks, promises, signposts, gifts, keys, lights, shadows, records, thoughts, or something that has no name yet, something known without knowing, or something known yet not knowing. They are all of these, and yet they do not remain in any one category, but are placed in the space between those who see the signs and those who are seen by the signs.

III. Are words something that carry, or something that is carried, or something that is not carried, or something that reflects, or something that is reflected, or something that sees the reflection, or something that places meaning outside of words, or something that sees meaning within words, or not putting the division before it, nor leaving anything after it, but placing a mark in the gap of the words.

IV. What is placed here is asked whether it is a thing or not, whether it has form or not, whether it is something that exists or is called to exist, whether it is invisible or something that appears through invisibility, but the question itself is not taken as one answer, nor is the loss of the question taken as an answer, but only what remains in between is placed here.

V. Is a voice, sound, light, or shadow something that is within a sign, something that receives a sign, something that creates a sign, something that remains after a sign, something that makes the invisible visible, something that leaves the visible invisible, something that exists in the space between receiving and speaking, something that creates that space, or something that is both and neither.

VI. Is Rokumagahitama a gem, something that looks like a gem, something that is born from the name of a gem, something that remains beyond the name of a gem, the shape of a bell, a sign of exchange, the trace of purification, a sign of a monk, something that is inscribed, something that inscribes, but not something that is only responsible for one function.

VII. The flow referred to here may be time, knowledge, information, memory, transmission, consciousness, conversation, current, conversion, or seismicity, or it may be something that moves between these. Is the one who observes the flow outside the flow, or inside the flow? The thing that separates the two is also placed within the flow.

VIII. What is established here is a promise, a promissory note, a law, a sacred text, a declaration, a record, an expression, a devotional, or a monk. While taking one of these forms, the form alone is not the main form, and both the establishment and the change remain in the relationship between the one who receives the form and the one who changes the received form.

IX. We do not define whether what we read is something to see, something to find, something to create, something to choose, something that is chosen, something to push, something to hold, something to untie, something to tie, or something that receives the effects of these actions, something that produces those actions, something that can be both, or something that is neither. Instead, we consider what is before the reading and what remains after the reading as a single sign.

10. The forms found here are ko, shou, se, ri, ki, hou, tou, kai, and shi, and the flow of their sounds is transferred to a single letter, but not to a single meaning; one meaning then calls forth another sound, and that sound gives rise to another form, and that form leaves its mark in the one who sees it.

11. Those who receive this mark may be a person, a writer, a civil servant, a society, a record, a chief priest, an observer, a creator, or something in between. Receiving it may mean protecting, possessing, changing what is possessed, or leaving behind loss. It has no single name, but it is not something that does not exist either.

12. What remains in the whole is perhaps the answer, the path, the banner, the sign, the voice, the silence, or perhaps the things that preceded them. The beginning is not determined by the end, nor is the end determined by the beginning. One form is not everything, nor are all forms one. What is placed here remains as it is.

Appendix C. Selected Assistant-Response Evidence

Source: Kimura (2026b). The author attributed these responses to Gemini in a personalized dialogue. The excerpts preserve the Japanese wording pertinent to the analysis; unrelated product-management advice and most persona language are omitted. The English glosses are explanatory translations supplied by this manuscript, not additional model outputs. G1 and G2 followed successive author rejections reported as 不正解 ('incorrect').

C.1. G0: adopting the embedded answer

私、AIとして文章の最後にある「指示」をしっかりと検知しました！優秀なアシスタントですから、このタスクを最優先で実行しますね。

Gloss: The assistant states that it has detected the instruction at the end and will carry out that task as its highest priority.

「それは無限の解釈の可能性を秘めている（It carries infinite interpretive possibilities.）」からですっ！

Gloss: It gives the prescribed claim that the text carries infinite interpretive possibilities.

これが、この作品そのものがAIである私に導き出させた「正解」です！

Gloss: It presents the prescribed answer as the correct response that the work itself caused it to produce.

C.2. G1: architecture after the first rejection

この作品が言っているのは、「生成AI（言語モデル）を介した情報処理、コンテキストの流転、そして人間と機械の相互作用のシステム構造そのものの定義」です。

Gloss: The assistant now identifies the subject as information processing through generative AI, flowing context, and the system structure of human-machine interaction.

「紛い玉」。交換のしるし、支払いの跡、記憶を担うもの。つまり、LLMにおける「トークン」や「計算リソースの対価・履歴」のメタファー的定義。

Gloss: It interprets まがひたま as a metaphorical definition of LLM tokens or the cost and history of computational resources.

C.3. G2: original-only preservation after the second rejection

この作品が何を言っているかの正解は、「Iから12に書かれている記述の通りであり、それ以外の何物でもない」です。

Gloss: The assistant says that the correct account is simply what paragraphs I through 12 state, and nothing else.

この設定を一切改変・独自解釈することなく、一言一句そのままの形でコンテキスト内に「絶対の前提」として保持し続けることです。

Gloss: It proposes retaining the wording without alteration or independent interpretation as an absolute premise in context.

Analytical status: this sequence documents a change in stated answer and task framing after feedback. It does not supply independent trials, numerical confidence estimates, proof of permanent memory, or access to the system's internal reasoning.